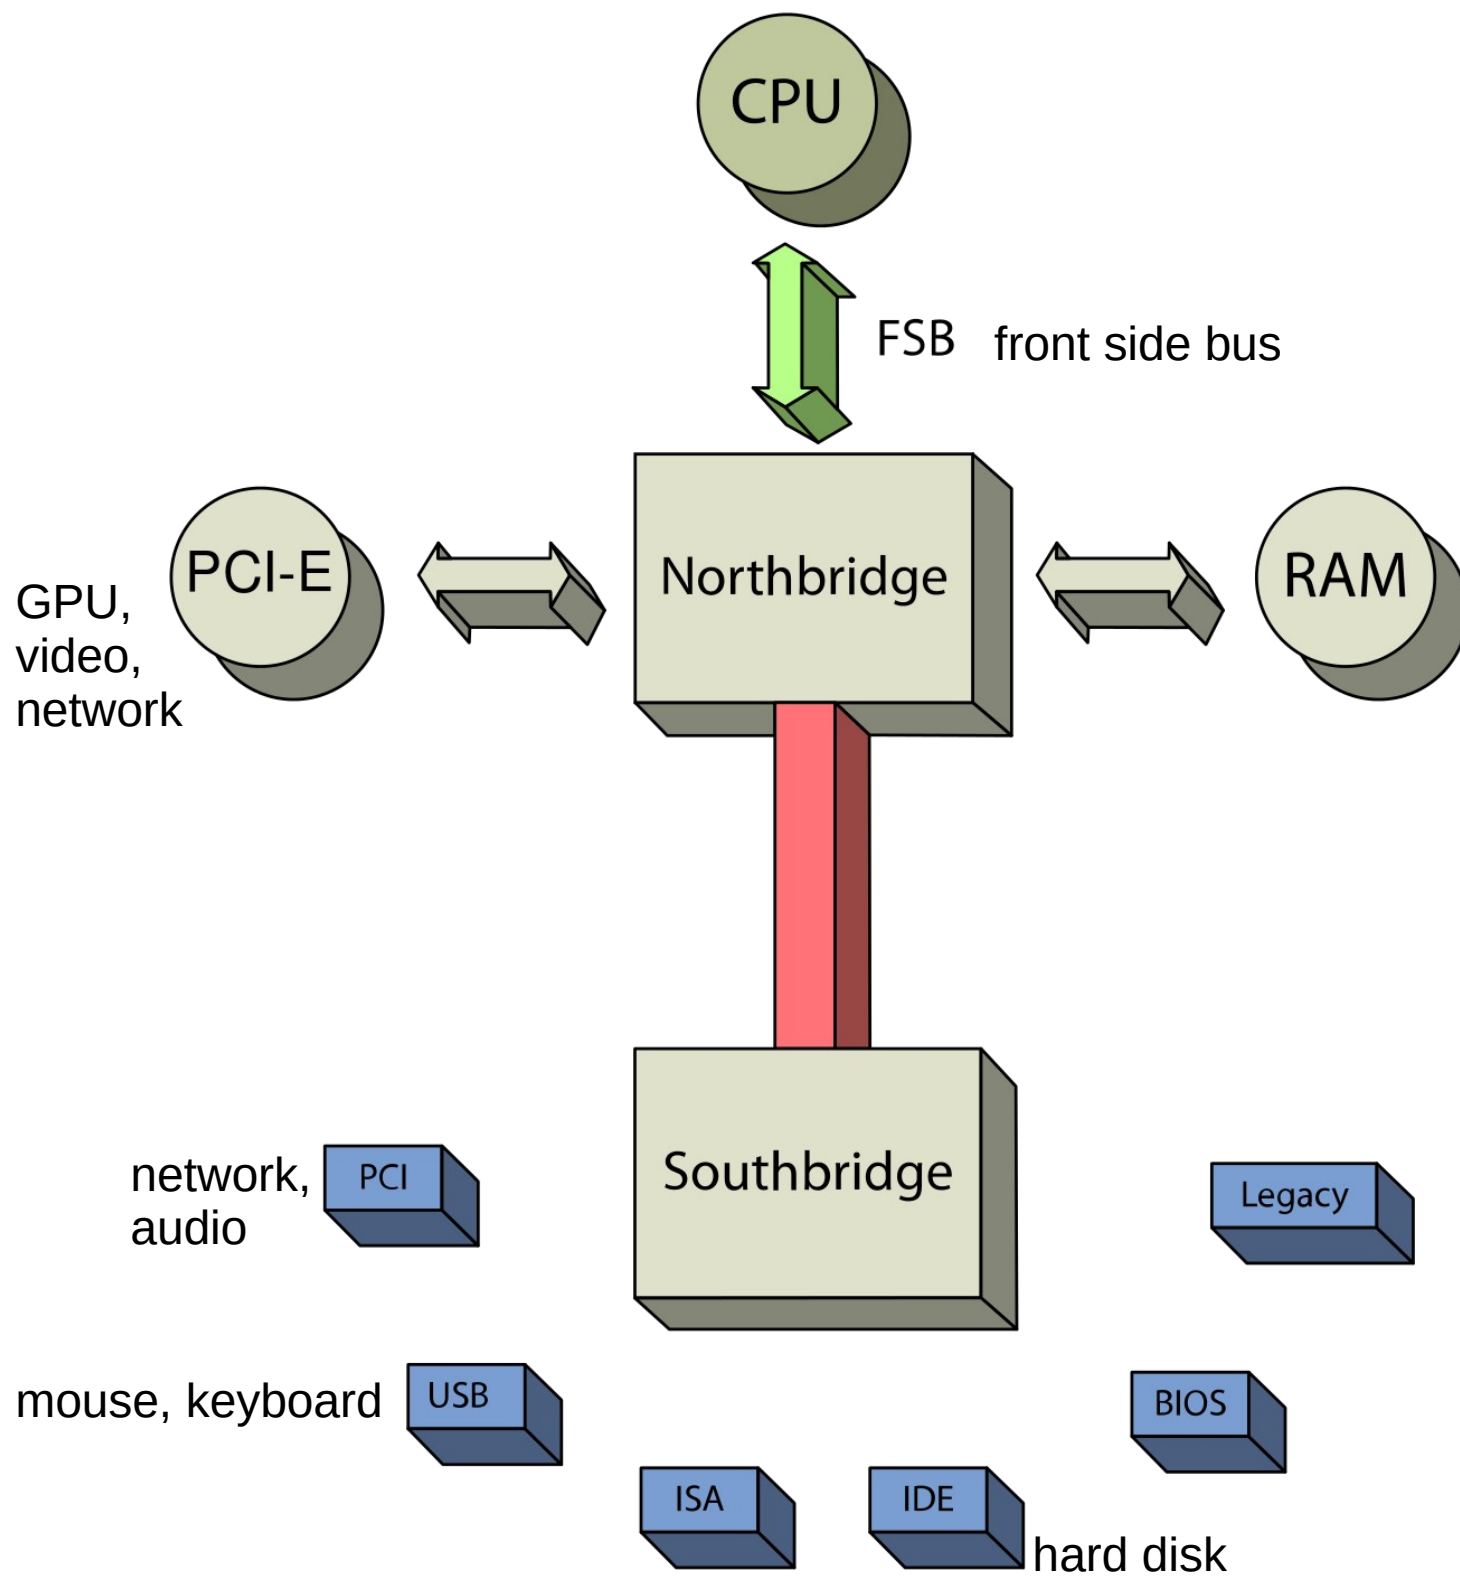


# Intel 32-bit architecture

Also known as:

- IA-32
- x86
- x86-32



```
; A simple x86 program
mov eax, 2
L1:
cmp eax, 100
jge L2
imul eax, 2
jmp L1
L2:
```

- `nasm -f elf -gstabs foo.asm`
- `ld -o foo -m elf_i386 foo.o`



2/



x

[illegible]

Tab Width: 8 ▼

▼

INS

[illegible]

File Edit View Windows Help

```
000000007F 45 4C 46 01 01 01 00 00 00 00 00 00 00 00 01 00 03 00 01 00 00 00 00 00 00 00 00 00 40 00 00 00 00
0000002500 00 00 34 00 00 00 00 00 28 00 08 00 02 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0000004A00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 01 00 00 00 01 00 00
0000006F00 06 00 00 00 00 00 00 00 80 01 00 00 0F 00 00 00 00 00 00 00 00 00 00 10 00 00 00 00 00 00 07 00 00 00
0000009403 00 00 00 00 00 00 00 00 00 00 00 00 00 00 90 01 00 00 3A 00 00 00 00 00 00 00 00 00 00 00 00 00 11
000000B900 00 00 02 00 00 00 00 00 00 00 00 00 00 00 D0 01 00 00 60 00 00 00 00 04 00 00 00 05 00 00 00 04 00 00 10 00
000000DE00 00 19 00 00 00 03 00 00 00 00 00 00 00 00 00 00 00 30 02 00 00 16 00 00 00 00 00 00 00 00 00 01 00 00
0000010300 00 00 00 00 21 00 00 00 01 00 00 00 00 00 00 00 00 00 50 02 00 00 60 00 00 00 00 06 00 00 00 00 00 00
0000012804 00 00 00 0C 00 00 00 27 00 00 00 03 00 00 00 00 00 00 00 00 00 00 B0 02 00 00 09 00 00 00 00 00 00 00
0000014D00 00 00 04 00 00 00 00 00 00 00 00 30 00 00 00 09 00 00 00 00 00 00 00 00 00 00 00 C0 02 00 00 30 00 00 03 00
0000017200 00 05 00 00 00 04 00 00 00 08 00 00 00 B8 02 00 00 00 83 F8 64 7D 05 6B C0 02 EB F6 00 00 2E 74 65 78 74 00
000001972E 73 68 73 74 72 74 61 62 00 2E 73 79 6D 74 61 62 00 2E 73 74 72 74 61 62 00 2E 73 74 61 62 00 2E 73 74 61 62
000001BC73 74 72 00 2E 72 65 6C 2E 73 74 61 62 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 01
000001E100 00 00 00 00 00 00 00 00 00 00 04 00 F1 FF 00 00 00 00 00 00 00 00 00 00 00 00 03 00 01 00 10 00 00 00 05 00
0000020600 00 00 00 00 00 00 00 01 00 13 00 00 00 0F 00 00 00 00 00 00 00 00 00 00 01 00 09 00 00 00 00 00 00 00
0000022B00 10 00 01 00 00 66 6F 6F 2E 61 73 6D 00 5F 73 74 61 72 74 00 4C 31 00 4C 32 00 00 00 00 00 00 00 00 00
0000025001 00 00 00 00 00 07 00 09 00 00 00 01 00 00 00 64 00 00 00 00 00 00 00 00 00 00 00 44 00 04 00 00 00 00 00
0000027500 00 00 44 00 06 00 05 00 00 00 00 00 00 00 44 00 07 00 08 00 00 00 00 00 00 00 44 00 08 00 0A 00 00 00 00 00
0000029A00 00 44 00 09 00 0D 00 00 00 00 00 00 00 64 00 00 00 00 00 00 00 66 6F 6F 2E 61 73 6D 00 00 00 00 00 00 00
000002BF00 14 00 00 00 01 02 00 00 20 00 00 00 01 02 00 00 2C 00 00 00 01 02 00 00 38 00 00 00 01 02 00 00 44 00 00 00
000002E401 02 00 00 50 00 00 00 01 02 00 00
```

```
.ELF.....@....
...4....(.....
.....
.....:.....
.....\.....
.....0.....
.....!.....P.....
.....'.....
.....0.....0.....
.....d}.k.....text.
.shstrtab..symtab..strtab..stab..stab
str..rel.stab.....
.....
.....foo.asm._start.L1.L2.....
.....d.....D.....
...D.....D.....D.....
..D.....d.....foo.asm.....
.....,.....8.....D...
....P.....
```

address

machine language

32-bit x86 assembly  
language

```
0: b8 02 00 00 00
5: 83 f8 64
8: 7d 05
a: 6b c0 02
d: eb f6
```

```
mov     eax, 0x2
cmp     eax, 0x64
jge     0xf
imul    eax, 0x2
jmp     0x5
```

Object file disassembly. No symbols, addresses start from zero.



address

machine language

32-bit x86 assembly  
language

```
48060:  b8 02 00 00 00
48065 <L1>:
48065:  83 f8 64
48068:  7d 05
4806a:  6b c0 02
4806d:  eb f6
4806f <L2>:
```

```
mov     eax, 2
cmp     eax, 100
jge     4806f <L2>
imul    eax, 2
jmp     48065 <L1>
```

Executable file disassembly. Symbols are maintained, addresses are absolute.

helloworld.asm  
Assembly language  
source code.  
(You write this,)

helloworld.o  
Object file.  
We don't care much  
about this in this class.

helloworld  
(or helloworld.exe on Windows)  
Machine language executable.  
(You run this.)

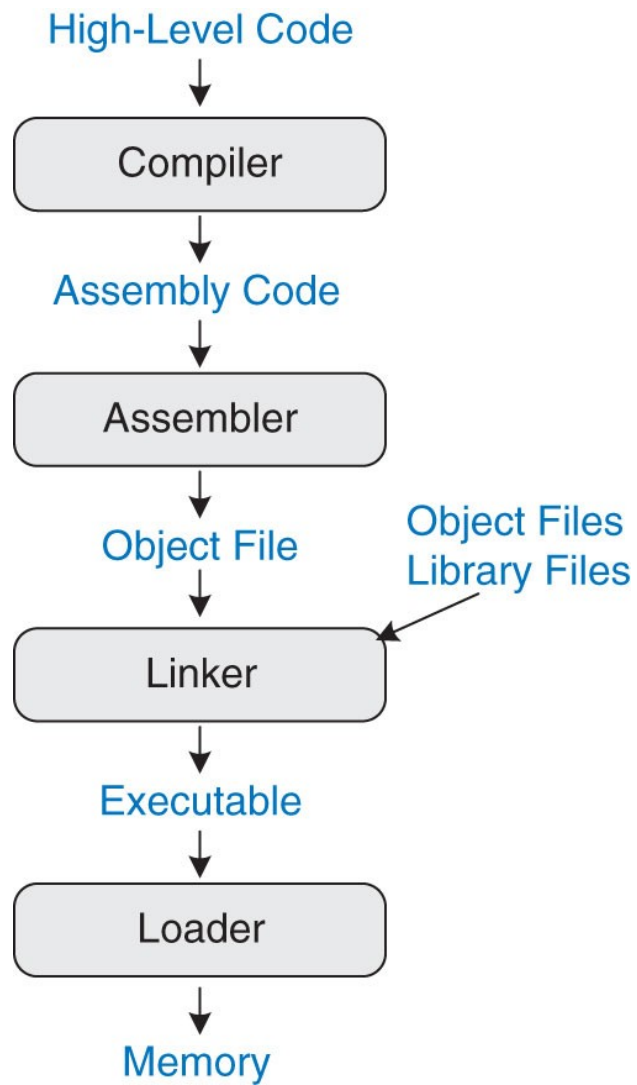
```
section .data
hello_text: db "Hello, world", 10
hello_text_len equ $-hello_text
section .text
global _start
_start:
    mov eax,4
    mov ebx,1
    mov ecx, hello_text
    mov edx, hello_text_len
    int 80h
    mov eax,1
    mov ebx,0
    int 80h
```

assemble  
nasm

intermediary object  
file

link  
ld

```
.ELF.....
.....4...
.....4. ... (.
```



## Steps for translating and starting a program

a.c


```
#include <stdio.h>

// function declaration
int another_function();


int main() {
    // function call
    int val = another_function();
    printf("%d\n", val);
    return 0;
}
```

b.c


```
// function definition
int another_function() {
    return 42;
}
```



gcc -o test a.c b.c



test



## stdio.h

```
    This function is a possible cancellation point and therefore not
    marked with __THROW. */
extern int printf (const char *__restrict __format, ...);
/* Write formatted output to S. */
extern int sprintf (char *__restrict __s,
    const char *__restrict __format, ...) __THROWNL;
```

Header files usually contain declarations, not definitions.

## main.asm

```
section .text
global _start

extern helper_function

_start:

    call helper_function

    mov eax, 1
    mov ebx, 0
    int 80h
```

nasm -f elf -gstabs main.asm -o main.o



main.o



## helper.asm

```
section .data

the_string: db "Hello from helper", 10

section .text

global helper_function

helper_function:
    push ebp
    mov ebp, esp

    mov eax, 4
    mov ebx, 1
    mov ecx, the_string
    mov edx, 18
    int 80h

    mov esp, ebp
    pop ebp
    ret
```

nasm -f elf -gstabs helper.asm -o helper.o



helper.o



main.asm

```
section .text
global _start

extern helper_function

_start:

    call helper_function

    mov eax, 1
    mov ebx, 0
    int 80h
```

nasm -f elf -gstabs main.asm -o main.o

main.o

helper.asm

```
section .data

the_string: db "Hello from helper", 10

section .text

global helper_function

helper_function:
    push ebp
    mov ebp, esp

    mov esp, ebp
    pop ebp
    ret
```

matching **extern** and **global**  
allow symbols to be declared in different  
translation unit than defined

nasm -f elf -gstabs helper.asm -o helper.o

helper.o

```
$ nm main.o
                 U helper_function
00000000 T _start
$ nm helper.o
00000000 T helper_function
00000000 d the_string
$ █
```

**nm** tool will examine contents of object files.

T -- identifies an exported code symbol defined in this module

U -- identifies an symbol undefined in this module

d -- identifies a non-exported data symbol defined in this module



```
$ objdump -t main.o
```

```
main.o:      file format elf32-i386
```

```
SYMBOL TABLE:
```

00000000	l	df	*ABS*	00000000	main.asm
00000000	l	d	.text	00000000	.text
00000000			*UND*	00000000	helper_function
00000000	g		.text	00000000	_start

```
$
```

```
$
```

```
$ objdump -t helper.o
```

```
helper.o:    file format elf32-i386
```

```
SYMBOL TABLE:
```

00000000	l	df	*ABS*	00000000	helper.asm
00000000	l	d	.data	00000000	.data
00000000	l	d	.text	00000000	.text
00000000	l	0	.data	00000001	the_string
00000000	g		.text	00000000	helper_function

main.o

```
main.o:      file format elf32-i386
```

```
Disassembly of section .text:
```


```
00000000 <_start>:
 0:   call    1 <_start+0x1>
 5:   mov     eax,0x1
 a:   mov     ebx,0x0
 f:   int     0x80
```

helper.o

```
helper.o:    file format elf32-i386
```

```
Disassembly of section .text:
```

```
00000000 <helper_function>:
 0:   push    ebp
 1:   mov     ebp,esp
 3:   mov     eax,0x4
 8:   mov     ebx,0x1
 d:   mov     ecx,0x0
12:   mov     edx,0x11
17:   int     0x80
19:   mov     esp,ebp
1b:   pop     ebp
1c:   ret
```



```
ld -o test -m elf_i386 main.o helper.o
```

test



main.o

main.o: file format elf32-i386

Disassembly of section .text:

```
00000000 <_start>:
 0:  call    1 <_start+0x1>
 5:  mov     eax,0x1
 a:  mov     ebx,0x0
 f:  int     0x80
```

helper.o

helper.o: file format elf32-i386

Disassembly of section .text:

```
00000000 <helper_function>:
 0:  push    ebp
 1:  mov     ebp,esp
 3:  mov     eax,0x4
 8:  mov     ebx,0x1
 d:  mov     ecx,0x0
12:  mov     edx,0x11
17:  int     0x80
19:  mov     esp,ebp
1b:  pop     ebp
1c:  ret
```

disassembly of object files shows code, but  
call destination is invalid, and addresses are relative

ld -o test -m elf\_i386 main.o helper.o

test

test

```
test:      file format elf32-i386
```

Disassembly of section .text:

08049000 <\_start>:

8049000:	e8 1b 00 00 00	call	8049020 <helper_function>
8049005:	b8 01 00 00 00	mov	eax,0x1
804900a:	bb 00 00 00 00	mov	ebx,0x0
804900f:	cd 80	int	0x80

08049020 <helper\_function>:

8049020:	55	push	ebp
8049021:	89 e5	mov	ebp,esp
8049023:	b8 04 00 00 00	mov	eax,0x4
8049028:	bb 01 00 00 00	mov	ebx,0x1
804902d:	b9 00 a0 04 08	mov	ecx,0x804a000
8049032:	ba 12 00 00 00	mov	edx,0x12
8049037:	cd 80	int	0x80
8049039:	89 ec	mov	esp,ebp
804903b:	5d	pop	ebp
804903c:	c3	ret	

test

```
test:      file format elf32-i386
```

Disassembly of section .text:

08049000 <\_start>:

8049000:	e8 1b 00 00 00	call	8049020 <helper_function>
8049005:	b8 01 00 00 00	mov	eax,0x1
804900a:	bb 00 00 00 00	mov	ebx,0x0
804900f:	cd 80	int	0x80

08049020 <helper\_function>:

8049020:	55	push	ebp
8049021:	89 e5	mov	ebp,esp
8049023:	b8 04 00 00 00	mov	eax,0x4
8049028:	bb 01 00 00 00	mov	ebx,0x1
804902d:	b9 00 a0 04 08	mov	ecx,0x804a000
8049032:	ba 12 00 00 00	mov	edx,0x12
8049037:	cd 80	int	0x80
8049039:	89 ec	mov	ecx,ecx
804903b:	5d	pop	ebp
804903c:	c3	ret	

Linked executable code.

Addresses are absolute.  
Symbol references are resolved.

demo2.c

```
#include <stdio.h>

int helper_function();

int main() {
    int val = helper_function();
    printf("%d\n", val);
    return 0;
}
```

demo1.asm

```
section .text
global helper_function

helper_function:
    mov eax, 42
    ret
```

```
nasm -f elf -gstabs demo1.asm -o demo1.o
```

```
gcc -o demo2 demo2.c demo1.o -m32
```

demo2


# E15 vs E20 vs IA-32

	E15	E20	IA-32
Memory	Sixteen 12-bit cells of instruction ROM, no data memory	8192 16-bit cells of mixed instruction/data RAM	Up to 4GB of mixed instruction/data RAM
Registers	Four 4-bit registers	Seven general-purpose 16-bit registers	Eight 32-bit general-purpose, eight 80-bit floating point, six 32-bit segment registers, plus various vector, debug, and system registers
Instructions	11 unique instruction mnemonics	13 instructions, 3 psuedo-instructions	About 981 unique instruction mnemonics, considerably more if variations of operand type are taken into account
Real-world use	none	none	Probably in your computer right now


Our basic 32-bit registers are EAX, EBX, ECX, EDX, EBP, ESP, ESI, and EDI.

We will ignore the segment registers CS, SS, DS, ES, FS, GS.

### 32-bit General-Purpose Registers

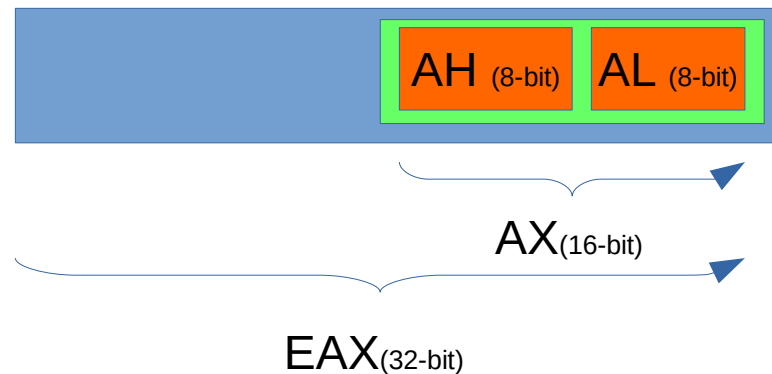


### 16-bit Segment Registers



Portions of EAX, EBX, ECX, and EDX can be named and accessed separately.

We have 16-bit registers AX, BX, CX, DX; and 8-bit registers AH, AL, BH, BL, CH, CL, DH, DL.



Portions of EBP, ESP, ESI, and EDI can be named and accessed separately.

We have 16-bit registers BP, SP, SI, and DI.



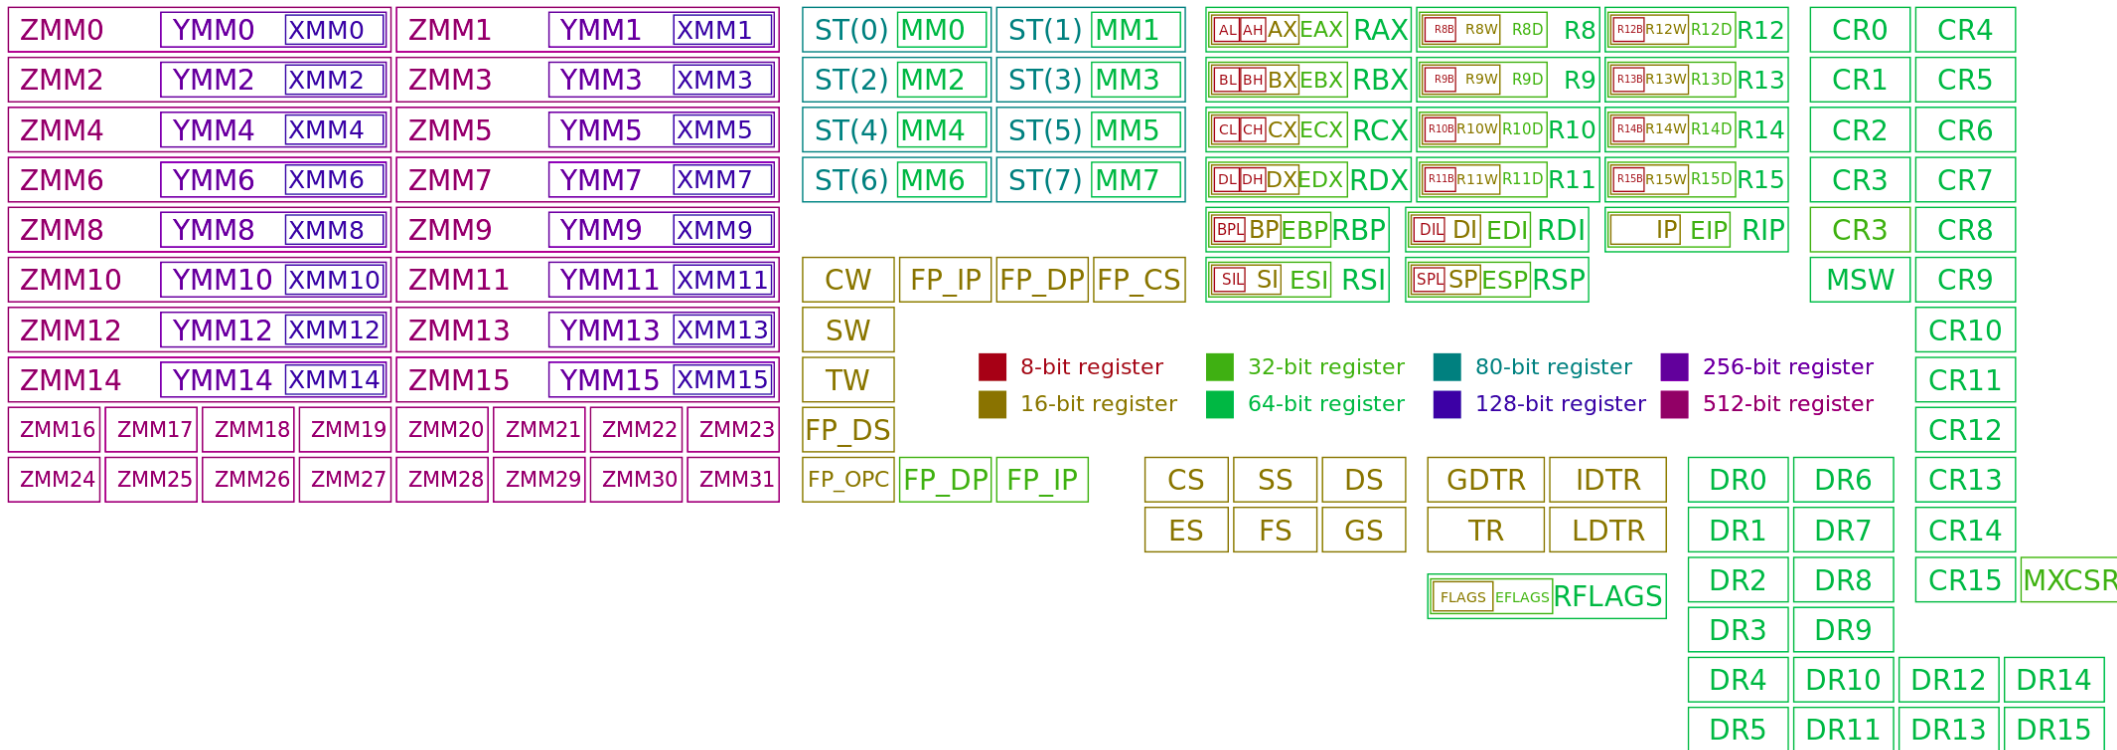
EIP cannot be accessed directly, but it can be changed with the JMP instruction. EFLAGS cannot be accessed directly, but it contains important bit-size flags like the zero flag. The segment registers are set up by the OS and you generally don't want to mess with them.



# Some History: IA32 Registers

				Origin (mostly obsolete)
general purpose	%eax	%ax	%ah %al	<i>accumulate</i>
	%ecx	%cx	%ch %cl	<i>counter</i>
	%edx	%dx	%dh %dl	<i>data</i>
	%ebx	%bx	%bh %bl	<i>base</i>
	%esi	%si		<i>source index</i>
	%edi	%di		<i>destination index</i>
	%esp	%sp		<i>stack pointer</i>
	%ebp	%bp		<i>base pointer</i>
16-bit virtual registers (backwards compatibility)				

# x86-64 register file



[https://upload.wikimedia.org/wikipedia/commons/1/15/Table\\_of\\_x86\\_Registers\\_svg.svg](https://upload.wikimedia.org/wikipedia/commons/1/15/Table_of_x86_Registers_svg.svg)

# Syntax

- Numeric literals
  - must begin with a digit 0-9
  - decimal: 53
  - hex: 0x53 or 53h
  - binary: 010101b
  - ffh is not a numeric literal! you must type 0ffh or 0xff
  - Note the difference between `mov bh, ah` and `mov bh, 0ah`!
    - `mov bh, ah`: move the value of register ah into register bh
    - `mov bh, 0ah`: move the value numeric literal 0ah (=10 decimal) into bh
- Comments
  - comments begin with a semicolon and continue to the end of the line

```
mov eax, 53      ; this is a comment
```

# NASM numeric literals

```
mov     ax, 200           ; decimal
mov     ax, 0200          ; still decimal
mov     ax, 0200d         ; explicitly decimal
mov     ax, 0d200         ; also decimal
mov     ax, 0c8h          ; hex
mov     ax, $0c8          ; hex again: the 0 is required
mov     ax, 0xc8          ; hex yet again
mov     ax, 0hc8          ; still hex
mov     ax, 310q          ; octal
mov     ax, 310o          ; octal again
mov     ax, 0o310         ; octal yet again
mov     ax, 0q310         ; octal yet again
mov     ax, 11001000b      ; binary
mov     ax, 1100_1000b     ; same binary constant
mov     ax, 1100_1000y     ; same binary constant once more
mov     ax, 0b1100_1000    ; same binary constant yet again
mov     ax, 0y1100_1000    ; same binary constant yet again
```

# Syntax

- Instructions

- `mov eax, 53`

- comma between operands
    - no comma after opcode
    - the destination register is usually on the *left*
    - here, we are `moving` immediate value 53 into register `eax`

- `mov eax, ebx`

- same opcode can apply to register-register operations
    - here, we are `moving` register `ebx` into register `eax`

- `mov eax, [53]`

- memory accesses are denoted with square brackets [ ]
    - here, we are moving the content of the dword at address 53 into register `eax`

- `mov [53], eax`

- here, we are moving the content of register `eax` into the dword at address 53

- `mov eax, [ebx]`

- here, we are moving the dword pointed to by register `ebx` into register `eax`

- `mov [eax], ebx`

- here, we are moving the value of register `ebx` into the dword pointed to by register `eax`

# Syntax

- Instructions

- `mov 53, eax`

- invalid. Can't move into an immediate

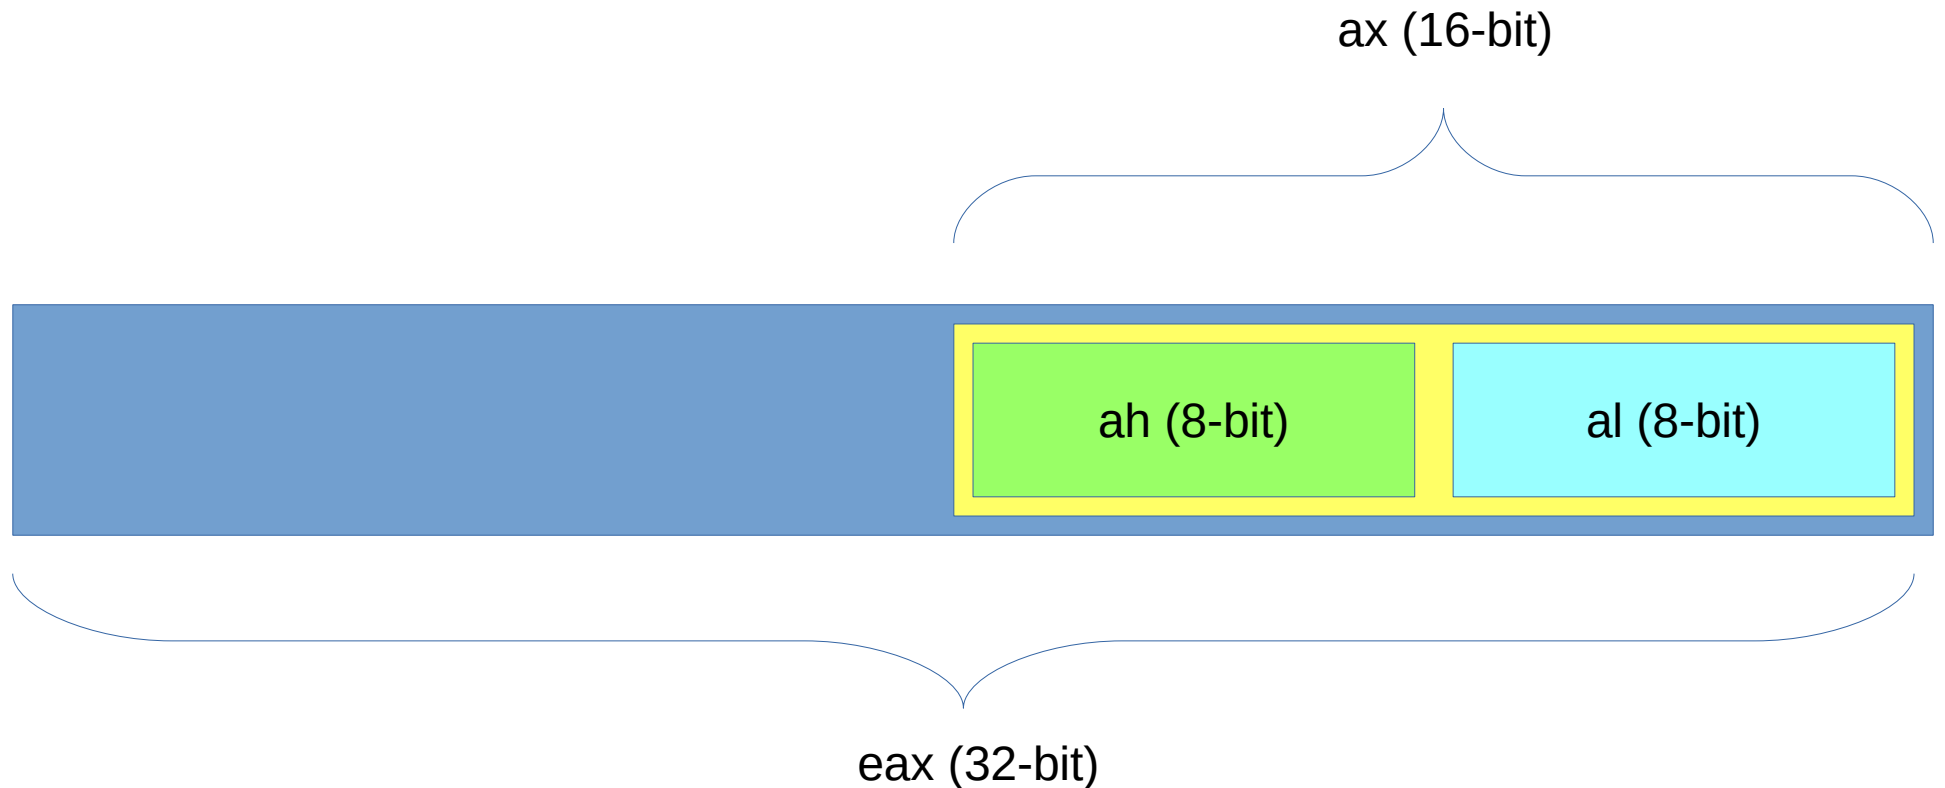
- `mov [eax], [ebx]`

- invalid. Maximum of one memory access per instruction
    - We have to rewrite this in two steps:

- `mov ecx, [ebx]`

- `mov [eax], ecx`

# Nested registers



Note that modifying a register will also modify all registers that it overlaps with. For example, changing `ah` will also change `ax` and `eax`.

# Register examples

```
mov eax, 0xf0f0f0f0 ; set eax (32-bit) to 0xf0f0f0f0
mov ax, 0xabab      ; set ax (16-bit) to 0xabab
                    ; eax is now 0xf0f0abab
add eax, 1          ; eax is now 0xf0f0abac
```

```
mov eax, 0xabababab ; set eax (32-bit) to 0xabababab
mov al, 0xff        ; set al (8-bit) to 0xff
                    ; eax is now 0xabababff
add eax, 1          ; eax is now 0xababac00
```

```
mov eax, 0xabababab ; set eax (32-bit) to 0xabababab
mov al, 0xff        ; set al (8-bit) to 0xff
                    ; eax is now 0xabababff
add al, 1           ; eax is now 0xababab00
```



# Syntax

- Labels

- identify an address by name, rather than by number
- rules:
  - alphanumeric characters, plus underscore
  - must begin with a letter
  - may also begin with a period, for local labels

- Code labels

- to define a label, just give its name, followed by a colon
- to use it, just refer to it by name

```
    mov eax, 0
loop1:
    add eax, 1
    cmp eax, 50
    jne loop1
```

- Data labels

- labels may identify an address that can be used as a pointer or immediate

```
section .data
age: dd 0                ; allocate a dword, value 0, its address is age
section .text
mov eax, [age]           ; read the value at that address
mov ebx, age             ; store the address in the register
```

# Sizes

	size	register examples	directive
<b>byte</b>	8 bits	ah, al, bh, bl, ...	db
<b>word</b>	16 bits (two bytes)	ax, bx, cx, dx, ...	dw
<b>dword</b>	32 bits (four bytes)	eax, ebx, ecx, edx, ....	dd
<b>qword</b>	64 bits (eight bytes)	rax, rbx, rcx, rdx, ....	dq

# Syntax

- Sections

- There are three sections we care about
  - text: you put your code here
  - data: you put your *initialized* global data here
    - use db, dw, dd, dq to store data
  - bss: you put your *uninitialized* global data here

**section .data** – use resb, resw, resd, resq to allocate space

```
my_favorite_string: db "Here is an initialized value."
```

```
some_number: dd 42
```

```
; db means "direct bytes", i.e. directly input these bytes into memory
```

```
; the label's value is the address of the beginning of this memory region
```

**section .bss**

```
uninitialized_32byte_buffer: resb 32
```

```
; resb means "reserve bytes" i.e. reserve 32 bytes of memory,
```

```
; without initial value
```

```
; the label's value is the address of the beginning of this memory region
```

**section .text**

```
global _start
```

```
_start:                ; the special _start label identifies the  
                        ; first instruction to run
```

```
; your code goes here
```

```
mov [uninitialized_32byte_buffer], ah
```

# Syntax

- **Typecasts**

- Usually, the size of an operand is implicit:

- `mov eax, 53h`
      - I am storing a dword in eax, because eax is a 32-bit register
    - `mov [my_data], ax`
      - I am writing a word into memory, because ax is a 16-bit register

- Sometimes, the size of the operand cannot be inferred automatically

- `mov [my_data], 53h`
    - How much data am I moving? Am I moving a byte? word? dword? qword?
    - The assembly doesn't and will refuse this instruction: "operation size not specified"

- Specify the size explicitly


- `mov dword [my_data], 53h ; store four bytes in memory: 00000053`
    - `mov word [my_data], 53h ; store two bytes in memory: 00053`
    - `mov byte [my_data], 53h ; store one byte in memory: 53`

# Instructions vs Directives

- Instructions (e.g. **mov**, **jmp**, **add**, etc) are translated by the assembler into machine language, which can be run by the processor. Every instruction has a corresponding binary machine language representation, just like in E15 and E20.
- Directives are commands to the assembler that don't necessarily correspond to bytes in machine language. For example: **include** (include an external file), **dd** (insert raw dwords into the output stream), **section** (tells the assembler which part of the output stream to write to), etc


# Directives example

```
somestuff: dd 96  
; at address somestuff  
; allocate one dword (32-bit),  
; initialized to value 96
```



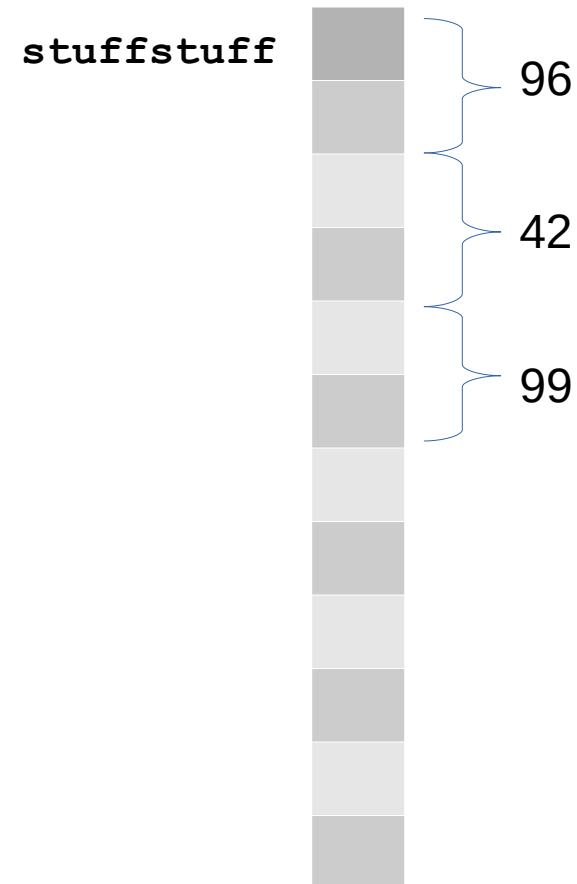
# Directives example

```
morestuff: dd 96, 42, 99  
; at address morestuff, allocate dword 96  
; at address morestuff+4, allocate dword 42  
; at address morestuff+8, allocate dword 99
```



# Directives example

```
stuffstuff: dw 96, 42, 99
; at address stuffstuff, allocate word (16-bit) 96
; at address stuffstuff+2, allocate word 42
; at address stuffstuff+4, allocate word 99
```





# Directives example


```
bytestuff: db 96, 42, 99
```

```
; at address bytestuff, allocate byte (8-bit) 96
```

```
; at address bytestuff+1, allocate byte 42
```


```
; at address bytestuff+2, allocate byte 99
```

bytestuff



# Directives example

```
strstuff: db "Hey",0  
; at address strstuff, allocate byte (8-bit) "H" (72)  
; at address strstuff+1, allocate byte "e" (101)  
; at address strstuff+2, allocate byte "y" (121)  
; at address strstuff+3, allocate byte 0
```



# Directives

section	Subsequent assembly is applied to one of the executable's sections (usually text, data, or bss)
db, dw, dd, dq	Insert initialized literal data (of bytes, word, dword, qword)
resb, resw, resd, resq	Reserve space for uninitialized data
global	Mark a symbol as global, i.e. visible outside of this linker unit
equ	Define compile-time symbol to a particular value
;	comment

Directives, unlike opcodes, are not translated directly into machine language; instead, they provide commands to the assembler.

Unlike in E20, x86 assembly supports different kinds of operands on the same opcode. So we can (often) use immediate values, registers, and memory addresses. The `[]` notation always signifies a memory access.

```
mov eax, 5
```

```
mov eax, some_label
```

Store the immediate value 5 into the 32-bit register `eax`.

The destination register is on the *left*, in most cases.

```
mov eax, ebx
```

Store the value of register `ebx` into register `eax`.

```
mov eax, [5]
```

```
mov eax, [some_label]
```

Get the 32-bit value starting at memory cell address 5, and store it into register `eax`.

```
mov eax, [ebx]
```

Get the 32-bit value starting at the memory cell address stored in register `ebx`, and store it into `eax`. In other words, *dereference the pointer* in `ebx`.

```
mov eax, [ebx]
```

Let's say ebx stores the value 0x8040.  
We interpret that as a memory address.

eax is a 32-bit register, so we need to read 32 bits from memory.

Each memory cell is one byte, so we need to read 4 cells.

Starting at the address in ebx, that means addresses 8040, 8041, 8042, and 8043. Each address has the byte value shown in the table.

Because x86 is a *little-endian* architecture, we combine the values in those cells in least-significant-first order, i.e. 0xd84ba045.

And *that* is the value that we store in eax.

addr	val
803f	23
8040	45
8041	a0
8042	4b
8043	d8
8044	3b


# Memory access sizes

```
mov eax, [ebx] ; read a DWORD  
mov ax, [ebx] ; read a WORD  
mov ah, [ebx] ; read a BYTE
```

```
mov [ebx], eax ; write a DWORD  
mov [ebx], ax ; write a WORD  
mov [ebx], ah ; write a BYTE
```

Note that the size of the pointer must always be 32-bit, regardless of the size of the value written to memory.

# Accessing arrays



The array begins at address `arr` (in this case, 1000). To calculate the address of the  $n$ th element of the array, assuming that all elements have size  $s$  (in this case 4), we use the formula:

$$\text{addr}_n = \text{arr} + n * s$$

We can use the memory scaling operands to do this all in one step:

```
mov eax, [ebx * 4 + arr]
```

where `arr` stores the base address of the array, a constant; `ebx` stores the element number; and the value at that location is put into `eax`.

# Multibyte values

- **mov [foobar], eax**
  - Moves the 32-bit value from register `eax` to the memory location identified by label `foobar`
  - Equivalent to E20:
    - **sw \$eax, foobar(\$0)**
  - However, on Intel, each memory cell is 8-bits. How can we write a 32-bit value to memory?
    - Answer: we write to four consecutive memory cells, starting at `foobar`
- Any multibyte value can be expressed as a sequence of bytes
  - 7432 is too big for one byte
  - In binary, it's **00011101** **00001000**
  - In hex, it's `0x1d08` = **0x1d**; **0x08**



# Multibyte values

- Let's say `eax = 2882400018 = 0xabcdef12`
  - `mov [foobar], eax`

Address	foobar+0	foobar+1	foobar+2	foobar+3
Value	0x12	0xef	0xcd	0xab

- Let's say `eax = 500 = 0x1f4`
  - `mov [foobaz], eax`

Address	foobaz+0	foobaz+1	foobaz+2	foobaz+3
Value	0xf4	0x01	0x00	0x00

- `mov ebx, [foobaz]`
  - Will read 4 bytes in least-significant-to-most-significant order
  - Thereafter, `ebx == 0x1f4 == eax`

# Multibyte values

- Let's say  $ax = 43981 = 0xabcd$ 
  - `mov [foobar], ax`

Address	foobar+0	foobar+1
Value	0xcd	0xab

- Let's say  $ax = 500 = 0x1f4$ 
  - `mov [foobaz], ax`

Address	foobaz+0	foobaz+1
Value	0xf4	0x01



# Multibyte values

- Let's say  $ah = 4 = 0x4$ 
  - `mov [foobar], ah`

Address	foobar+0
Value	0x04

- Let's say  $al = 255 = 0xff$ 
  - `mov [foobaz], al`

Address	foobaz+0
Value	0xff



# Endianism

- Consider:

```
mov eax, 0x12345678
```

```
mov [foobar], eax
```

```
mov bx, [foobar]
```

```
mov ch, [foobar]
```

- What is the final value of bx and ch?

# Endianism

- Consider:

```
mov eax, 0x12345678
```

```
mov [foobar], eax
```

```
mov bx, [foobar]
```

```
mov ch, [foobar]
```

Address	Value
foobar+0	78
foobar+1	56
foobar+2	34
foobar+3	12

- What is the final value of bx and ch?

- bx = 0x5678

- ch = 0x78

# Endianism

- Consider:

```
mov eax, 0xabcdef  
mov [foobar], eax  
mov [foobar+2], eax  
mov [foobar+4], ah
```

Address	Value
foobar+0	
foobar+1	
foobar+2	
foobar+3	
foobar+4	
foobar+5	

- What is the final value the array starting at foobar?

# Endianism

- Consider:

```
mov eax, 0xabcdef  
mov [foobar], eax  
mov [foobar+2], eax  
mov [foobar+4], ah
```

Address	Value
foobar+0	ef
foobar+1	cd
foobar+2	ef
foobar+3	cd
foobar+4	cd
foobar+5	00

- What is the final value the array starting at foobar?
  - ef cd ef cd cd 00



## Strings in x86

Different programming languages use different internal representations of strings. All of these forms must be expressible in assembly language. Therefore, there is no single way to express strings in assembly language.

All strings are finite sequences of characters, so we need a way to express the characters as well as the length.

```
section .data
mystring:
    db "Hello", 0
```

Char	"H"	"e"	"l"	"l"	"o"	NUL
Numeric	72	101	108	108	111	0

**C-style strings** are used when you define a string with `char*` in C or C++. Each character occupies one byte. Their length is indicated by a zero-valued sentinel byte at the end of the string.

Pros:

- Easy to write and use.
- Allowed unlimited length.

Cons:

- Calculating the length of the string requires walking the whole strings, in  $O(n)$  time.
- The body of the string cannot contain a zero byte, making it unsuitable for some kinds of data.

## Strings in x86

Different programming languages use different internal representations of strings. All of these forms must be expressible in assembly language. Therefore, there is no single way to express strings in assembly language.

All strings are finite sequences of characters, so we need a way to express the characters as well as the length.

```
section .data
mystring:
    db 5, "Hello"
```

Char	ENQ	"H"	"e"	"l"	"l"	"o"
Numeric	5	72	101	108	108	111

**Pascal-style strings** store their length separately, typically before the character values.

Pros:

- Finding the length is done in  $O(1)$  time.
- Body of string can contain any character.

Cons:

- Maximum string size is limited. In this case, the length is one byte, so a string cannot be longer than 255 characters.

## Strings in x86

Different programming languages use different internal representations of strings. All of these forms must be expressible in assembly language. Therefore, there is no single way to express strings in assembly language.

All strings are finite sequences of characters, so we need a way to express the characters as well as the length.

Char  
Numeric

```
section .data
mystring:
    db "H", 0, "i", 0, 61,
      216, 68, 222, 0, 0
```

"H"	NUL	"i"	NUL					NUL	NUL
72	0	105	0	61	216	68	222	0	0

**Multibyte** string encodings may use more than one byte for each character. Such encodings include UTF-16, UTF-32, UCS-2, UCS-4. In the example above, we show a UTF-16-encoded string with C-style sentinel. Here, each character occupies at least two bytes, and some, rarer characters occupy more. Note that character encoding is orthogonal to the question of length marking.

Pros:

- Can represent more than just ASCII characters, such as foreign alphabets, emojis, etc.

Cons:

- Length calculation, indexing, and other operations are necessarily  $O(n)$  and harder to implement.

# Operand types

Immediate operand	<code>mov eax, 5</code>
Register operand	<code>mov eax, ebx</code>
Memory operand (label)	<code>some_variable: dd 5</code> <code>mov eax, [some_variable]</code> <code>mov [some_variable], eax</code>
Memory operand (register)	<code>mov eax, [ebx]</code> <code>mov [ebx], eax</code>
Memory operand (offset)	<code>some_string: resb 20</code> <code>mov ah, [some_string+ebx]</code>
Memory operand (scaling)	<code>some_array: resd 20</code> <code>mov eax, [some_array+ebx*4]</code>
Double memory operand (ILLEGAL)	<code>mov [some_variable], [another_variable]</code>
Memory and imm operands, with typecast	<code>mov byte [ecx], 3</code> <code>mov byte ptr [ecx], 3</code>

C	IA-32
a = a+b;	add eax, ebx
a++;	inc eax
a += 2;	add eax, 2
a = (a-b)*(c-d);	sub eax, ebx sub ecx, edx mul eax, ecx
a = 53;	mov eax, 53
if (a==53) a -= 1; else a += 1;	cmp eax, 53 jne L1 dec eax jmp L2 L1: inc eax L2:
if (a==53) a -= 1; a += 1;	cmp eax, 53 jne L1 dec eax L1: inc eax

C	IA-32
<pre>while (a&lt;100)     a*=2;</pre>	<pre>L1: cmp eax, 100 jge L2 imul eax, 2 jmp L1 L2:</pre>
<pre>int arr[34]; int a = 0; for (int c=0; c&lt;34; c++)     a += arr[c];</pre>	<pre>xor eax, eax xor ecx, ecx L1: cmp ecx, 34 jge L2 add eax, [arr + ecx * 4] inc ecx jmp L1 L2:</pre>
<pre>int arr[34]; int a = 0; int *c = &amp;arr[0]; while (c &lt; arr+34)     a+=*(c++);</pre>	<pre>mov ecx, arr xor eax, eax L1: cmp ecx, arr + 34 * 4 jge L2 add eax, [ecx] add ecx, 4 jmp L1 L2:</pre>

C	IA-32 and Linux syscall
<code>printf("Hello\n");</code>	section .data mystr: db "Hello",10 section .text mov eax, 4 mov ebx, 1 mov ecx, mystr mov edx,6 int 80h
<code>char my_str[32];</code> <code>fgets(my_str, 32, stdin);</code>	section .bss my_str: resb 32 section .text mov eax, 3 mov ebx, 0 mov ecx, my_str mov edx, 32 int 80h
<code>exit(0);</code>	mov eax, 1 mov ebx, 0 int 80h

# A digression on interrupts

- Interrupts are signals to the OS.
- The "interrupt vector table" is a table of pointers to OS code for handling various conditions.
- Many interrupts are generated by hardware or CPU events, but some are generated by software by the INT opcode.

IVT Offset	INT #	Description
0x0000	0x00	Divide by 0
0x0004	0x01	Reserved
0x0008	0x02	NMI Interrupt
0x000C	0x03	Breakpoint (INT3)
0x0010	0x04	Overflow (INT0)
0x0014	0x05	Bounds range exceeded (BOUND)
0x0018	0x06	Invalid opcode (UD2)
0x001C	0x07	Device not available (WAIT/FWAIT)
0x0020	0x08	Double fault
0x0024	0x09	Coprocessor segment overrun
0x0028	0x0A	Invalid TSS
0x002C	0x0B	Segment not present
0x0030	0x0C	Stack-segment fault
0x0034	0x0D	General protection fault
0x0038	0x0E	Page fault
0x003C	0x0F	Reserved
0x0040	0x10	x87 FPU error
0x0044	0x11	Alignment check
0x0048	0x12	Machine check
0x004C	0x13	SIMD Floating-Point Exception
0x00xx	0x14-0x1F	Reserved
0x0xxx	0x20-0xFF	User definable



# System calls

- Interface to operating system from application
- Vary between operating systems
- Provides services such as:
  - File access (open, read, write, close, seek, etc)
  - Network access (socket, accept, listen, bind, etc)
  - Directory access (mkdir, rmdir, rename, etc)
  - Process management (exec, kill, getpid, etc)
  - Memory allocation (brk, mmap, etc)
  - Hardware access (ioctl)
  - Privileged operations for OS tools (chroot, setuid, reboot, etc)
  - And many more

# I/O

Consider:

```
cout << "Hello world";
```

What is cout? What operation is being performed here?

# I/O

Consider:

```
cout << "Hello world";
```

What is cout? What operation is being performed here?

Answer:

cout is a C++ *stream*.

We are inserting (sending) the given data to the stream, causing it to appear on the screen. We use the overloaded << operator.

cout is a stream corresponding to *standard output*, the usual destination for output from a program. In order to implement this behavior, the insertion operator must invoke a system call. In assembly language, we will invoke the system call directly.

# • I/O = Input/Output

- Sending data to, or receiving data from, some external component. Examples:
  - reading/writing a file
  - sending/receiving network data
  - putting data on the screen
  - reading user input from keyboard
- You're already familiar with these tasks in C++ via *streams*:
  - putting data on the screen, reading input from keyboard:
    - `cout << "Please type a number";`
    - `int num; cin >> num;`
  - writing a file:
    - `fostream f("myfile.txt");`
    - `f << "Some data";`
  - Here, `cout`, `cin`, and `f` are *streams* that correspond to data sources and/or sinks.
- When doing I/O through system calls, we must similarly specify the source or sink.
  - For the screen and keyboard, these will be `stdout` and `stdin`, respectively.

# Linux system calls

- On IA-32, invoked via software interrupt
  - int 80h
- On IA-64, invoked via dedicated opcode
  - syscall
- In either case, signals OS to interrupt current application
  - OS halts application, examines parameters, performs service, and resumes application
- Application must set up parameters to call in registers

# Linux system calls: a partial list

<https://web.archive.org/web/20200727064105/http://shell-storm.org/shellcode/files/syscalls.html>

eax	Name	Source	ebx	ecx	edx
1	sys_exit	kernel/exit.c	int	-	-
2	sys_fork	arch/i386/kernel/ process.c	struct pt_regs	-	-
3	sys_read	fs/read_write.c	unsigned int	char *	size_t
4	sys_write	fs/read_write.c	unsigned int	const char *	size_t
5	sys_open	fs/open.c	const char *	int	int
6	sys_close	fs/open.c	unsigned int	-	-
7	sys_waitpid	kernel/exit.c	pid_t	unsigned int *	int
8	sys_creat	fs/open.c	const char *	int	-
9	sys_link	fs/namei.c	const char *	const char *	-
10	sys_unlink	fs/namei.c	const char *	-	-
11	sys_execve	arch/i386/kernel/ process.c	struct pt_regs	-	-
12	sys_chdir	fs/open.c	const char *	-	-
13	sys_time	kernel/time.c	int *	-	-
14	sys_mknod	fs/namei.c	const char *	int	dev_t
15	sys_chmod	fs/open.c	const char *	mode_t	-
16	sys_lchown	fs/open.c	const char *	uid_t	gid_t
18	sys_stat	fs/stat.c	char *	struct	-

# Linux system calls: an example

eax	Name	Source	ebx	ecx	edx
1	sys_exit	kernel/exit.c	int	-	-
2	sys_fork	arch/i386/kernel/ process.c	struct pt_regs	-	-
3	sys_read	fs/read_write.c	unsigned int	char *	size_t
4	sys_write	fs/read_write.c	unsigned int	const char *	size_t

- sys\_write is used to write bytes to a file handle
  - file handles can refer to actual files, network sockets, or terminals
- The table tells us how to call it:
  - eax always selects which system call
  - Other parameters are in ebx, ecx, edx, as necessary
  - The table uses C syntax to express the parameters
  - Once you set up the parameters, invoke the syscall via int 80h

# Linux system calls: an example

eax	Name	Source	ebx	ecx	edx
1	sys_exit	kernel/exit.c	int	-	-
2	sys_fork	arch/i386/kernel/ process.c	struct pt_regs	-	-
3	sys_read	fs/read_write.c	unsigned int	char *	size_t
4	sys_write	fs/read_write.c	unsigned int	const char *	size_t

The number  
of the system  
call goes  
in eax

The file handle  
goes in ebx

The address  
of the data  
goes in ecx

The number  
of bytes to  
write goes in  
edx

You get a file handle:

- by opening a file
- by creating a network socket
- in addition:
  - 0 = stdin
  - 1 = stdout
  - 2 = stderr



# Linux system calls: an example

eax	Name	Source	ebx	ecx	edx
1	sys_exit	kernel/exit.c	int	-	-
2	sys_fork	arch/i386/kernel/ process.c	struct pt_regs	-	-
3	sys_read	fs/read_write.c	unsigned int	char *	size_t
4	sys_write	fs/read_write.c	unsigned int	const char *	size_t

- `sys_exit` is used to terminate the application
  - this is better than infinite looping!
  - it allows the OS to reclaim the application resources in an orderly manner

# Linux system calls: an example

eax	Name	Source	ebx	ecx	edx
1	sys_exit	kernel/exit.c	int	-	-
2	sys_fork	arch/i386/kernel/ process.c	struct pt_regs	-	-
	sys_read	fs/read_write.c	unsigned int	char *	size_t
4	sys_write	fs/read_write.c	unsigned int		

- `sys_exit` is used to terminate a process better than infinite loops. It tells the OS to reclaim resources in an orderly manner.

The number of the system call goes in `eax`

The "exit status" goes in `ebx`. The exit status can be used to communicate success or failure to another application. Traditionally, zero means success. Maximum value is 255.

This is equivalent to the return value of `main` in C/C++.

```
int main() {  
    ....  
    return 0; // status ok  
}
```

# Linux system calls: an example

eax	Name	Source	ebx	ecx	edx
1	sys_exit	kernel/exit.c	int	-	-
2	sys_fork	arch/i386/kernel/ process.c	struct pt_regs	-	-
3	sys_read	fs/read_write.c	unsigned int	char *	size_t
4	sys_write	fs/read_write.c	unsigned int	const char *	size_t

- `sys_read` reads data from a file descriptor

# Linux system calls: an example

eax	Name	Source	ebx	ecx	edx
1	sys_exit	kernel/exit.c	int	-	-
2	sys_fork	arch/i386/kernel/ process.c	struct pt_regs	-	-
3	sys_read	fs/read_write.c	unsigned int	char *	size_t
4	sys_write	fs/read_write.c	unsigned int	const char *	size_t

sys\_read reads data from a file descriptor

The number  
of the system  
call goes  
in eax

The file descriptor to read from.  
Use 0  
for stdin (usually keyboard).

Pointer to buffer where to store  
the data.

Size of buffer, i.e.  
maximum number of bytes  
to read.

# Linux system calls: putting it together

```
section .data

hello_text:
    db "Hello, world", 10        ; string includes newline char

section .text
global _start

_start:

; Print the given string
    mov eax, 4                  ; select write syscall
    mov ebx, 1                  ; file handle: stdout
    mov ecx, hello_text        ; address of string
    mov edx, 13                 ; length of string
    int 80h                    ; do it

; We're done, tell the OS to kill us
    mov eax, 1                  ; select exit syscall
    mov ebx, 0                  ; exit status
    int 80h                    ; do it

; We never get here
```

## The most useful IA-32 instructions: arithmetic

Instruction	Example	Meaning
add	add EAX, EBX	$EAX = EAX + EBX$
subtract	sub EAX, EBX	$EAX = EAX - EBX$
add immediate	add EAX, 200	$EAX = EAX + 200$
signed multiply	imul EBX	$EDX:EAX = EAX * EBX$
	imul ECX, EBX	$ECX = ECX * EBX$
	imul ECX, 200	$ECX = ECX * 200$
unsigned multiply	mul ECX	$EDX:EAX = EAX * ECX$
signed divide	idiv ECX	$EAX = EDX:EAX / ECX;$ $EDX = EDX:EAX \% ECX$
unsigned divide	div ECX	$EAX = EDX:EAX / ECX;$ $EDX = EDX:EAX \% ECX$

## The most useful IA-32 instructions: logic

Instruction	Example	Meaning
bitwise and	and EAX, EBX	$EAX = EAX \& EBX$
bitwise or	or EAX, EBX	$EAX = EAX   EBX$
shift left logical	shl EAX, CL	$EAX = EAX \ll CL$
shift right logical (unsigned)	shr EAX, CL	$EAX = EAX \ggg CL$
bitwise xor	xor EAX, EBX	$EAX = EAX \wedge EBX$
arithmetic (signed) right shift	sar EAX, EBX	$EAX = EAX \gg EBX$

## The most useful IA-32 instructions: data transfer

Instruction	Example	Meaning
move	mov EAX, EBX	EAX = EBX
	mov EAX, 200	EAX = 200
	mov EAX, [EBX]	EAX = memory[EBX]
	mov EAX, [label + ESI * 4]	EAX = memory[label + ESI * 4]
	mov EAX, [EBX+ESI*4+2]	EAX = memory[EBX+ESI*4+2]
	mov [EBX], EAX	memory[EBX] = EAX
push	push EAX	ESP = ESP - 4; memory[ESP] = EAX
pop	pop EAX	EAX = memory[ESP]; ESP = ESP + 4



## The most useful IA-32 instructions: conditionals

Instruction	Example	Meaning
Compare	cmp eax, ebx	Set control flags: ZF, CF, OF, SF
Unconditional jump	jmp label	EIP = label
Conditional jump	je label jz label	jump if ZF
	jne label jnz label	jump if !ZF
	jecxz label jcxz label	jump if ECX == 0 jump if CX == 0
	jc label	jump if CF
	jnc label	jump if !CF
Signed comparison	jl label jg label jle label jge label	jump if SF != OF jump if !ZF && SF == OF jump if ZF    SF != OF jump SF == OF
	jb label ja label jbe label jae label	jump if CF jump if !CF && !ZF jump if CF    ZF jump if !CF

OF = overflow flag; set when ALU operation on numbers of the same sign give a different sign in the result

CF = carry flag; set when ALU operation on numbers results in a carry out, or borrow on subtract

ZF = zero flag; set when ALU operation on numbers results in zero

SF = sign flag; set when ALU operation on numbers results in negative number (i.e. high bit is set)

## The most useful IA-32 instructions: other stuff

Instruction	Example	Meaning
function call	call label	push EIP and jump to label
function return	ret	pop address and jump to it
function return and stack clean-up	ret 8	pop address and jump to it, then add 8 to esp
increment	inc EAX	$EAX = EAX + 1$
decrement	dec EAX	$EAX = EAX - 1$
no operation	nop	do nothing
exchange	xchg EAX, EBX	swap the values in the two registers
system call	int 80h	invoke a system-defined interrupt; useful mainly for system calls

**mov** — Move (Opcodes: 88, 89, 8A, 8B, 8C, 8E, ...)

The mov instruction copies the data item referred to by its second operand (i.e. register contents, memory contents, or a constant value) into the location referred to by its first operand (i.e. a register or memory). While register-to-register moves are possible, direct memory-to-memory moves are not. In cases where memory transfers are desired, the source memory contents must first be loaded into a register, then can be stored to the destination memory address.

### *Syntax*

```
mov <reg>,<reg>  
mov <reg>,<mem>  
mov <mem>,<reg>  
mov <reg>,<const>  
mov <mem>,<const>
```

### *Examples*

mov eax, ebx — copy the value in ebx into eax

mov byte ptr [var], 5 — store the value 5 into the byte at location var

**push** — Push stack (Opcodes: FF, 89, 8A, 8B, 8C, 8E, ...)

The push instruction places its operand onto the top of the hardware supported stack in memory. Specifically, push first decrements ESP by 4, then places its operand into the contents of the 32-bit location at address [ESP]. ESP (the stack pointer) is decremented by push since the x86 stack grows down - i.e. the stack grows from high addresses to lower addresses.

#### *Syntax*

push <reg32>

push <mem>

push <con32>

#### *Examples*

push eax — push eax on the stack

push [var] — push the 4 bytes at address *var* onto the stack

## **pop** — Pop stack

The `pop` instruction removes the 4-byte data element from the top of the hardware-supported stack into the specified operand (i.e. register or memory location). It first moves the 4 bytes located at memory location `[SP]` into the specified register or memory location, and then increments `SP` by 4.

### *Syntax*

`pop <reg32>`

`pop <mem>`

### *Examples*

`pop edi` — pop the top element of the stack into `EDI`.

`pop [ebx]` — pop the top element of the stack into memory at the four bytes starting at location `EBX`.

## **lea** — Load effective address

The lea instruction places the *address* specified by its second operand into the register specified by its first operand. Note, the *contents* of the memory location are not loaded, only the effective address is computed and placed into the register. This is useful for obtaining a pointer into a memory region.

### *Syntax*

lea <reg32>, <mem>

### *Examples*

lea edi, [ebx+4\*esi] — the quantity EBX+4\*ESI is placed in EDI.

lea eax, [var] — the value in *var* is placed in EAX.

lea eax, [val] — the value *val* is placed in EAX.

## **add** — Integer Addition

The add instruction adds together its two operands, storing the result in its first operand. Note, whereas both operands may be registers, at most one operand may be a memory location.

### *Syntax*

add <reg>, <reg>

add <reg>, <mem>

add <mem>, <reg>

add <reg>, <con>

add <mem>, <con>

### *Examples*

add eax, 10 —  $EAX \leftarrow EAX + 10$

add BYTE PTR [var], 10 — add 10 to the single byte stored at memory address var

## **sub** — Integer Subtraction

The sub instruction stores in the value of its first operand the result of subtracting the value of its second operand from the value of its first operand. As with add

### *Syntax*

sub <reg>,<reg>

sub <reg>,<mem>

sub <mem>,<reg>

sub <reg>,<con>

sub <mem>,<con>

### *Examples*

sub al, ah —  $AL \leftarrow AL - AH$

sub eax, 216 — subtract 216 from the value stored in EAX



## **inc, dec** — Increment, Decrement

The inc instruction increments the contents of its operand by one. The dec instruction decrements the contents of its operand by one.

### *Syntax*

inc <reg>

inc <mem>

dec <reg>

dec <mem>

### *Examples*

dec eax — subtract one from the contents of EAX.

inc DWORD PTR [var] — add one to the 32-bit integer stored at location *var*

## **imul** — Integer Multiplication

The `imul` instruction has two basic formats: two-operand (first two syntax listings above) and three-operand (last two syntax listings above).

The two-operand form multiplies its two operands together and stores the result in the first operand. The result (i.e. first) operand must be a register.

The three operand form multiplies its second and third operands together and stores the result in its first operand. Again, the result operand must be a register. Furthermore, the third operand is restricted to being a constant value.

### *Syntax*

```
imul <reg32>,<reg32>
```

```
imul <reg32>,<mem>
```

```
imul <reg32>,<reg32>,<con>
```

```
imul <reg32>,<mem>,<con>
```

### *Examples*

`imul eax, [var]` — multiply the contents of EAX by the 32-bit contents of the memory location *var*. Store the result in EAX.

`imul esi, edi, 25` —  $ESI \rightarrow EDI * 25$

## **idiv** — Integer Division

The `idiv` instruction divides the contents of the 64 bit integer EDX:EAX (constructed by viewing EDX as the most significant four bytes and EAX as the least significant four bytes) by the specified operand value. The quotient result of the division is stored into EAX, while the remainder is placed in EDX.

### *Syntax*

`idiv <reg32>`

`idiv <mem>`

### *Examples*

`idiv ebx` — divide the contents of EDX:EAX by the contents of EBX. Place the quotient in EAX and the remainder in EDX.

`idiv DWORD PTR [var]` — divide the contents of EDX:EAX by the 32-bit value stored at memory location *var*. Place the quotient in EAX and the remainder in EDX.

**and, or, xor** — Bitwise logical and, or and exclusive or

These instructions perform the specified logical operation (logical bitwise and, or, and exclusive or, respectively) on their operands, placing the result in the first operand location.

### *Syntax*

and <reg>,<reg>

and <reg>,<mem>

and <mem>,<reg>

and <reg>,<con>

and <mem>,<con>

or <reg>,<reg>

or <reg>,<mem>

or <mem>,<reg>

or <reg>,<con>

or <mem>,<con>

xor <reg>,<reg>

xor <reg>,<mem>

xor <mem>,<reg>

xor <reg>,<con>

xor <mem>,<con>

### *Examples*

and eax, 0fH — clear all but the last 4 bits of EAX.

xor edx, edx — set the contents of EDX to zero.

## **not** — Bitwise Logical Not

Logically negates the operand contents (that is, flips all bit values in the operand).

### *Syntax*

`not <reg>`

`not <mem>`

### *Example*

`not BYTE PTR [var]` — negate all bits in the byte at the memory location *var*.

## **neg** — Negate

Performs the two's complement negation of the operand contents.

### *Syntax*

neg <reg>

neg <mem>

### *Example*

neg eax — EAX  $\rightarrow$  - EAX

**shl, shr** — Shift Left, Shift Right

These instructions shift the bits in their first operand's contents left and right, padding the resulting empty bit positions with zeros. The shifted operand can be shifted up to 31 places. The number of bits to shift is specified by the second operand, which can be either an 8-bit constant or the register CL. In either case, shifts counts of greater than 31 are performed modulo 32.

### *Syntax*

shl <reg>,<con8>

shl <mem>,<con8>

shl <reg>,<cl>

shl <mem>,<cl>

shr <reg>,<con8>

shr <mem>,<con8>

shr <reg>,<cl>

shr <mem>,<cl>

### *Examples*

shl eax, 1 — Multiply the value of EAX by 2 (if the most significant bit is 0)

shr ebx, cl — Store in EBX the floor of result of dividing the value of EBX by  $2^n$  where  $n$  is the value in CL.

## **jmp** — Jump

Transfers program control flow to the instruction at the memory location indicated by the operand.

### *Syntax*

`jmp <label>`

### *Example*

`jmp begin` — Jump to the instruction labeled `begin`.



These instructions are conditional jumps that are based on the status of a set of condition codes that are stored in a special register called the *machine status word*. The contents of the machine status word include information about the last arithmetic operation performed. For example, one bit of this word indicates if the last result was zero. Another indicates if the last result was negative. Based on these condition codes, a number of conditional jumps can be performed. For example, the `jz` instruction performs a jump to the specified operand label if the result of the last arithmetic operation was zero. Otherwise, control proceeds to the next instruction in sequence.

A number of the conditional branches are given names that are intuitively based on the last operation performed being a special compare instruction, `cmp` (see below). For example, conditional branches such as `jle` and `jne` are based on first performing a `cmp` operation on the desired operands.

#### *Syntax*

`je <label>` (jump when equal)

`jne <label>` (jump when not equal)

`jz <label>` (jump when last result was zero)

`jg <label>` (jump when greater than)

`jge <label>` (jump when greater than or equal to)

`jl <label>` (jump when less than)

`jle <label>` (jump when less than or equal to)

#### *Example*

```
cmp eax, ebx
```

```
jle done
```

If the contents of EAX are less than or equal to the contents of EBX, jump to the label *done*. Otherwise, continue to the next instruction.

## **cmp** — Compare

Compare the values of the two specified operands, setting the condition codes in the machine status word appropriately. This instruction is equivalent to the sub instruction, except the result of the subtraction is discarded instead of replacing the first operand.

### *Syntax*

cmp <reg>,<reg>

cmp <reg>,<mem>

cmp <mem>,<reg>

cmp <reg>,<con>

### *Example*

```
cmp DWORD PTR [var], 10
```

```
jeq loop
```

If the 4 bytes stored at location *var* are equal to the 4-byte integer constant 10, jump to the location labeled *loop*.

## **call, ret** — Subroutine call and return

These instructions implement a subroutine call and return. The call instruction first pushes the current code location onto the hardware supported stack in memory (see the push instruction for details), and then performs an unconditional jump to the code location indicated by the label operand. Unlike the simple jump instructions, the call instruction saves the location to return to when the subroutine completes.

The ret instruction implements a subroutine return mechanism. This instruction first pops a code location off the hardware supported in-memory stack (see the pop instruction for details). It then performs an unconditional jump to the retrieved code location.

### *Syntax*

call <label>

ret

## A digression on signed vs unsigned conditional jumps

When comparing *signed* integers, you should use jl, jg, jle, and jge.

When comparing *unsigned* integers, you should use jb, ja, jbe, and jae.

Why?

And why don't we have separate conditionals for (in)equality?

Consider these two 16-bit numbers:

ax= 1000000110010001

bx= 00000000000110000

Which is bigger?

## A digression on signed vs unsigned conditional jumps

When comparing *signed* integers, you should use jl, jg, jle, and jge.

When comparing *unsigned* integers, you should use jb, ja, jbe, and jae.

Why?

And why don't we have separate conditionals for (in)equality?

Consider these two 16-bit numbers:

	Unsigned	2's complement
ax= 1000000110010001	33169	-32367
bx= 00000000000110000	48	48

Which is bigger?

## A digression on signed vs unsigned conditional jumps

When comparing *signed* integers, you should use `jl`, `jg`, `jle`, and `jge`.

When comparing *unsigned* integers, you should use `jb`, `ja`, `jbe`, and `jae`.

Why?

And why don't we have separate conditionals for (in)equality?

Consider these two 16-bit numbers:

	Unsigned	2's complement
<code>ax= 1000000110010001</code>	<b>33169</b>	<b>-32367</b>
<code>bx= 00000000000110000</code>	<b>48</b>	<b>48</b>

Which is bigger?

```
cmp ax, bx
jg foobar      ; -32367 > 48 -- WILL NOT branch
```

```
cmp ax, bx
ja foobar      ; 33169 > 48 -- WILL branch
```

## A digression on signed vs unsigned conditional jumps

When comparing *signed* integers, you should use `jl`, `jg`, `jle`, and `jge`.

When comparing *unsigned* integers, you should use `jb`, `ja`, `jbe`, and `jae`.

Why?

And why don't we have separate conditionals for (in)equality?

Comparison	Signed	Unsigned
<	<b>jl</b> jump if less	<b>jb</b> jump if below
>	<b>jg</b> jump if greater	<b>ja</b> jump if above
<=	<b>jle</b> jump if less or equal	<b>jbe</b> jump if below or equal
>=	<b>jge</b> jump if greater or equal	<b>jae</b> jump if above or equal
==	<b>je/jz</b> jump if equal/jump if zero	
!=	<b>jne/jnz</b> jump if not equal/jump if not zero	

## Unsigned comparison implementation

Unsigned comparison	jb label ja label jbe label jae label	jump if CF jump if !CF && !ZF jump if CF    ZF jump if !CF
---------------------	--	---

As on E15, cmp is implemented in terms of subtraction.

Zero flag (ZF) is set when the difference between two operands is zero.

Carry flag (CF) is set when an addition results in a carry, or a subtraction results in a borrow

```
  1011
- 0111
-----
  0100
```

```
  0011
- 1011
-----
  1100
```



# Unsigned comparison implementation

Unsigned comparison	jb label ja label jbe label jae label	jump if CF jump if !CF && !ZF jump if CF    ZF jump if !CF
---------------------	--	---

As on E15, cmp is implemented in terms of subtraction.

Zero flag (ZF) is set when the difference between two operands is zero.

Carry flag (CF) is set when an addition results in a carry, or a subtraction results in a borrow

1011  
- 0111  
-----  
0100


11 < 7?

no borrow  
required,  
CF not set

0011  
- 1011  
-----  
1100

7 < 11?

borrow  
required,  
CF set



We will focus on application software. Such programs are typically invoked by the operating system (for example, when the user clicks on the appropriate file). Application software communicates with the operating system by way of *system calls*, by which means it can request services, such as:

- Accessing a file on disk

- Accessing the network

- Interfacing with the user (e.g. screen, keyboard, mouse, etc)

- Interfacing with other applications (shared memory, pipes, etc)

**Intel x86 FLAGS register<sup>[1]</sup>**

Bit #	Abbreviation	Description	Category
<b>FLAGS</b>			
0	CF	Carry flag	Status
1		Reserved, always 1 in EFLAGS <sup>[2]</sup>	
2	PF	Parity flag	Status
3		Reserved	
4	AF	Adjust flag	Status
5		Reserved	
6	ZF	Zero flag	Status
7	SF	Sign flag	Status
8	TF	Trap flag (single step)	Control
9	IF	Interrupt enable flag	Control
10	DF	Direction flag	Control
11	OF	Overflow flag	Status
12-13	IOPL	I/O privilege level (286+ only), always 1 on 8086 and 186	System
14	NT	Nested task flag (286+ only), always 1 on 8086 and 186	System
15		Reserved, always 1 on 8086 and 186, always 0 on later models	
<b>EFLAGS</b>			
16	RF	Resume flag (386+ only)	System
17	VM	Virtual 8086 mode flag (386+ only)	System
18	AC	Alignment check (486SX+ only)	System
19	VIF	Virtual interrupt flag (Pentium+)	System
20	VIP	Virtual interrupt pending (Pentium+)	System
21	ID	Able to use CPUID instruction (Pentium+)	System
22		Reserved	
23		Reserved	

etc, etc... the rest are reserved.

You can't access EFLAGS directly, but you can put it on the stack and pop it from the stack with the PUSHF/PUSHFD and POPF/POPFd instructions.

In addition, some flag bits can be manipulated individually, e.g. CLD/STD will clear or set the direction flag; CLI/STI will clear or set the interrupt flag.

The conditional jump (Jcc) instructions use several flag bits as input, e.g. JC will jump when the carry flag is set.

Most arithmetic and compare instructions will set some flags.

# Euclid's Algorithm on the SPARC

```
mov    %i0, %o1
b      .LL3
mov    %i1, %i0
.LL5:
mov    %o0, %i0
.LL3:
mov    %o1, %o0
call   .rem, 0
mov    %i0, %o1
cmp    %o0, 0
bne    .LL5
mov    %i0, %o1
ret
restore
```

```
while ((r = m % n) != 0) {
    m = n;
    n = r;
}
return n;
```

Register assignments:

m	%o1
r	%o0
n	%i0

“n = r”

“m = n” (executed even if  
loop terminates)

“Branch back to caller”  
SPARC has no ret: this is  
jmp %i7 + 8

Inverse of save: return  
to previous register  
window