


Intel stack

Stack instructions

Instruction	What it does (conceptually)
<code>JMP addr</code>	<code>MOV EIP, addr</code>
<code>PUSH reg</code>	<code>SUB ESP, 4</code> <code>MOV [ESP], reg</code>
<code>POP reg</code>	<code>MOV reg, [ESP]</code> <code>ADD ESP, 4</code>
<code>CALL addr</code>	<code>PUSH EIP</code> <code>JMP addr</code>
<code>RET</code>	<code>POP EIP</code>
<code>RET n</code>	<code>POP EIP</code> <code>ADD ESP, n</code>

addr	val
10a4	...
10a0	...
109c	...
1098	...
1094	...
1090	...



addr	val
10a4	...
10a0	...
109c	...
1098	5
1094	...
1090	...

esp

esp = 109c

esp = 1098

addr	val
10a4	...
10a0	...
109c	...
1098	5
1094	...
1090	...

← esp

esp = 1098

pop eax
→

addr	val
10a4	...
10a0	...
109c	...
1098	5
1094	...
1090	...

← esp

esp = 109c
eax = 5

addr	val
10a4	...
10a0	...
109c	...
1098	...
1094	...
1090	...

esp
←

```

2030  foo:
...
2050  ret
...
2090  call foo
2095  ....

```



addr	val
10a4	...
10a0	2095
109c	...
1098	...
1094	...
1090	...

esp
←

esp = 10a4
eip = 2090

esp = 10a0
eip = 2030

addr	val
10a4	...
10a0	2095
109c	...
1098	...
1094	...
1090	...

esp
←

```

2030  foo:
...
2050  ret
...
2090  call foo
2095  ....

```



addr	val
10a4	...
10a0	2095
109c	...
1098	...
1094	...
1090	...

esp
←

esp = 10a0
eip = 2050

esp = 10a4
eip = 2095

A function

```
section .data
message: db "Hello",10
```

```
section .text
global _start
```

```
print_message:
    mov eax, 4
    mov ebx, 1
    mov ecx, message
    mov edx, 6
    int 80h
    ret
```

```
_start:
    call print_message
    call print_message

    mov eax, 1
    mov ebx, 0
    int 80h
```

A function

```
section .text  
global _start
```

```
times_two:  
    add eax, eax  
    ret
```

```
_start:  
    mov eax, 4  
    call times_two  
    call times_two  
  
    mov eax, 1  
    mov ebx, 0  
    int 80h
```


A broken Fibonacci

```
section .text
global _start
```

```
fib:
    cmp eax, 2
    jae recursive_case
    mov eax, 1
    jmp done

recursive_case:
    mov ebx, eax

    dec eax          ; call fib(n-1)
    call fib         ; result in eax
    mov ecx, eax

    mov eax, ebx
    sub eax, 2
    call fib         ; call fib(n-2)

    add eax, ecx     ; fib(n-1) + fib(n-2)
done:
    ret
```

```
_start:
    mov eax, 5
    call fib

    mov ebx, eax
    mov eax, 1
    int 80h
```

Typical prologue/epilogue

Typical function prologue:

```
push ebp      ; save the caller's base pointer
mov ebp, esp  ; set our base pointer
sub esp, X    ; allocate X bytes for local vars
```

Typical function epilogue:

```
mov esp, ebp  ; de-allocate local vars
pop ebp       ; restore caller's base pointer
ret N         ; return, and pop N bytes of params
```

Function call with two parameters

addr	val
10a4	0
10a0	99
109c	42
1098	...
1094	...
1090	...

ebp
←

esp
←

```
2030  foo:
      push ebp
      mov ebp, esp
2038  ...
      mov esp, ebp
      pop ebp
2050  ret 8
...
208a  push 99h
208c  push 42h
2090  call foo
2095  ....
```



addr	val
10a4	0
10a0	99
109c	42
1098	2095
1094	10a4
1090	...

esp,ebp
←

esp = 109c
ebp = 10a4
eip = 2090

Within `foo` (around address 2038),
the first parameter is accessible as
`[ebp+8]` and the second is
`[ebp+12]`

The `ret 8` instruction will clear both
parameters off the stack before
returning

esp = 1094
ebp = 1094
eip = 2038

Function call with two parameters and a local variable

addr	val
10a4	0
10a0	99
109c	42
1098	...
1094	...
1090	...

ebp
←

esp
←

```

2030  foo:
      push ebp
      mov ebp, esp
      sub esp, 4
2040  ...
      mov esp, ebp
      pop ebp
2050  ret 8
...
208a  push 99h
208c  push 42h
2090  call foo
2095  ....

```



addr	val
10a4	0
10a0	99
109c	42
1098	2095
1094	10a4
1090	...

ebp
←

esp
←

Parameters available as before, and now the local variable is accessible as `[ebp-4]`. The `sub` instruction allocates space on the stack for local variables; allocate as much space as you need. Local variables are removed from stack by the `mov esp, ebp` instruction.

esp = 109c
ebp = 10a4
eip = 2090

esp = 1090
ebp = 1094
eip = 2040

A recursive call		addr	val	
8060	foof:	10a4	...	
	push ebp	10a0	9	
8061	mov ebp, esp			
8063	mov eax, [ebp+8]	109c	808d	
8066	cmp eax, 0	1098	0	original EBP
8069	je .done			
806b	dec eax	1094	8	
806c	push eax			
806d	call foof	1090	8072	
8072	.done:	108c	1098	
	pop ebp			
8073	ret 4	1088	7	
8076	_start:	1084	8072	
	push 9			
8078	call foof	1080	108c	
808d	...	107c	6	
		1078	8072	
		1074	1080	
		1070	5	
		106c	8072	
		1068	...	

A better Fibonacci

```
section .text
global _start
```

```
_start:
    push 10
    call fib
```

```
fib:
    push ebp        ; prologue
    mov ebp, esp

    mov eax, [ebp+8] ; parameter
    cmp eax, 2
    jae recursive_case
    mov eax, 1
    jmp done
```

```
    mov ebx, eax
    mov eax, 1
    int 80h
```

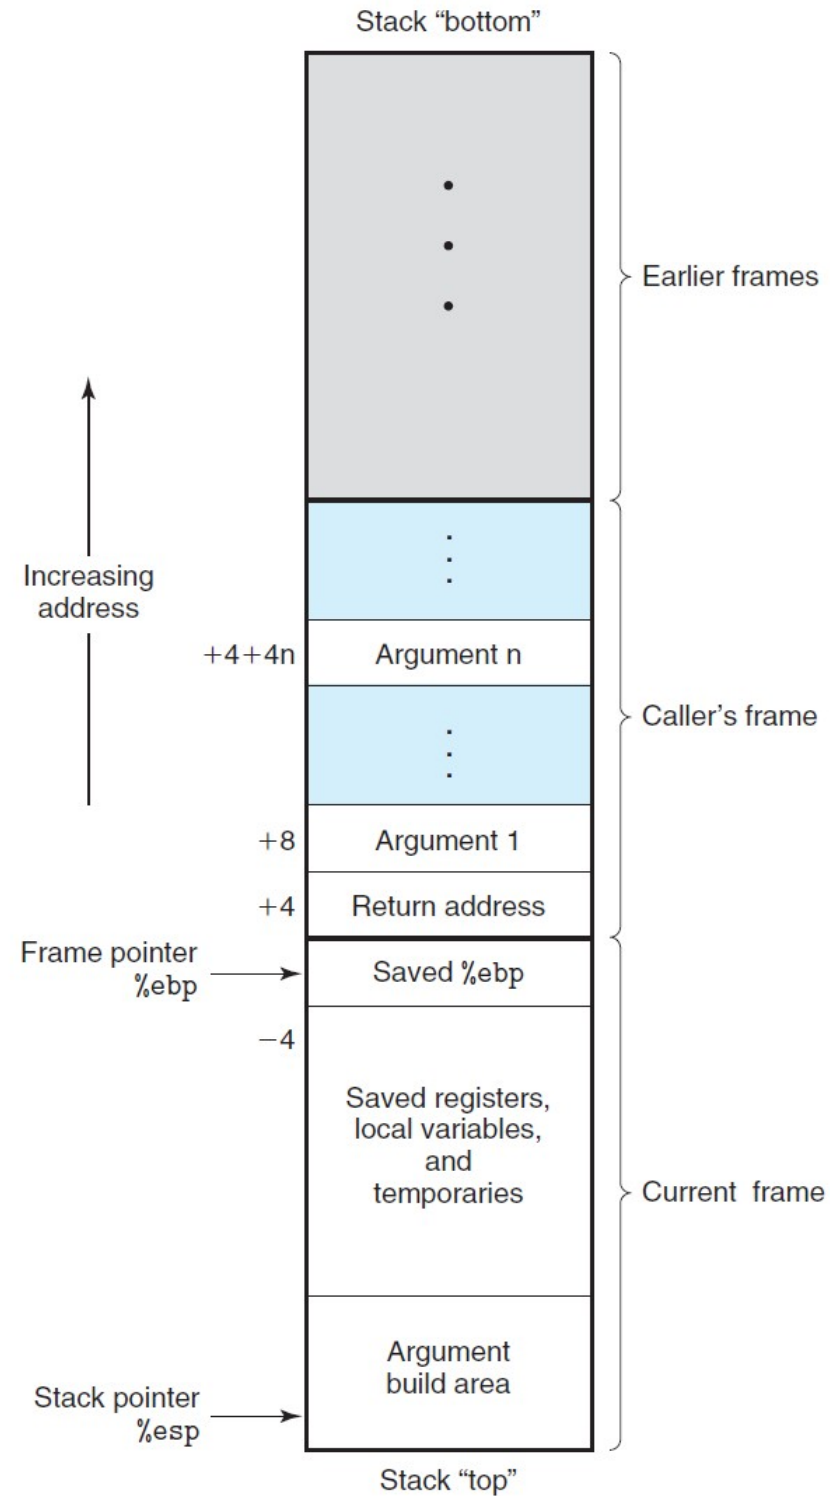
```
recursive_case:
    dec eax
    push eax        ; call fib(n-1)
    call fib        ; result in eax
    push eax

    mov eax, [ebp+8]
    sub eax, 2
    push eax        ; call fib(n-2)
    call fib        ; result in eax

    pop ebx
    add eax, ebx    ; fib(n-1)+fib(n-2)
done:
    pop ebp        ; epilogue
    ret 4
```

Figure 3.21

Stack frame structure. The stack is used for passing arguments, for storing return information, for saving registers, and for local storage.



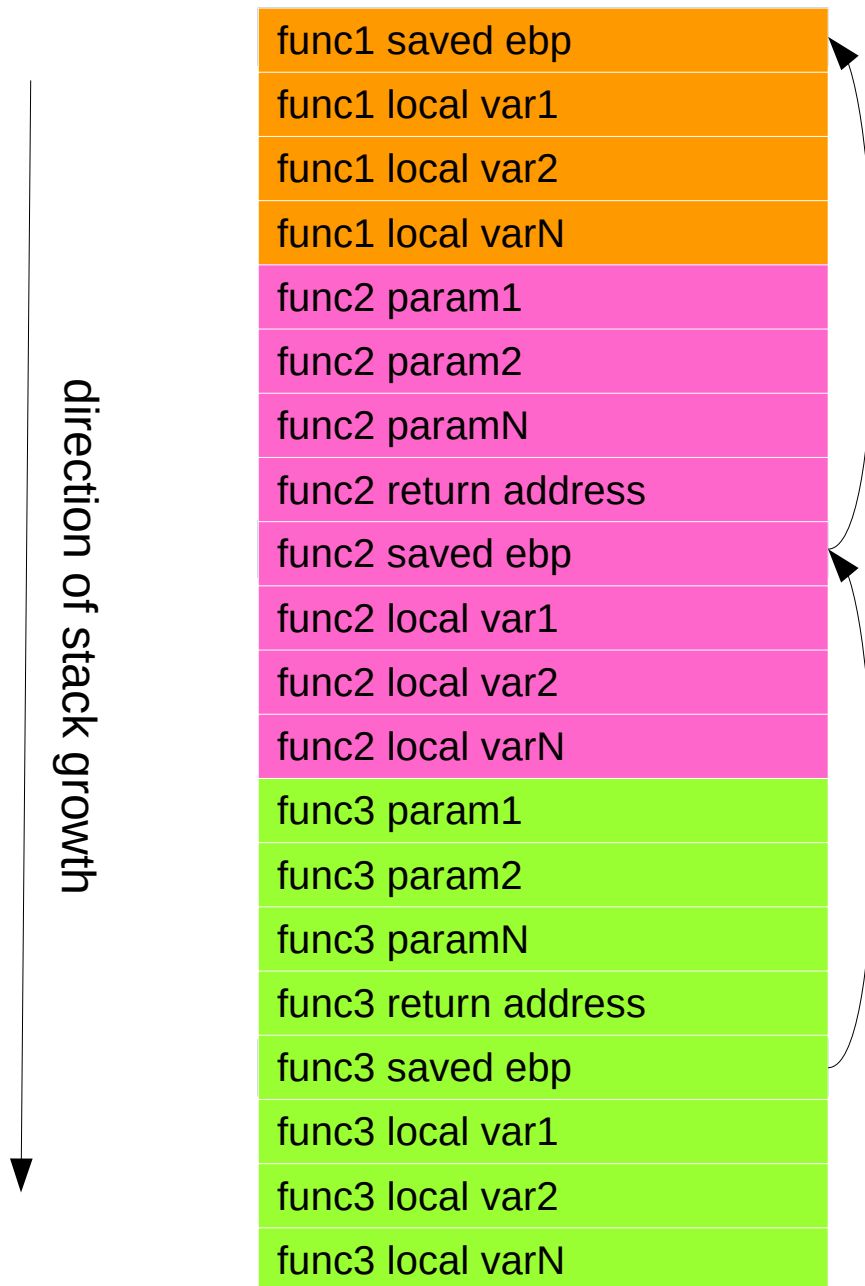
How to access params and local vars

- Given the stack layout on the previous slide, we know this:
 - EBP is the address of the caller's base pointer
 - EBP+4 is the address of this function's return address
 - EBP+8, EBP+12, EBP+18, etc... are the addresses of this function's parameters, in the opposite order that they were pushed (assuming all parameters are dwords)
 - EBP-4, EBP-8, EBP-12, etc.... are the addresses of this function's local variables (assuming that all variables are dwords)

How to access params and local vars

- In practice:
 - `mov eax, [ebp+8]` ; load first parameter
 - `mov ebx, [ebp-4]` ; load first local variable
- Parameters larger than a dword may be passed:
 - by *reference*: the (dword-sized) pointer to the actual value is pushed on the stack; the full value is not placed on the stack
 - by *value*: the full value is placed on the stack, in a sequence of pushes

Chain of stack frames

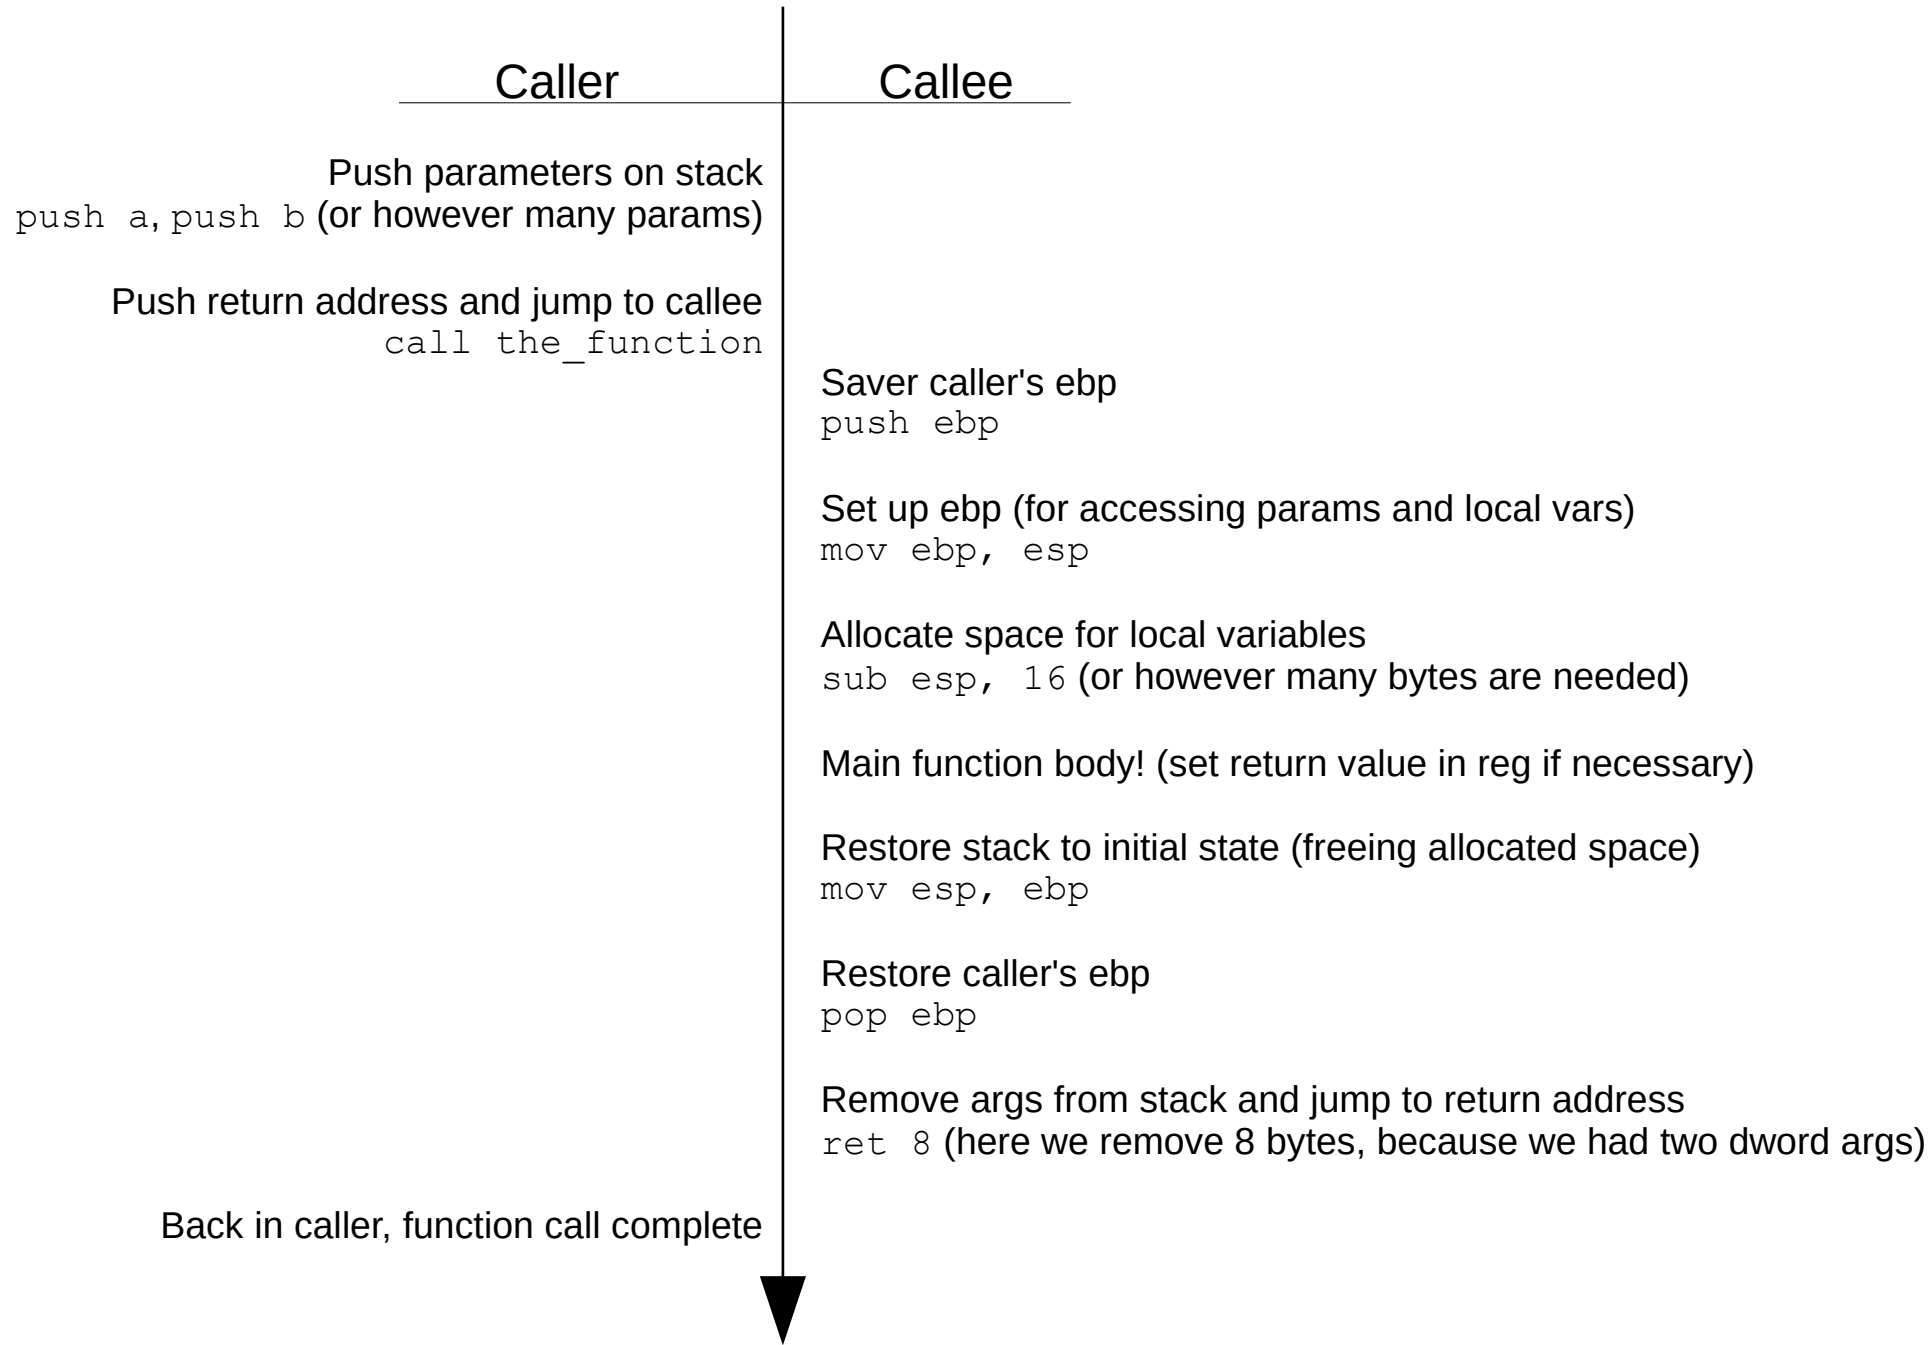


```
def func1():  
    var1=1  
    var2=2  
    varN=3  
    func2(4, 5, 6)
```

```
def func2(param1, param2, paramN):  
    var1=7  
    var2=8  
    var3=9  
    func3(10, 11, 12)
```

```
def func3(param1, param2, paramN):  
    var1=13  
    var2=14  
    var3=15  
    print_stack_trace()
```

Timeline of a function call



A note on a common point of confusion

`ret`

This instruction pops the return address and jumps to it.

`ret n`

This instruction pops the return address, jumps to it, *and removes an additional *n* bytes from the stack.*

`ret n`, despite its appearance, does *not* specify the return value of a function! That is typically given in the `eax` register.

Removing the additional *n* bytes from the stack is called *stack clean-up*: this removes the parameters that the caller pushed before calling the function. We need to do this, or else the height of the stack would be different after calling.

We've addressed a style called "callee clean-up," wherein the called function (through the `ret n` instruction) removes the parameters before returning. An alternative style, "caller clean-up" is similar, but the called function does a simple `ret` and then the caller must adjust the stack pointer manually (typically with `add esp, n`).

It doesn't matter which style is used, as long as caller and callee agree. For this class, we are only concerned with callee clean-up.

Uppercase a single letter

```
toupper:
push ebp
mov ebp, esp
mov eax, [ebp+8]    ; param is a single ASCII char
cmp al, 97          ; 'a'
jb done
cmp al, 122         ; 'z'
ja done
sub al, 32

done:
pop ebp
ret 4
```

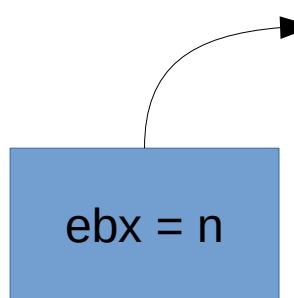
Given a single character, return its uppercase form in eax. If the parameter is not a lowercase letter, return the character unchanged.

Strings

- What do strings look like in memory?
- We can't store them in a register: strings can be of arbitrary size, but registers are small
- Instead, strings are arrays of bytes, where each byte represents a character
- Then, we store the address of the first byte of the string in a register: this is a `char*`
- How do we know how long the string is?
 - Option 1: The byte after the last character in the string is 0. When we reach that byte, we know we've found the end of the string
 - Option 2: The length of the string is stored separately as a dword

Strings

Option 1: register stores address of first byte. String is terminated by NUL.

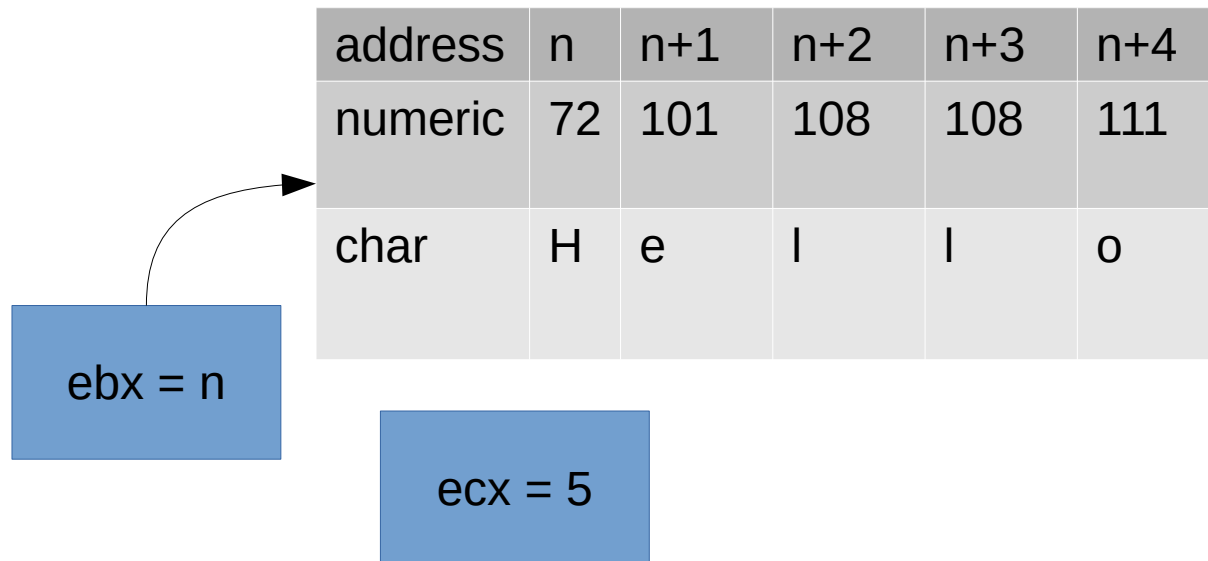


A blue box labeled 'ebx = n' has a curved arrow pointing to the first column of the table below.

address	n	n+1	n+2	n+3	n+4	n+5
numeric	72	101	108	108	111	0
char	H	e	l	l	o	NUL

Strings

Option 2: no terminated NULL byte. Instead, length is carried separately.



Uppercase a whole string

```
toupper_string:
push ebp
mov ebp, esp
mov ebx, [ebp+8]    ; param is address of a null-terminated string

loop:
xor ecx,ecx
mov cl, [ebx]      ; get one char from string
cmp cl, 0          ; if it's null, we're done
je done
push ecx
call toupper       ; returns output char in al
mov [ebx], al      ; store it back into string
inc ebx            ; go to next char
jmp loop

done:
pop ebp
ret 4
```

Given a pointer to a null-terminated string, adjust it to be uppercase. We call the toupper function defined earlier. We do not return a useful value: instead the parameter is mutated.

This is okay
because we know
that toupper doesn't
modify ebx

$$\text{gcd}(x, y) = \begin{cases} x & \text{if } y = 0 \\ \text{gcd}(y, \text{remainder}(x, y)) & \text{if } y > 0 \end{cases}$$

```
int gcd(int x, int y)
{
    if (y == 0)
        return x;
    else
        return gcd(y, x % y);
}
```

Call it like this:

```
push 15
push 25
call gcd
```

gcd:

```
    push ebp
    mov ebp, esp
    ; assume x is at ebp+8, y is at ebp+12

    ;; YOUR CODE HERE
    ;; put return value in eax

    mov esp, ebp
    pop ebp
    ret 8
```

```

gcd:
    push ebp
    mov ebp, esp
    ; assume x is at ebp+8, y is at ebp+12

    cmp [ebp+12], 0
    je returnx

    xor edx, edx
    mov eax, [ebp+8]
    div [ebp+12]          ; divides edx:eax by [ebp+12]
                          ; leaves remainder in edx

    push edx              ; push args right-to-left
    push [ebp+12]
    call gcd
    jmp justreturn

returnx:
    mov eax, [ebp+8]

justreturn:
    mov esp, ebp
    pop ebp
    ret 8

```

Debugging with gdb: an example

- We want to sum an array of dwords

```
section .data
array_len:
    dd 10
array:
    dd 1, 2, 3, 4, 5, 6, 7, 8, 9, 10
```

Debugging with gdb: an example

- Almost correct solution:

```
section .data
array_len:
    dd 10
array:
    dd 1, 2, 3, 4, 5, 6, 7, 8, 9, 10
sum:
    dd 0
section .text
_start:
    mov ecx, array_len
    mov eax, array
loop:
    mov ebx, [eax]
    add [sum], ebx
    add eax, 4
    dec ecx      ; reduce length by 1 to keep track of # of iterations
    cmp ecx, 1   ; stop if we're at the end
    jne loop
```

Debugging with gdb: an example

- Problem:

```
eppie@kmotr:~/Downloads$ ./sum-array  
Segmentation fault
```

Debugging with gdb: an example

- Let's run the program in gdb

Output/messages

```
Program received signal SIGSEGV, Segmentation fault.  
0x08049010 in loop ()
```

Assembly

```
0x08049010 loop+0  mov     ebx,DWORD PTR [eax]  
0x08049012 loop+2  add     DWORD PTR ds:0x804a024,ebx  
0x08049018 loop+8  add     eax,0x4  
0x0804901b loop+11 dec     ecx
```

Expressions

History

Memory

Registers

eax	0x0804b000	ecx	0x08049c01	edx	0x00000000	ebx	0x00000000
esp	0xffffd290	ebp	0x00000000	esi	0x00000000	edi	0x00000000
eip	0x08049010	eflags	[PF IF RF]	cs	0x00000023	ss	0x0000002b
ds	0x0000002b	es	0x0000002b	fs	0x00000000	gs	0x00000000

Source

Stack

Debugging with gdb: an example

- Why is this instruction problematic?

Output/messages

```
Program received signal SIGSEGV, Segmentation fault.  
0x08049010 in loop ()
```

Assembly

```
0x08049010 loop+0 mov     ebx,DWORD PTR [eax]  
0x08049012 loop+2 add     DWORD PTR ds:0x804a024,ebx  
0x08049018 loop+8 add     eax,0x4  
0x0804901b loop+11 dec     ecx
```

Expressions

History

Memory

Registers

eax	0x0804b000	ecx	0x08049010	edx	0x00000000	ebx	0x00000000
esp	0xffffd290	edx	0x00000000	esi	0x00000000	edi	0x00000000
eip	0x08049010	ecx	0x00000023	cs	0x00000023	ss	0x0000002b
ds	0x0000002b	fs	0x00000000	gs	0x00000000		

eax should store the address of an array element. Evidently, this address is invalid. But why?

Source

Debugging with gdb: an example

- Why is this instruction problematic?

Output/messages

```
Program received signal SIGSEGV, Segmentation fault.  
0x08049010 in loop ()
```

Assembly

```
0x08049010 loop+0 mov     ebx,DWORD PTR [eax]  
0x08049012 loop+2 add     DWORD PTR ds:0x804a024,ebx  
0x08049018 loop+8 add     eax,0x4  
0x0804901b loop+11 dec     ecx
```

Expressions

History

Memory

Registers

eax	0x0804b000	ecx	0x08049c01	edx	0x00000000	ebx	0x00000000
esp	0xffffd290	ebp	0x00000000			edi	0x00000000
eip	0x08049010	eflags	[PF IF RF			ss	0x0000002b
ds	0x0000002b	es	0x0000002b			gs	0x00000000

Source

ecx should store the length of the array. This value is clearly wrong. What happened?

Debugging with gdb: an example

- Almost correct solution:

```
section .data
array_len:
    dd 10
array:
    dd 1, 2, 3, 4, 5, 6, 7, 8, 9, 10
sum:
    dd 0
section .text
_start:
    mov ecx, array_len
    mov eax, array
loop:
    mov ebx, [eax]
    add [sum], ebx
    add eax, 4
    dec ecx      ; reduce length by 1 to keep track of # of iterations
    cmp ecx, 1   ; stop if we're at the end
    jne loop
```

Debugging with gdb: an example

- Almost correct solution:

```
section .data
array_len:
    dd 10
array:
    dd 1, 2, 3, 4, 5, 6, 7, 8, 9, 10
sum:
    dd 0
section .text
_start:
    mov ecx, [array_len]
    mov eax, array
loop:
    mov ebx, [eax]
    add [sum], ebx
    add eax, 4
    dec ecx          ; reduce length by 1 to keep track of # of iterations
    cmp ecx, 1       ; stop if we're at the end
    jne loop
```

We're storing the *address* of the array length in ecx, not the array length itself. Therefore ecx's value is too high, and the loop executes too many iterations, causing eax to be incremented past the range of valid memory.