

# Virtual memory

## Questions

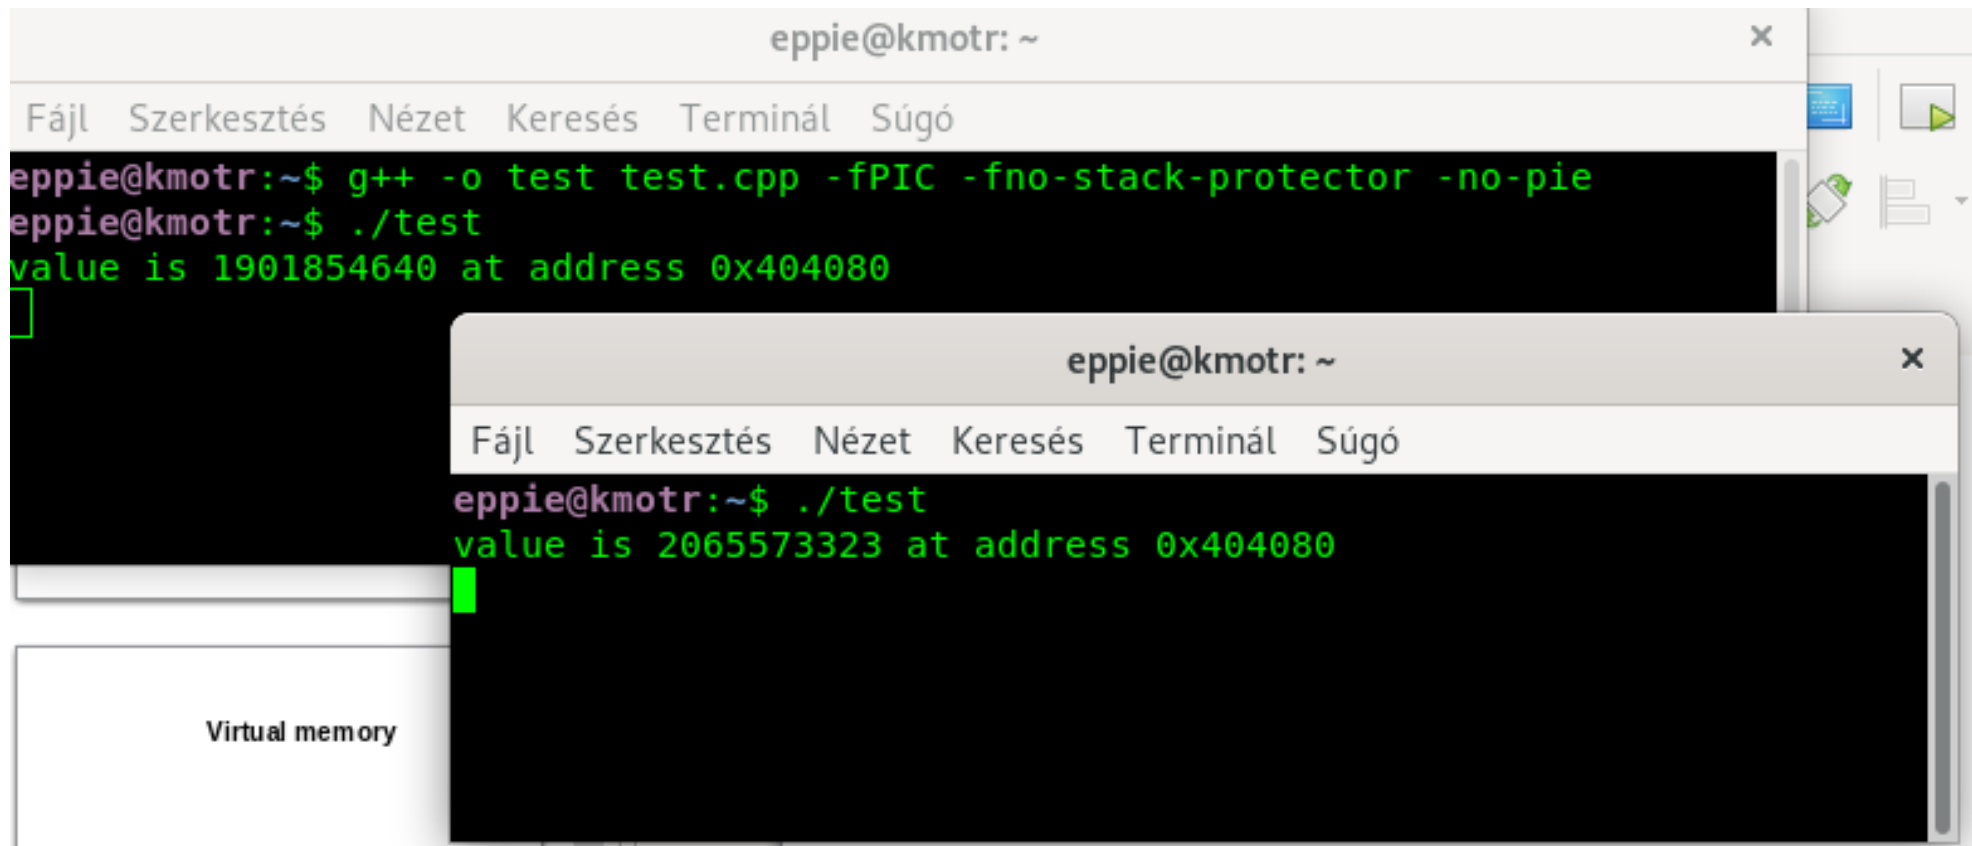
- What if you wanted to run a program that needs more memory than you have?
- What if you want to run 10 programs, which collectively need more memory than you have?
- What if multiple different programs want to store something at the same address?
- How do we protect one program's data from being read or written by another program, either intentionally or accidentally?
- How do we protect the operating system's data from being read or written by another program, either intentionally or accidentally?

```
#include <iostream>
#include <cstdlib>
#include <ctime>
#include <unistd.h>
```

```
// compile with: g++ -o test test.cpp -fPIC -fno-stack-protector -no-pie
```

```
using namespace std;
int main() {
    static int value;
    srand(time(NULL));
    value = rand();
    while(1) {
        cout << "value is " << value << " at address "
              << &value << endl;

        sleep(2);
    }
    return 0;
}
```



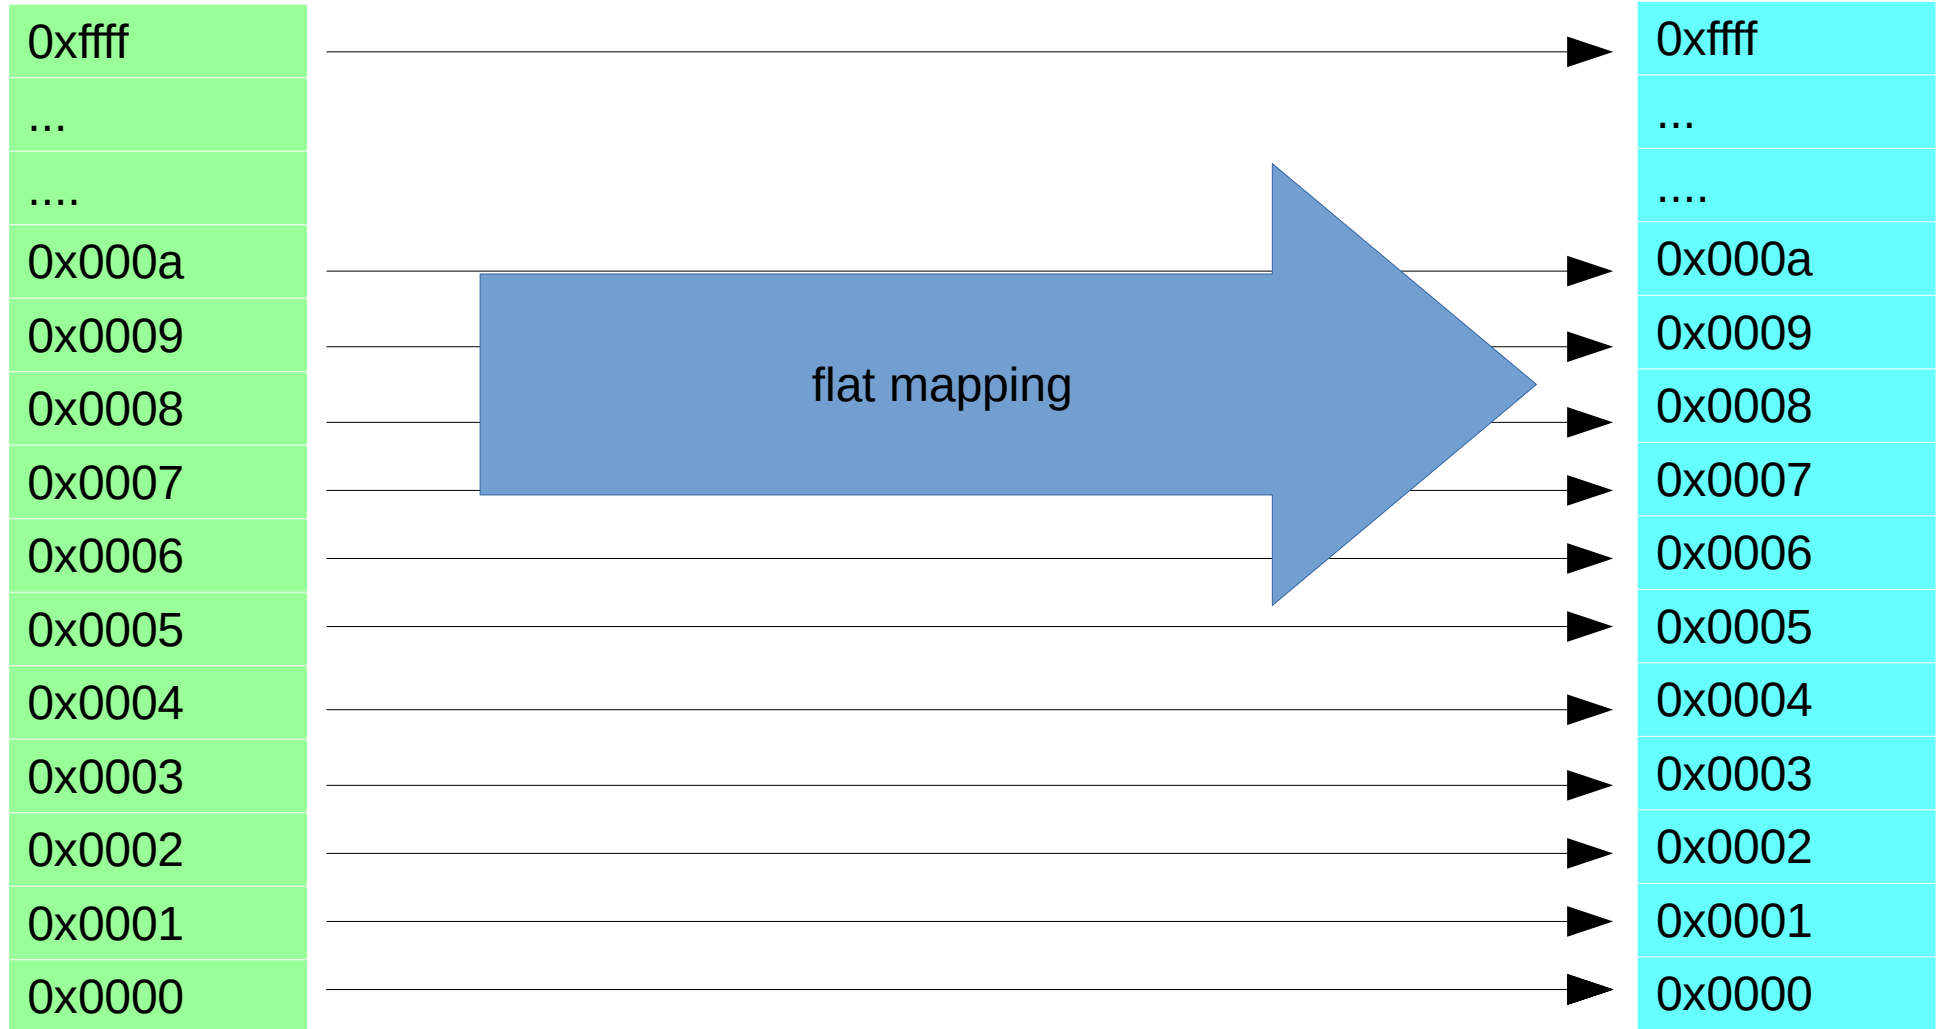
# Address space

*Address space* is the set of memory addresses, and their corresponding memory cells, that are available to a program.

On E20, the address space is expressed as a 16-bit pointer value, so it has  $2^{16}$  elements, starting from zero.

# Flat address space mapping

lw \$1, foo(\$0)

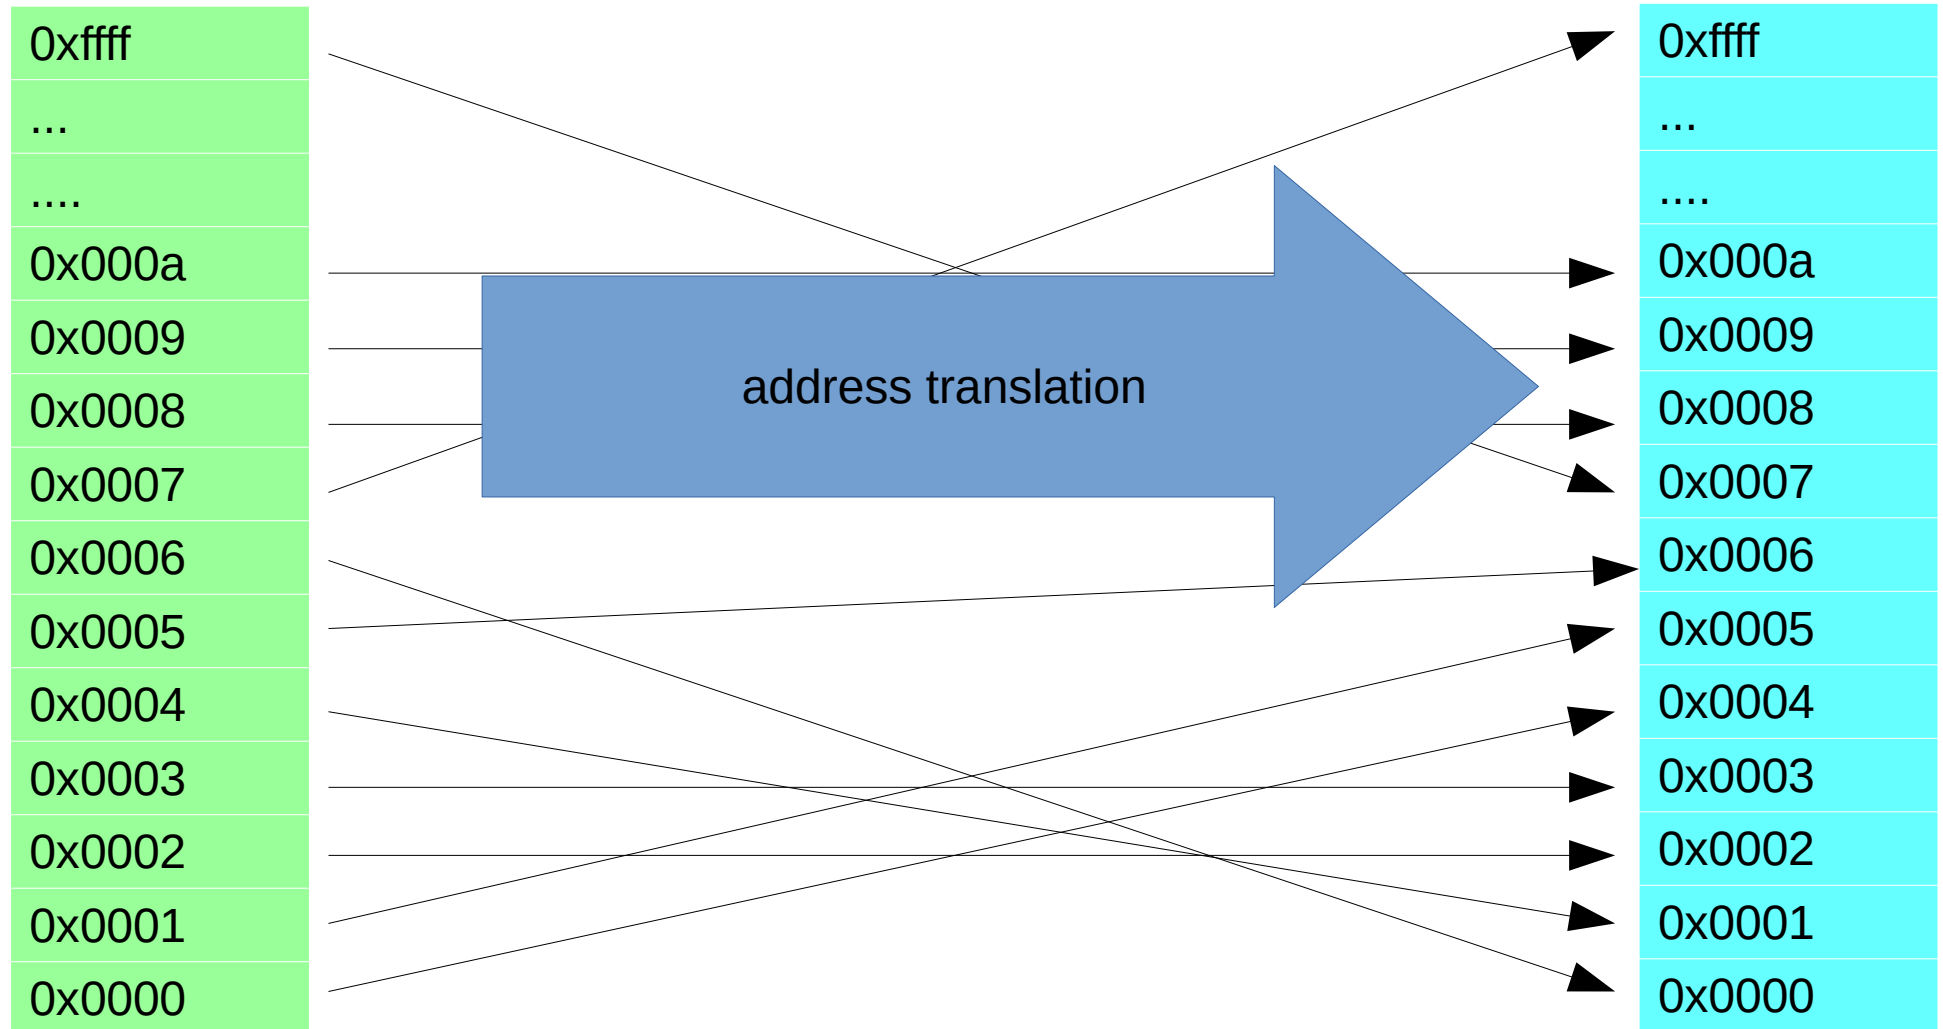


application address space

physical memory

# Virtual address space mapping


lw \$1, foo(\$0)



application address space


physical memory

# Virtual address space mapping



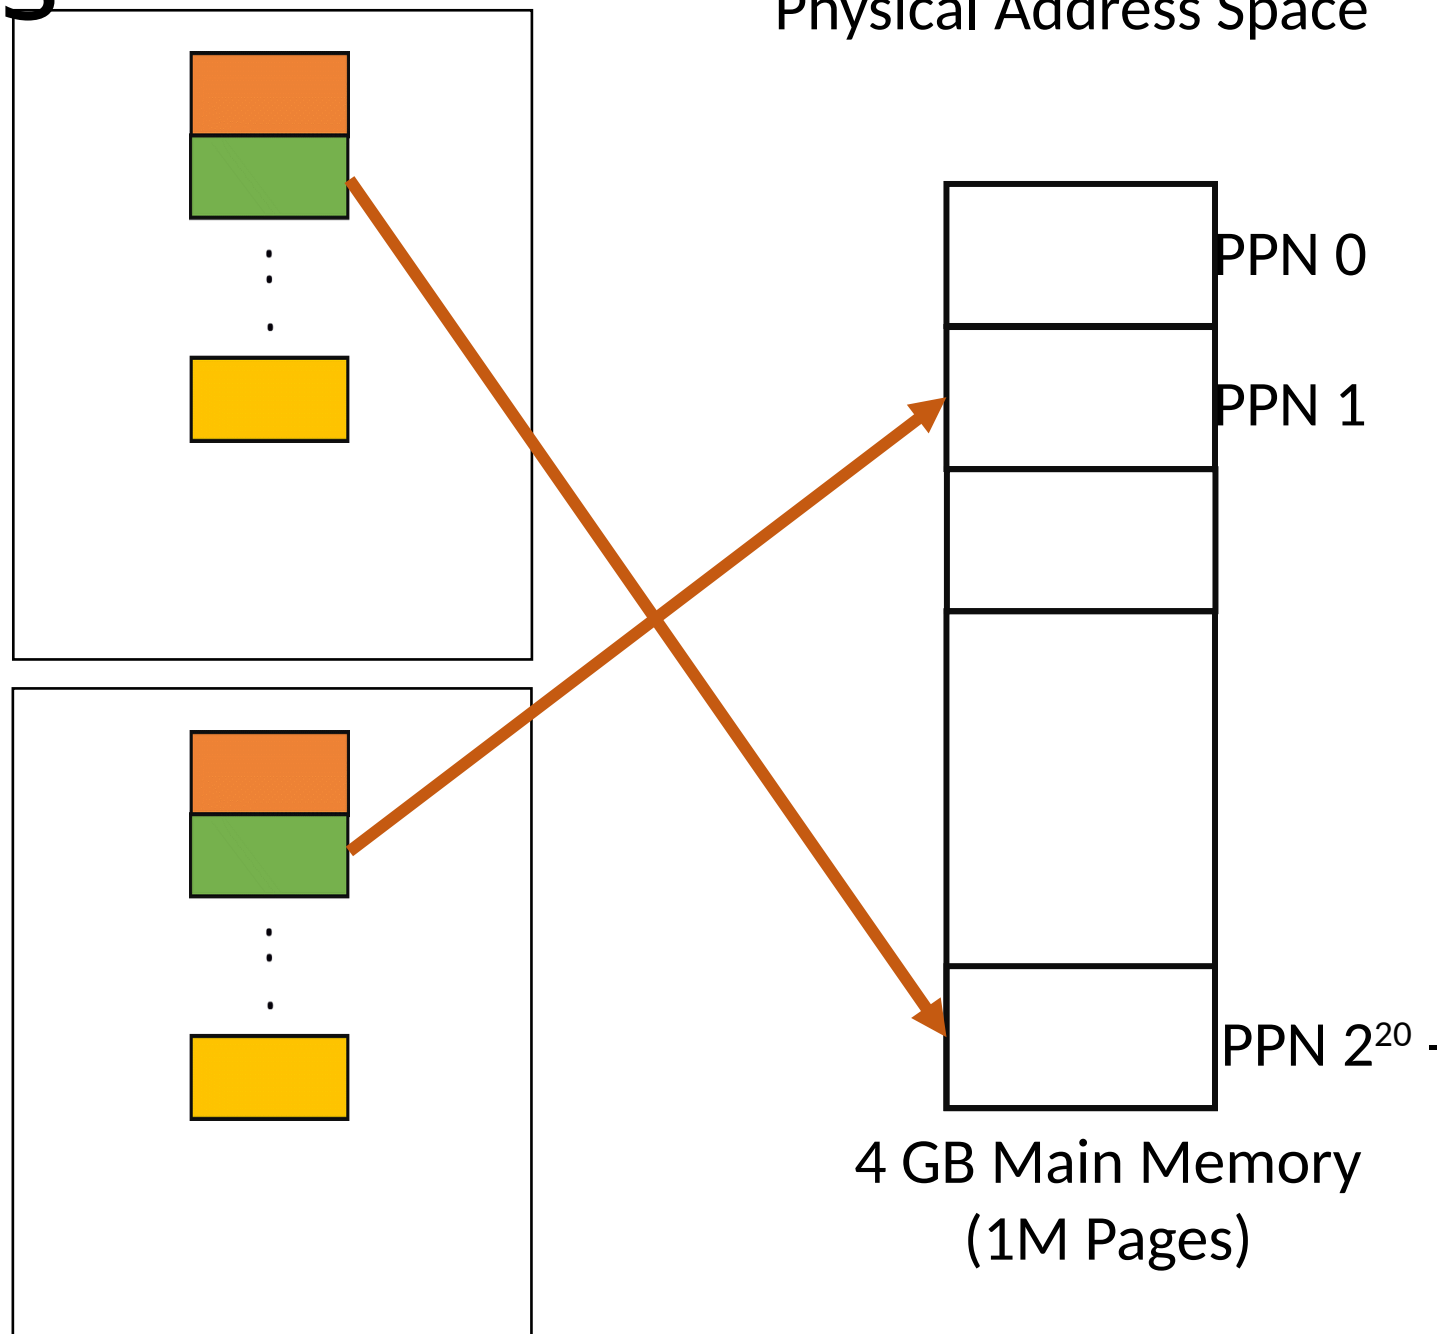


- Because different processes will have different mappings from virtual to physical addresses, two programs can freely use the same virtual address.
- By allocating distinct regions of physical memory to A and B, they are prevented from reading/writing each others data.




# Homonyms

- The same VA (in different processes) can map to two different physical addresses

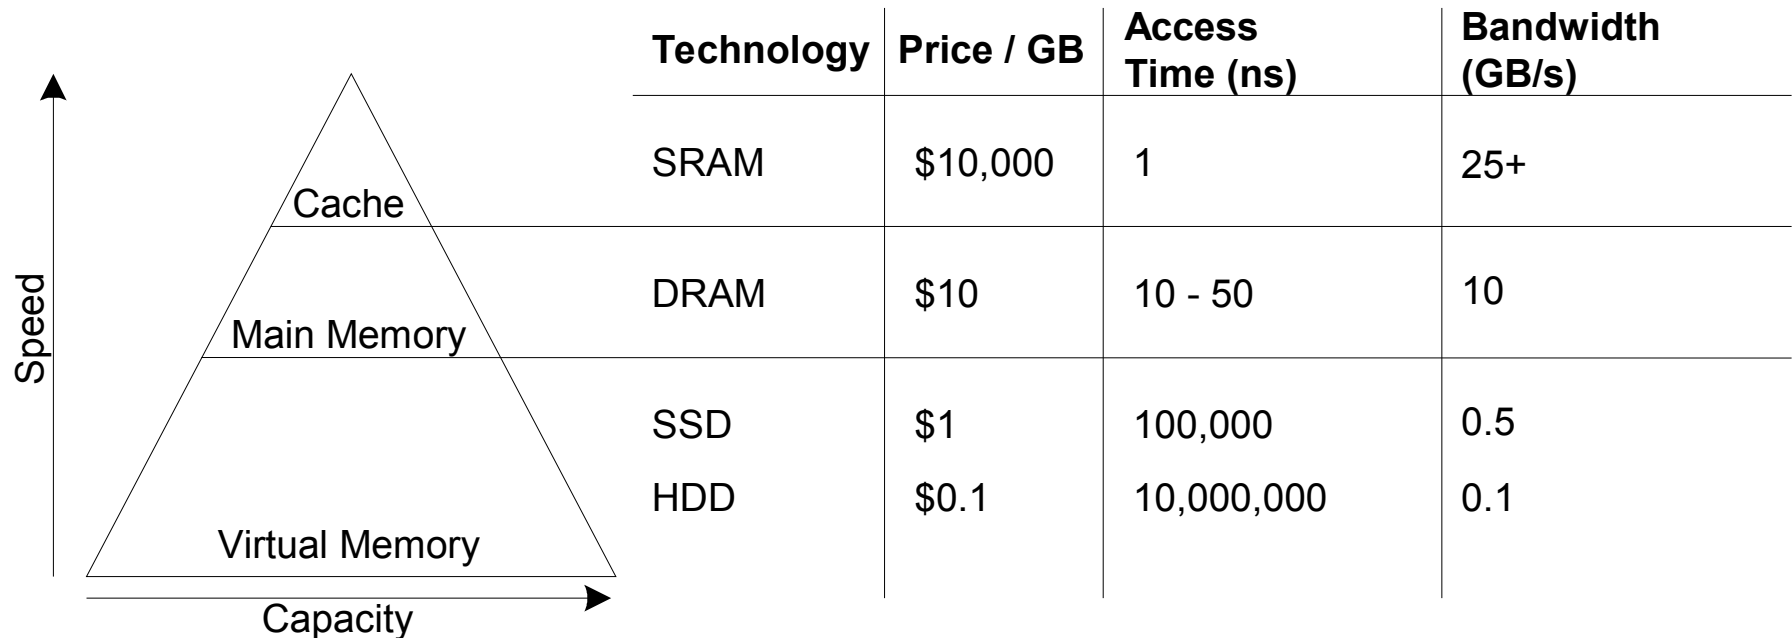


# Synonyms

- Different VAs can map to the same physical address
- Shared code and data



# Memory Hierarchy



SRAM, DRAM – volatile  
SSD, HDD – non-volatile

# Virtual address space mapping

lw \$1, foo(\$0)

Virtual address space  
can be larger than physical

address translation

0xffffffff

...

....

0x0000000a

0x00000009

0x00000008

0x00000007

0x00000006

0x00000005

0x00000004

0x00000003

0x00000002

0x00000001

0x00000000

0xffff

...

....

0x000a

0x0009

0x0008

0x0007

0x0006

0x0005

0x0004

0x0003


0x0002

0x0001

0x0000

application address space

physical memory




Virtual address	Physical address
Visible to applications	Not visible to applications; may be accessible to operating system
Used by <code>lw</code> and <code>sw</code> (or their equivalents)	
Used for fetching instructions from memory	

For every memory access, the virtual address must be translated to a physical address. This translation is done by special address translation hardware.

The translation happens automatically: the application program has no idea.

We divide address space (both virtual and physical) into *pages*. A typical size for pages is 4KB (4096, or 0x1000 bytes).

Address range	Page
0x0b000..0x0bfff	11
0x0a000..0x0afff	10
0x09000..0x09fff	9
0x08000..0x08fff	8
0x07000..0x07fff	7
0x06000..0x06fff	6
0x05000..0x05fff	5
0x04000..0x04fff	4
0x03000..0x03fff	3
0x02000..0x02fff	2
0x01000..0x01fff	1
0x00000..0x00fff	0

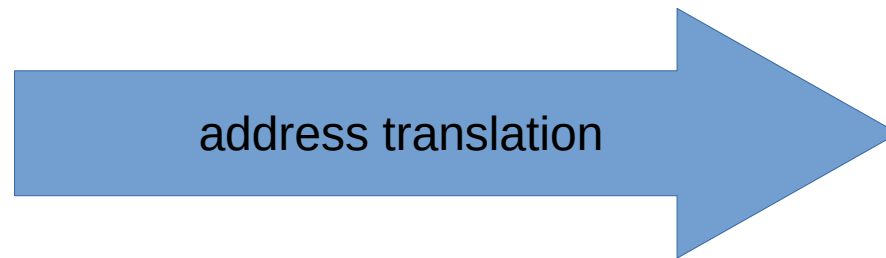


In other words:

`page_num = addr // page_size`

`page_offset = addr % page_size`

virtual address



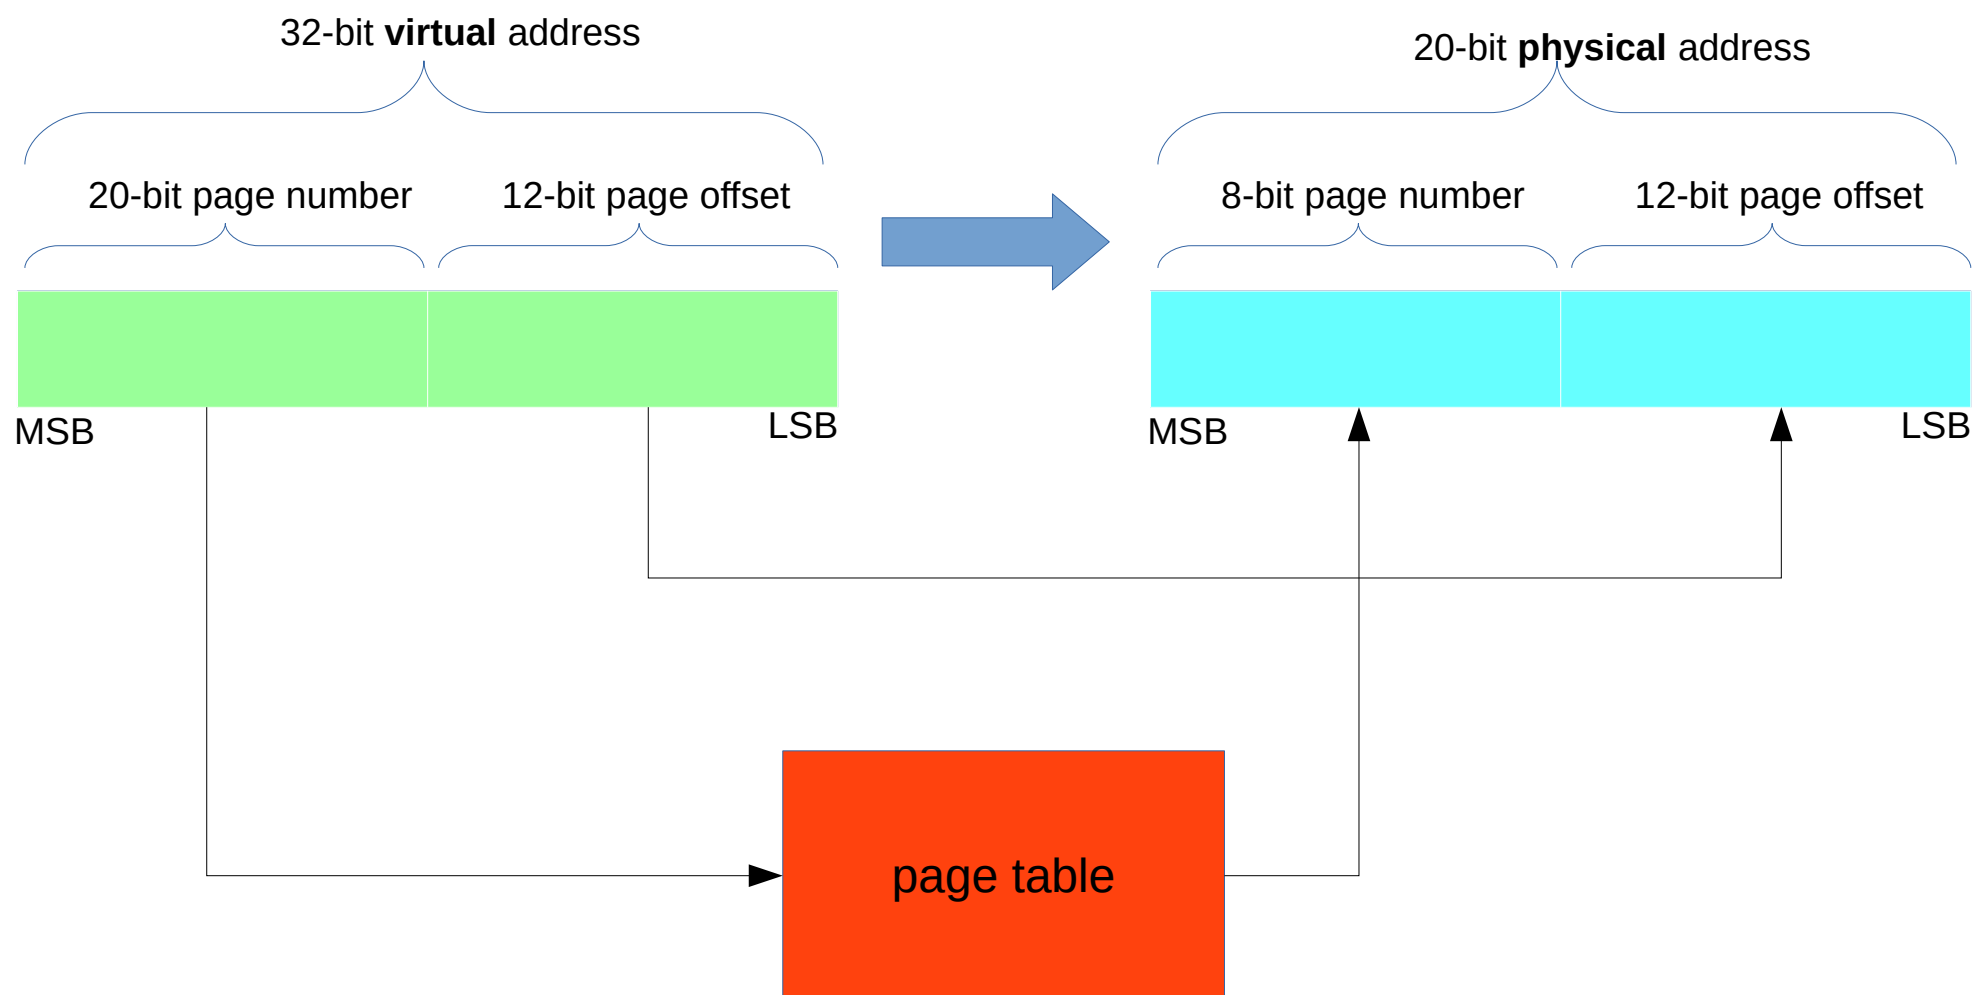
address translation

physical address

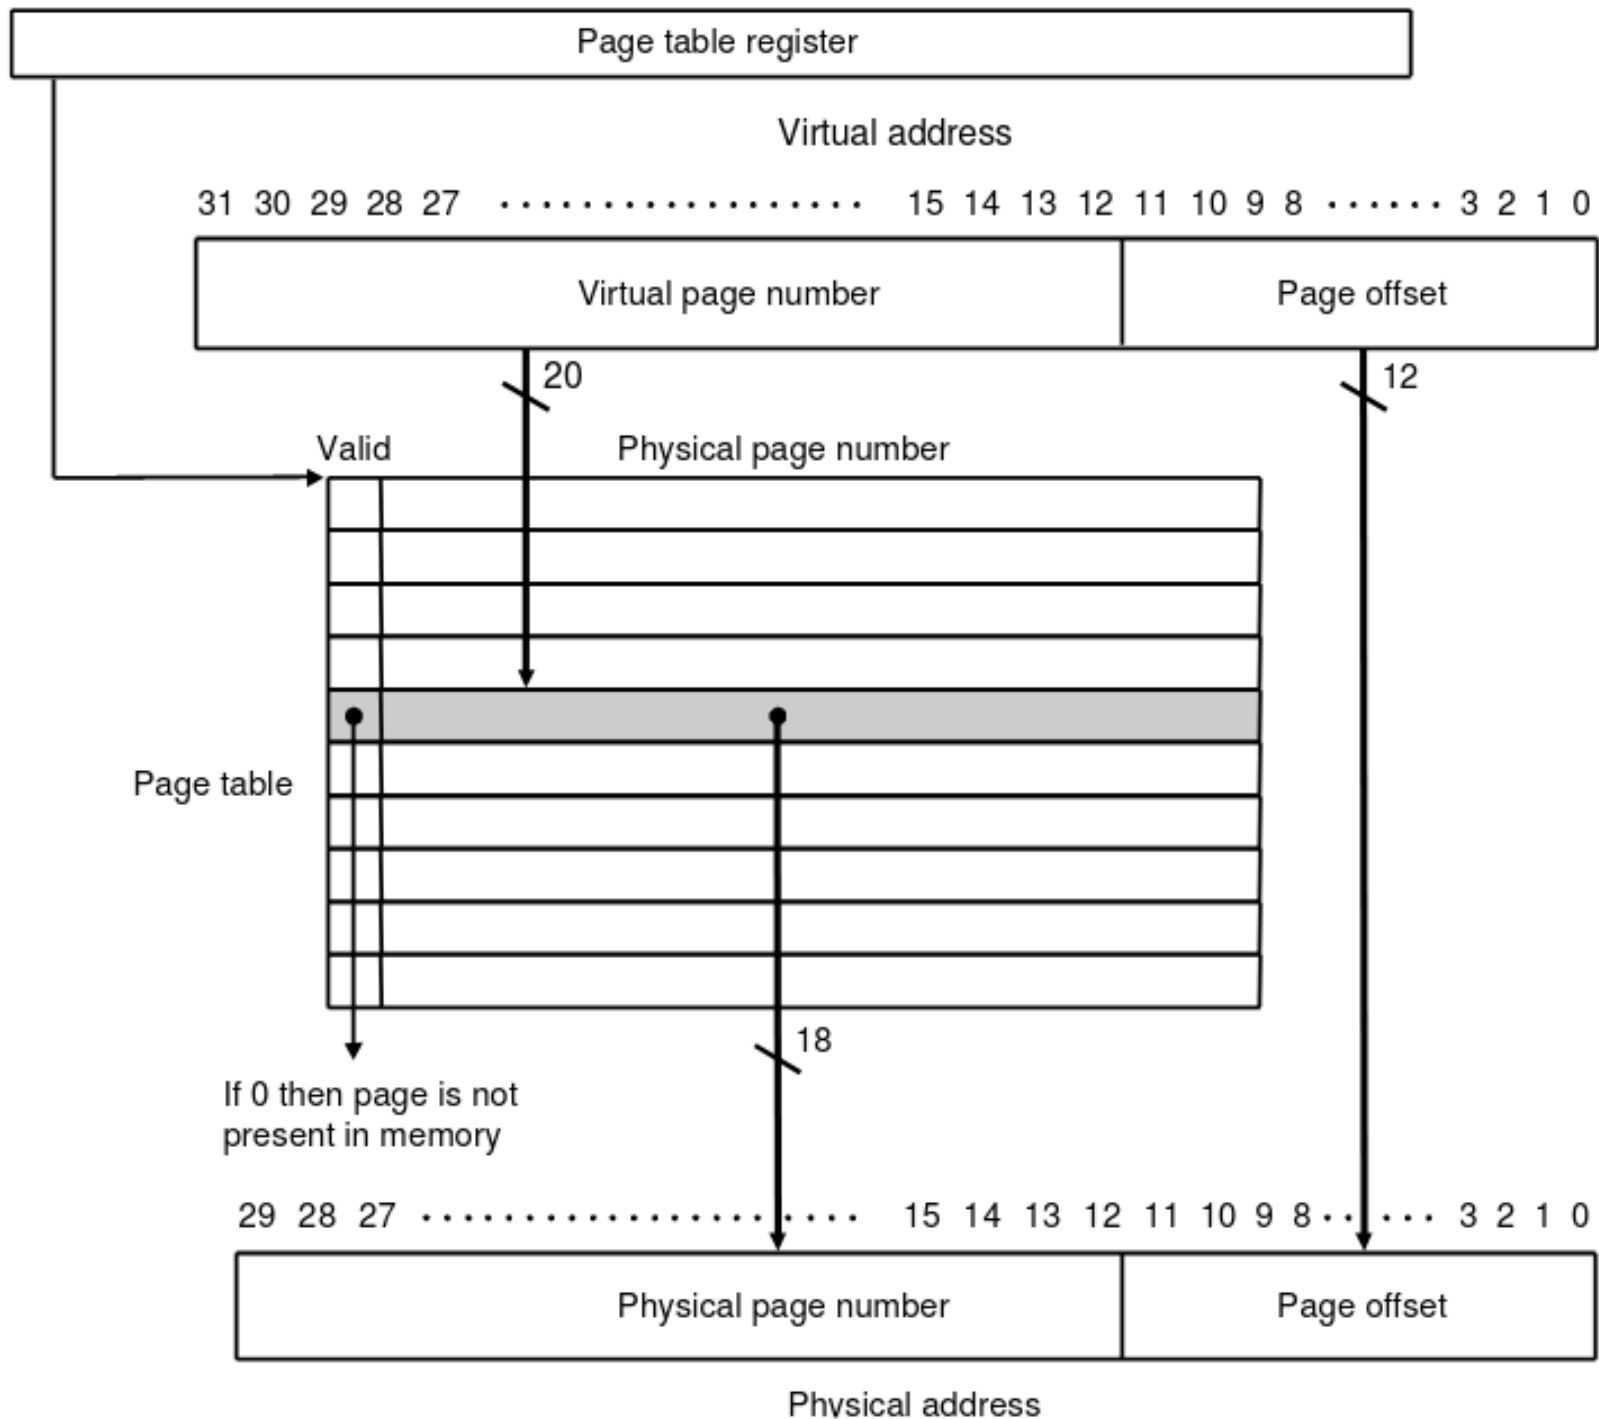


# Address translation

When translating a virtual address to a physical address, we are really translating a *virtual page* to *physical page*. The offset stays the same.



The *page table* is a data structure mapping virtual page numbers to physical page numbers.



# A page table

A page table is just an array. The index is the virtual page. Each virtual matches to a physical page.

Virtual Page	Physical Page
0	34
1	12
2	0
3	50
4	13
5	14
....	...

The page table is created and managed by the operating system.

Typically, each running program (process) will have its own page table, allowing each process to have its own address space.

There's usually a *special register* that holds the physical address of the page table. The CPU performs a page table lookup as part of its address translation on every memory access.

```
page_offset = virtual_addr % page_size
virtual_page = virtual_addr // page_size
physical_page = page_table[virtual_page]
physical_addr = physical_page * page_size + page_offset
```

# A page table

Caveat: not all virtual pages have a corresponding physical page. If an application accesses such a page, the hardware signals to the OS that the memory access cannot proceed: this is a *page fault*.

Thus we had a "valid" bit, which is checked before each lookup.

Virtual Page	V	Physical Page
0	1	34
1	0	
2	1	0
3	1	50
4	1	13
5	1	14
....		...

```
page_offset = virtual_addr % page_size
```

```
virtual_page = virtual_addr // page_size
```

```
physical_page = page_table[virtual_page]["physical_page"] if page_table[virtual_page]["v"]
```

```
physical_addr = physical_page * page_size + page_offset
```

What is the total size (in bits) of a page table for this system?

Assume a **32-bit virtual address**, consisting of a 20-bit virtual page number and a 12-bit offset.

Assume a **20-bit physical address**, consisting of an 8-bit physical page number and a 12-bit offset.

Assume the page table contains a physical page and valid bit for each entry.

What is the total size (in bits) of a page table for this system?

Assume a **32-bit virtual address**, consisting of a 20-bit virtual page number and a 12-bit offset.

Assume a **20-bit physical address**, consisting of an 8-bit physical page number and a 12-bit offset.

Assume the page table contains a physical page and valid bit for each entry.

How many entries are in the page table?

$$2^{**}20 = 1048576$$

What is the total size (in bits) of a page table for this system?

Assume a **32-bit virtual address**, consisting of a 20-bit virtual page number and a 12-bit offset.

Assume a **20-bit physical address**, consisting of an 8-bit physical page number and a 12-bit offset.

Assume the page table contains a physical page and valid bit for each entry.

How many entries are in the page table?

$$2^{**}20 = 1048576$$

How many bits is each entry?

1 valid bit + 8-bit physical page number

What is the total size (in bits) of a page table for this system?

Assume a **32-bit virtual address**, consisting of a 20-bit virtual page number and a 12-bit offset.

Assume a **20-bit physical address**, consisting of an 8-bit physical page number and a 12-bit offset.

Assume the page table contains a physical page and valid bit for each entry.

How many entries are in the page table?

$$2^{**}20 = 1048576$$

How many bits is each entry?


1 valid bit + 8-bit physical page number

What is the total size of the page table?

$$2^{**}20 * (1 + 8) = 9437184 \text{ bits}$$



# Virtual Memory

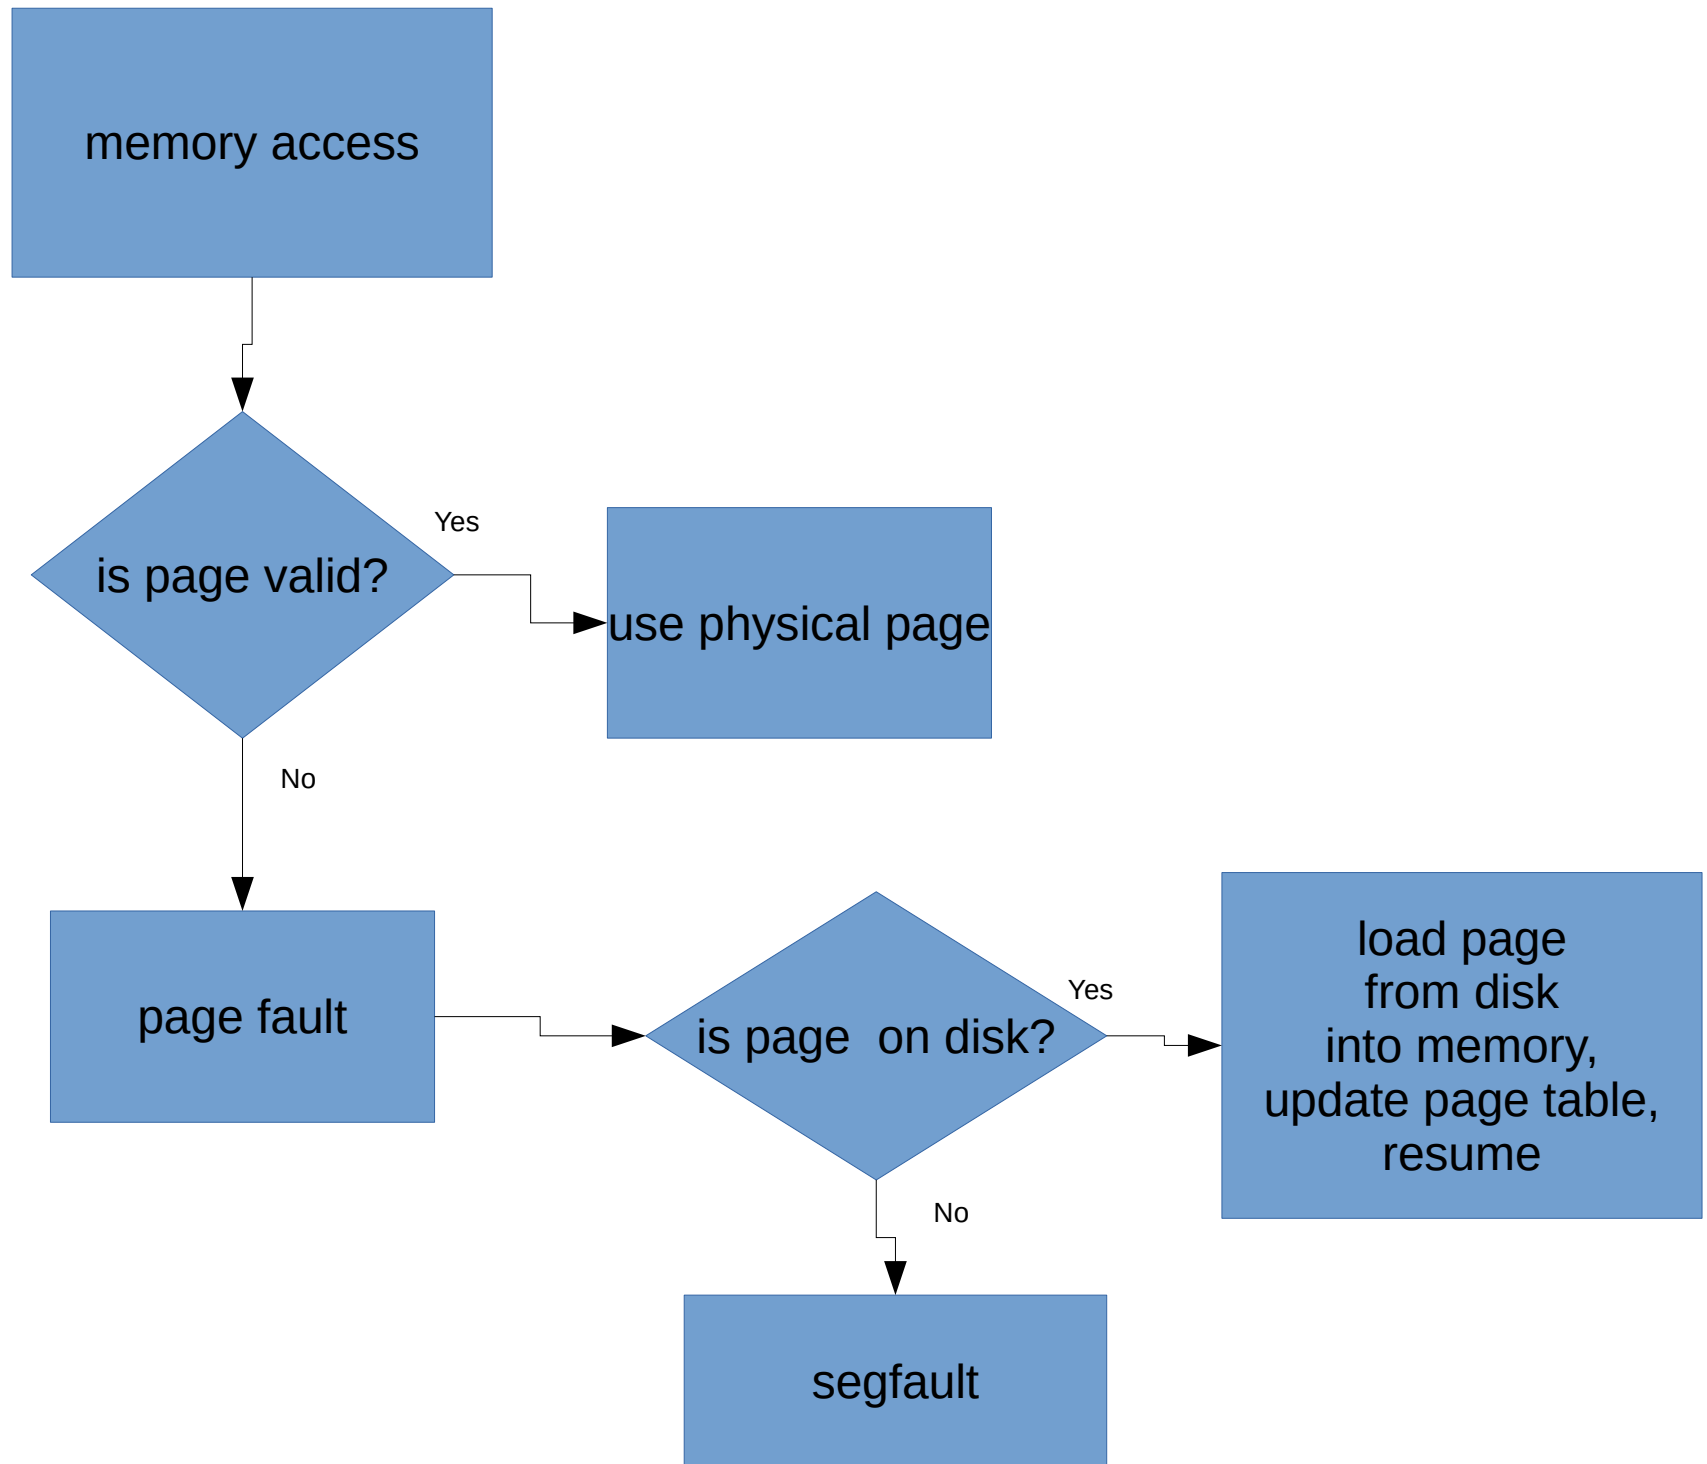


Anything does not fit in main memory is stored on disk

The swapping mechanism hopefully will:  
give us the **size** of the disk,  
with the **speed** of memory.

The design of virtual memory is based on the slowness of disk access.

We want to keep reduce misses.



# Swapping

# Swapping

The OS is responsible for allocating memory to applications.

If there is memory pressure (i.e. programs are using all available physical memory), what can the OS do?

Reject memory requests. Can't start programs, existing processes fail.

Or, the OS can choose to evict pages from memory and *swap* them to disk.

The corresponding page will be marked as invalid, so that when it's accessed, it will generate a page fault, whereupon it will be swapped back into memory.

The OS should prefer to swap seldom-used pages.

# A page table

When we evict a page from memory, we don't always have to save it to disk:

- The page may be empty or unused.
- The page may have been loaded from a file (e.g. an executable), so there's no reason to save it to disk again.
- The page may already be on disk, from when it was previously swapped out.

In none of the above cases apply, we have to write the page to disk.

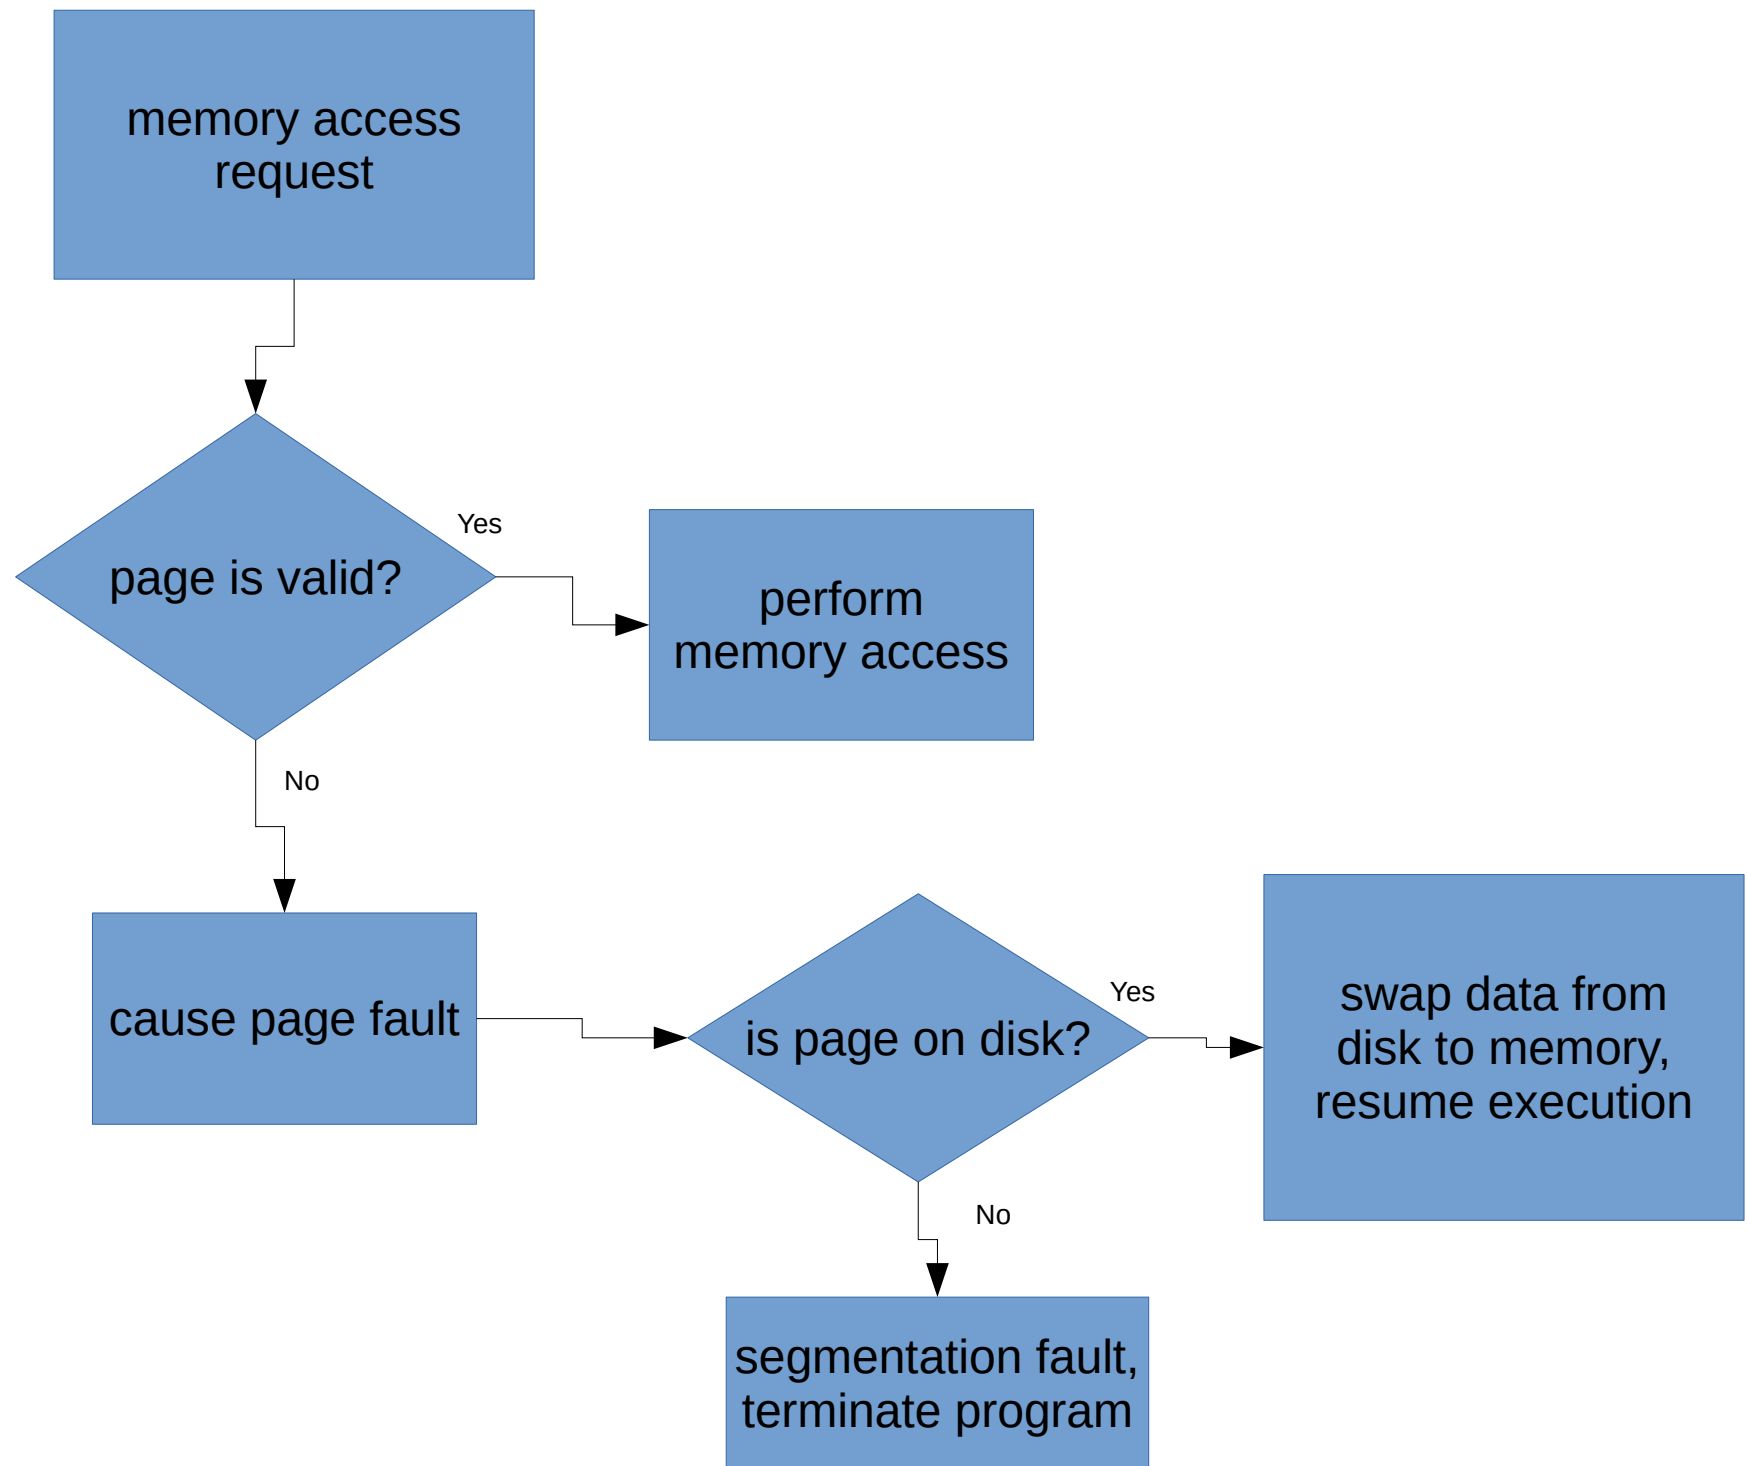
Virtual Page	V	Dirty	Physical Page
0	1	1	34
1	0		
2	1	0	0
3	1	1	50
4	1	0	13
5	1	0	14
....			...

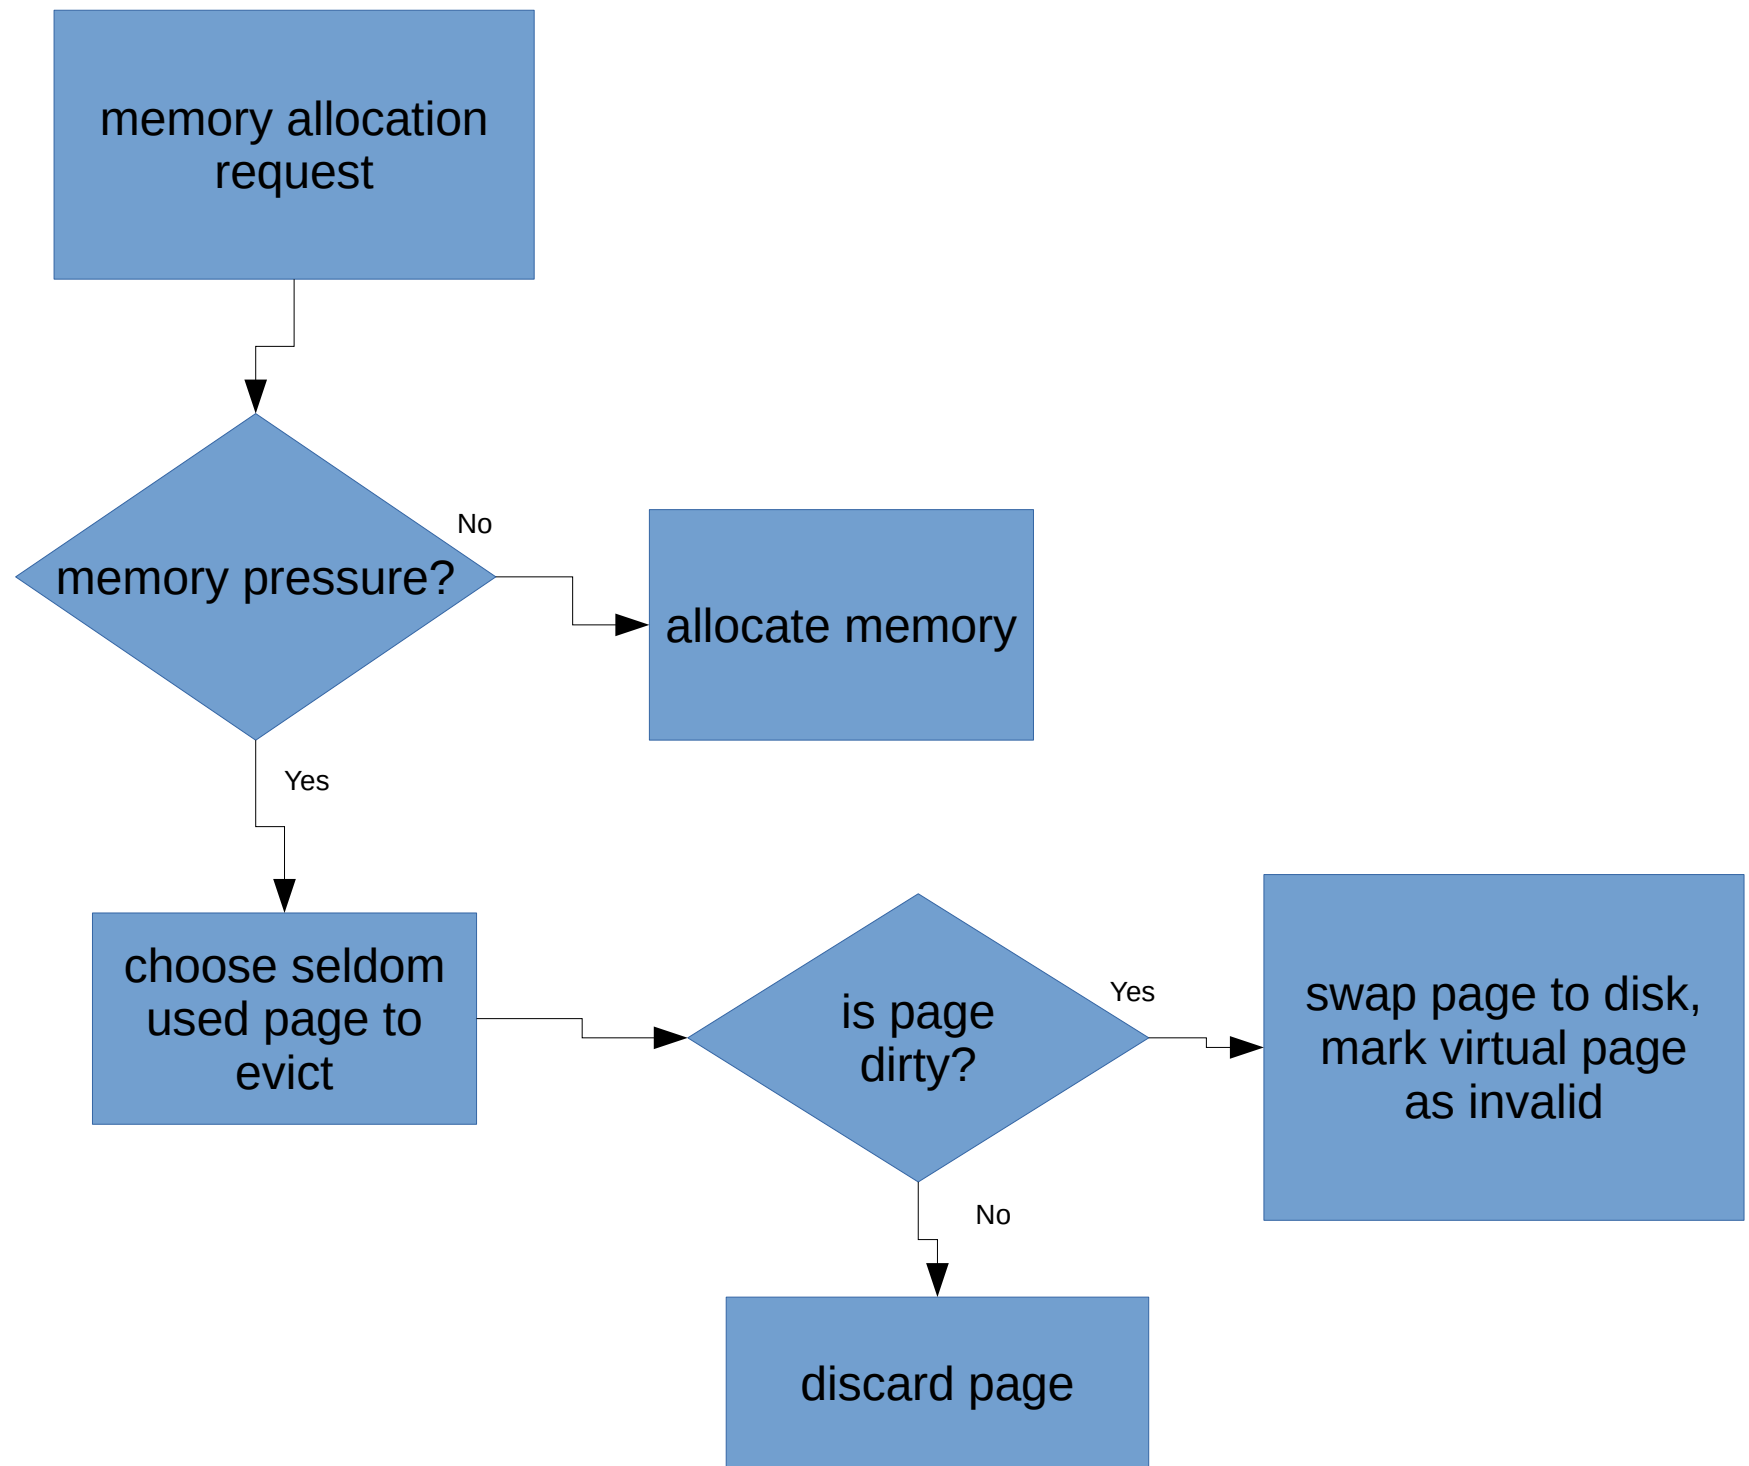
Every page initially has its dirty bit set to zero.

When a program does a write to that page, the processor sets the dirty bit to 1.

The OS will write the page to disk when swapping only if the dirty bit is 1.

This should all sound very familiar! With paging, *memory is a write-back cache of disk*.



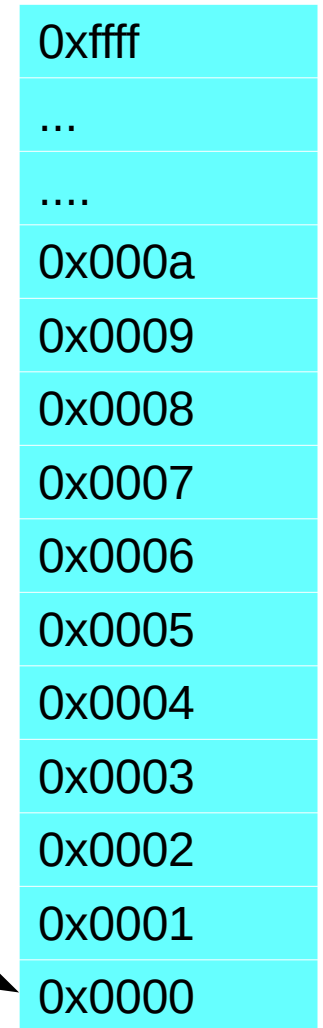






OS and applications may share address space:  
here, both address spaces overlap at physical  
page 0.

Reasons for this: the OS exposes code to the  
application for *system calls*.

What can go wrong?



physical memory

# A page table

To prevent the application from changing OS data, we introduce a *protect bit*. Pages marked with it cannot be written.

Virtual Page	V	Dirty	Protect	Physical Page
0	1	1	1	34
1	0			
2	1	0	0	0
3	1	1	0	50
4	1	0	0	13
5	1	0	0	14
....				...

# A page table

Assume an 8-byte page size. For each of the following address, given a corresponding physical page, or indicate if a page fault will occur. If the page is not writeable, indicate so.

43, 32, 10, 5, 30

Virtual Page	V	Dirty	Protect	Physical Page
0	1	1	1	34
1	0			
2	1	0	0	0
3	1	1	0	50
4	1	0	0	13
5	1	0	0	14
....				...

# A page table

Assume an 8-byte page size. For each of the following address, given a corresponding physical page, or indicate if a page fault will occur. If the page is not writeable, indicate so.

43, 32, 10, 5, 30

Virtual Page	V	Dirty	Protect	Physical Page
0	1	1	1	34
1	0			
2	1	0	0	0
3	1	1	0	50
4	1	0	0	13
5	1	0	0	14
....				...

virtual address 43. Page 5, offset 3. Physical address 115.

virtual address 32. Page 4, offset 0. Physical address 104.

virtual address 10. Page 1, offset 2. Page fault!

virtual address 5. Page 0, offset 5. Physical address 277. Read only.

virtual address 30. Page 3, offset 6. Physical address 406.

# Program walkthrough

- 1) When you run a program from an executable file, the content of the file is copied into memory. In this case, the program machine occupies virtual pages 0 and 1. They are valid, because they are in memory, and they are not dirty, because their content is directly copied from the file.

Virtual Page Num	Valid?	Dirty?
0	1	0
1	1	0
2	0	0
3	0	0
4	0	0
5	0	0
6	0	0

On disk:  
program.exe -> {0,1}

# Program walkthrough

- 1) When you run a program from an executable file, the content of the file is copied into memory. In this case, the program machine occupies virtual pages 0 and 1. They are valid, because they are in memory, and they are not dirty, because their content is directly copied from the file.
- 2) As the program runs, it requests additional memory from the operating system. The OS handles this by mapping additional pages 2 and 3 into its address space. They are now marked as valid. They are empty, and so are not dirty.

Virtual Page Num	Valid?	Dirty?
0	1	0
1	1	0
2	1	0
3	1	0
4	0	0
5	0	0
6	0	0

On disk:  
program.exe -> {0,1}

# Program walkthrough

- 1) When you run a program from an executable file, the content of the file is copied into memory. In this case, the program machine occupies virtual pages 0 and 1. They are valid, because they are in memory, and they are not dirty, because their content is directly copied from the file.
- 2) As the program runs, it requests additional memory from the operating system. The OS handles this by mapping additional pages 2 and 3 into its address space. They are now marked as valid. They are empty, and so are not dirty.
- 3) When the program writes to these new pages, they are marked dirty.

Virtual Page Num	Valid?	Dirty?
0	1	0
1	1	0
2	1	1
3	1	1
4	0	0
5	0	0
6	0	0

On disk:  
program.exe -> {0,1}

# Program walkthrough

- 1) When you run a program from an executable file, the content of the file is copied into memory. In this case, the program machine occupies virtual pages 0 and 1. They are valid, because they are in memory, and they are not dirty, because their content is directly copied from the file.
- 2) As the program runs, it requests additional memory from the operating system. The OS handles this by mapping additional pages 2 and 3 into its address space. They are now marked as valid. They are empty, and so are not dirty.
- 3) When the program writes to these new pages, they are marked dirty.
- 4) Now, some other application causes memory pressure, and the OS decides to swap out our program's pages 2 and 3 to disk. They are now marked invalid and not dirty.

Virtual Page Num	Valid?	Dirty?
0	1	0
1	1	0
2	0	0
3	0	0
4	0	0
5	0	0
6	0	0

On disk:  
program.exe -> {0,1}  
swapfile -> {2,3}



# Program walkthrough

- 1) When you run a program from an executable file, the content of the file is copied into memory. In this case, the program machine occupies virtual pages 0 and 1. They are valid, because they are in memory, and they are not dirty, because their content is directly copied from the file.
- 2) As the program runs, it requests additional memory from the operating system. The OS handles this by mapping additional pages 2 and 3 into its address space. They are now marked as valid. They are empty, and so are not dirty.
- 3) When the program writes to these new pages, they are marked dirty.
- 4) Now, some other application causes memory pressure, and the OS decides to swap out our program's pages 2 and 3 to disk. They are now marked invalid and not dirty.

Why would the OS swap out 2 and 3, instead of (non-dirty) 0 and 1?

It's possible the program was still running, resulting in reads to machine code in pages 0 and 1. Therefore pages 0 and 1 may be used more recently than 2 and 3.

Virtual Page Num	Valid?	Dirty?
0	1	0
1	1	0
2	0	0
3	0	0
4	0	0
5	0	0
6	0	0

On disk:  
program.exe -> {0,1}  
swapfile -> {2,3}

# Program walkthrough

- 1) When you run a program from an executable file, the content of the file is copied into memory. In this case, the program machine occupies virtual pages 0 and 1. They are valid, because they are in memory, and they are not dirty, because their content is directly copied from the file.
- 2) As the program runs, it requests additional memory from the operating system. The OS handles this by mapping additional pages 2 and 3 into its address space. They are now marked as valid. They are empty, and so are not dirty.
- 3) When the program writes to these new pages, they are marked dirty.
- 4) Now, some other application causes memory pressure, and the OS decides to swap out our program's pages 2 and 3 to disk. They are now marked invalid and not dirty.
- 5) When our program tries to access those pages, it will cause a page fault. The OS will swap them back in and mark them valid. They are still marked not dirty, because the copy on disk is still present.

Virtual Page Num	Valid?	Dirty?
0	1	0
1	1	0
2	1	0
3	1	0
4	0	0
5	0	0
6	0	0

On disk:  
program.exe -> {0,1}  
swapfile -> {2,3}

# Program walkthrough

- 1) When you run a program from an executable file, the content of the file is copied into memory. In this case, the program machine occupies virtual pages 0 and 1. They are valid, because they are in memory, and they are not dirty, because their content is directly copied from the file.
- 2) As the program runs, it requests additional memory from the operating system. The OS handles this by mapping additional pages 2 and 3 into its address space. They are now marked as valid. They are empty, and so are not dirty.
- 3) When the program writes to these new pages, they are marked dirty.
- 4) Now, some other application causes memory pressure, and the OS decides to swap out our program's pages 2 and 3 to disk. They are now marked invalid and not dirty.
- 5) When our program tries to access those pages, it will cause a page fault. The OS will swap them back in and mark them valid. They are still marked not dirty, because the copy on disk is still present.
- 6) If the OS decides again detects memory pressure, it is safe to simply discard these pages, so copies already exist on disk.

Virtual Page Num	Valid?	Dirty?
0	1	0
1	1	0
2	0	0
3	0	0
4	0	0
5	0	0
6	0	0

On disk:  
program.exe -> {0,1}  
swapfile -> {2,3}

## How big is a page table?

Consider a system with 32-bit virtual addresses, 32-bit physical addresses, and 4KB pages. How big will the page tables be (including metadata)? Assume each memory cell is one byte.

# How big is a page table?

Consider a system with 32-bit virtual addresses, 32-bit physical addresses, and 4KB pages. How big will the page tables be (including metadata)? Assume each memory cell is one byte.

There will be  $2^{20}$  page table entries.

Each entry will contain the 20-bit physical page, plus the dirty, protect, and valid bits.

So the total size is  $2^{20} * (20 + 3) = 24,117,248$  bits, or 3,014,656 bytes.

Each process will have a page table. That adds up!

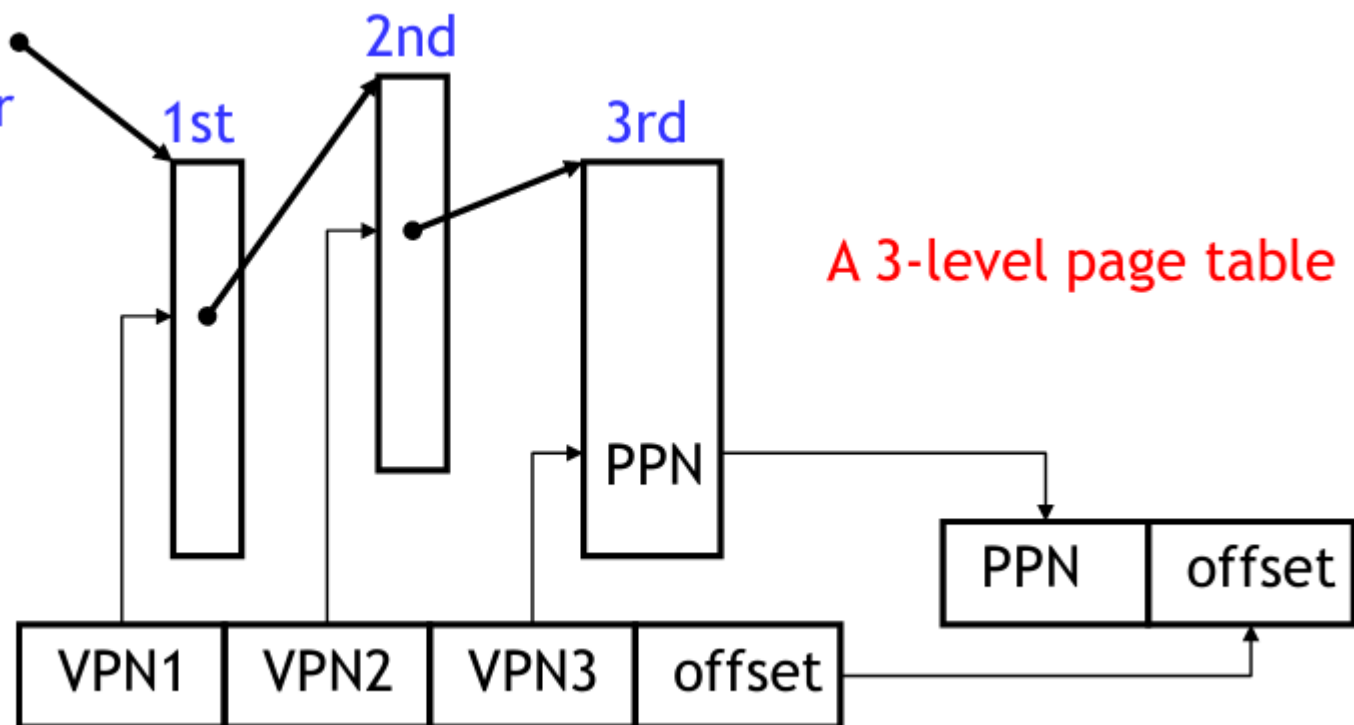
Further observation: most programs use only a small part of their address space. Most page tables are mostly empty.

Another observation: with a page table, every memory access by an application entails *two* (2) round-trip accesses to memory: first to translate the address in the page table, then to access the actual address. So we've just cut the speed of our memory unit in half.

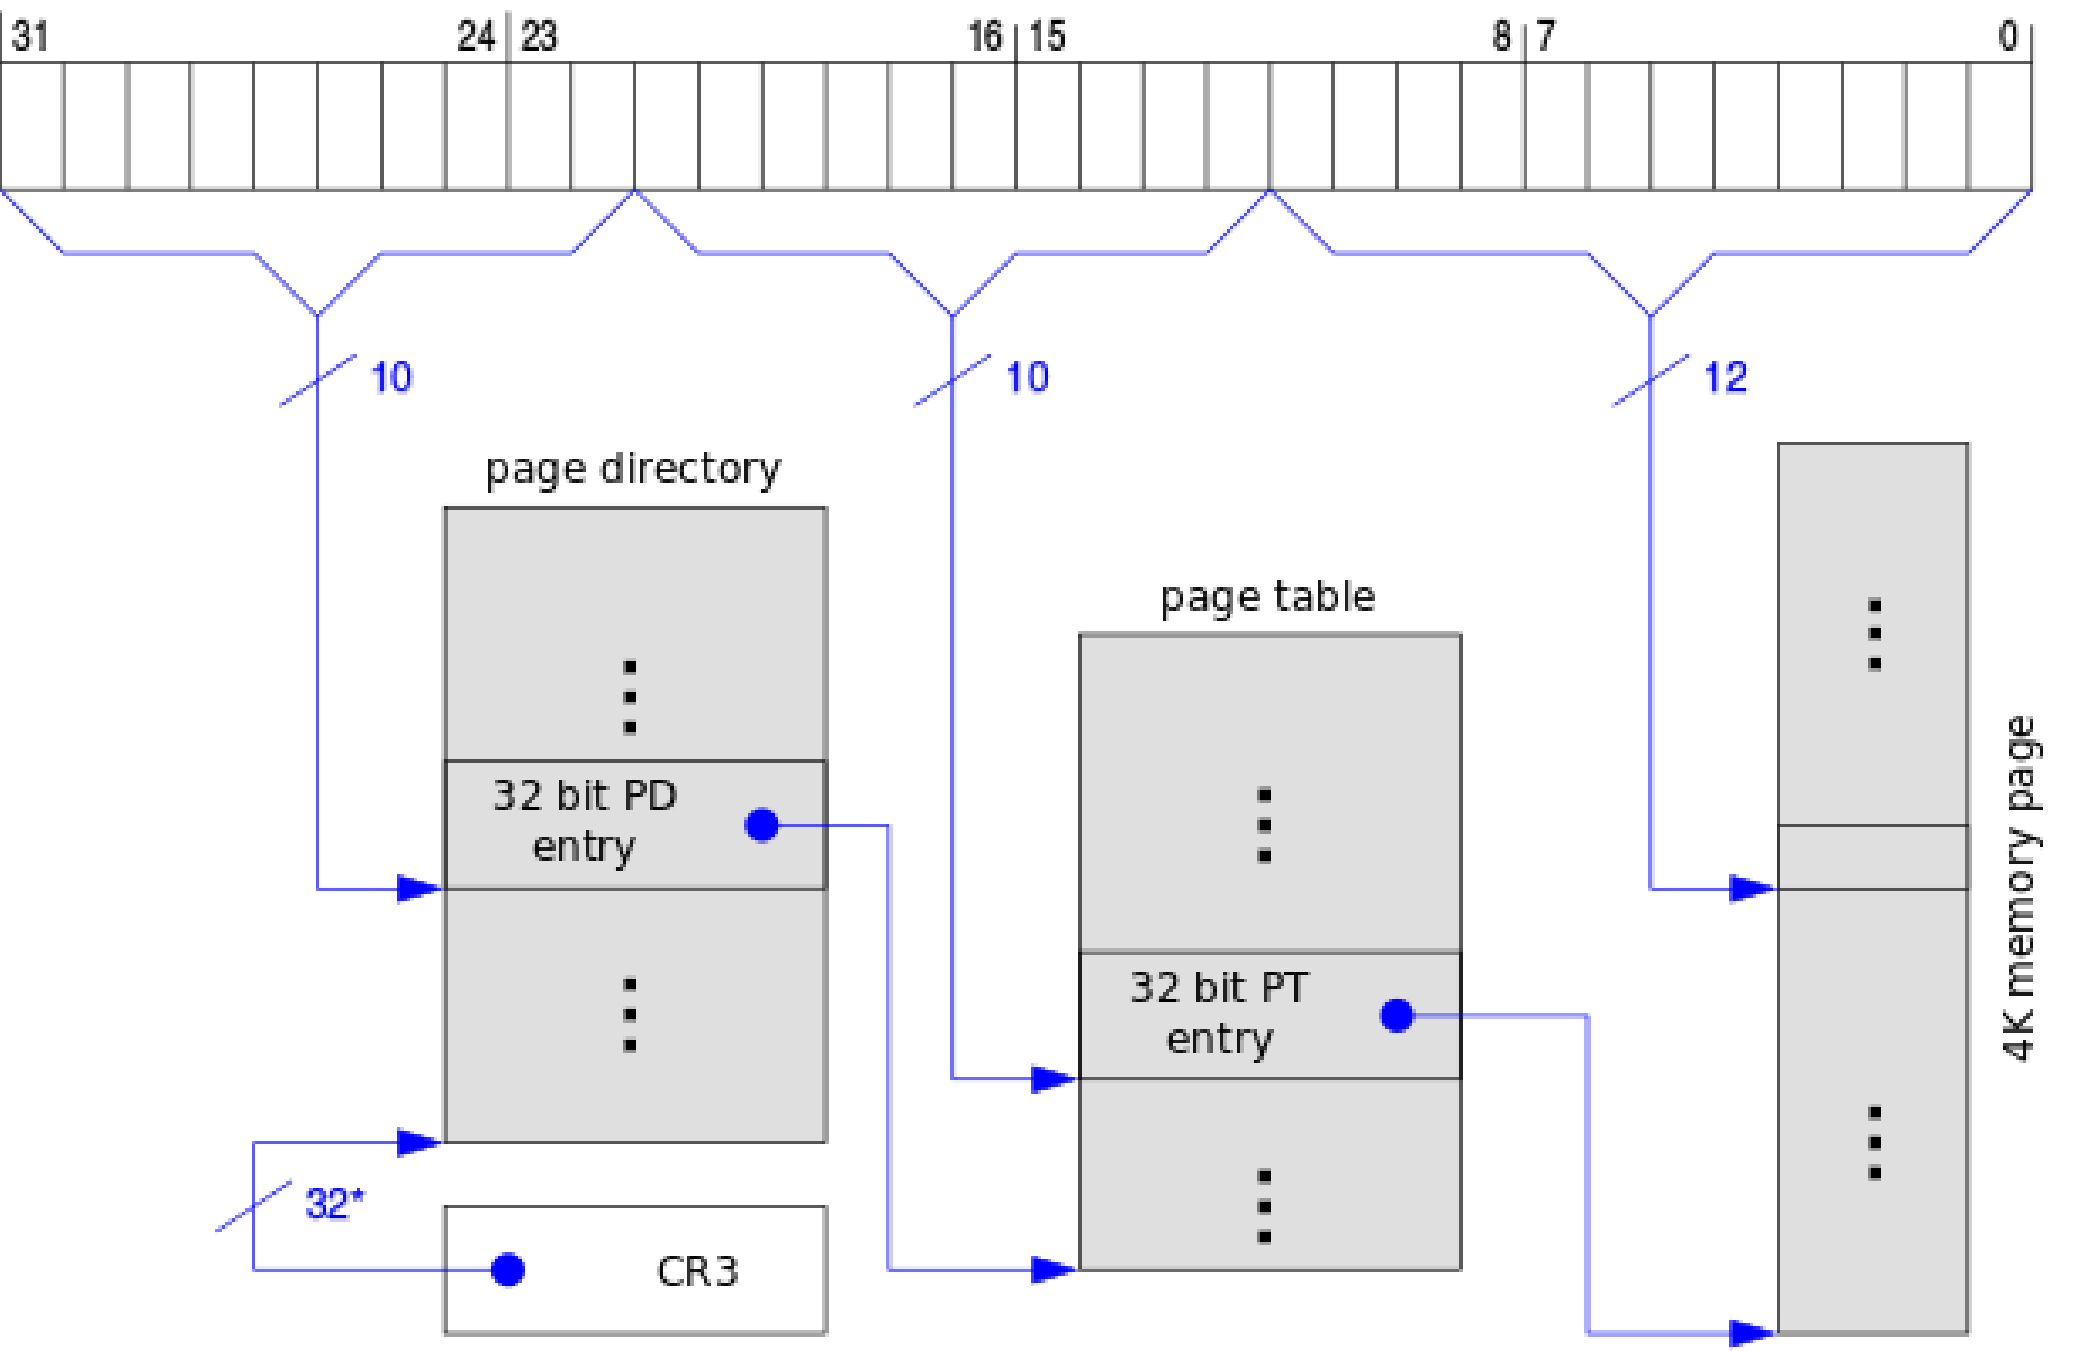
# Page tables

- Solutions to too-big page table
  - We swap it out, of course
  - The location of the page table is stored by another page table
  - We thus have a hierarchy of page tables, much as we had a hierarchy of caches
- Solutions to too-slow page table
  - We cache it, of course
  - Virtual-to-physical page translations are cached in the Translation Lookaside Buffer (TLB)

Page Table  
Base Pointer



Linear address:

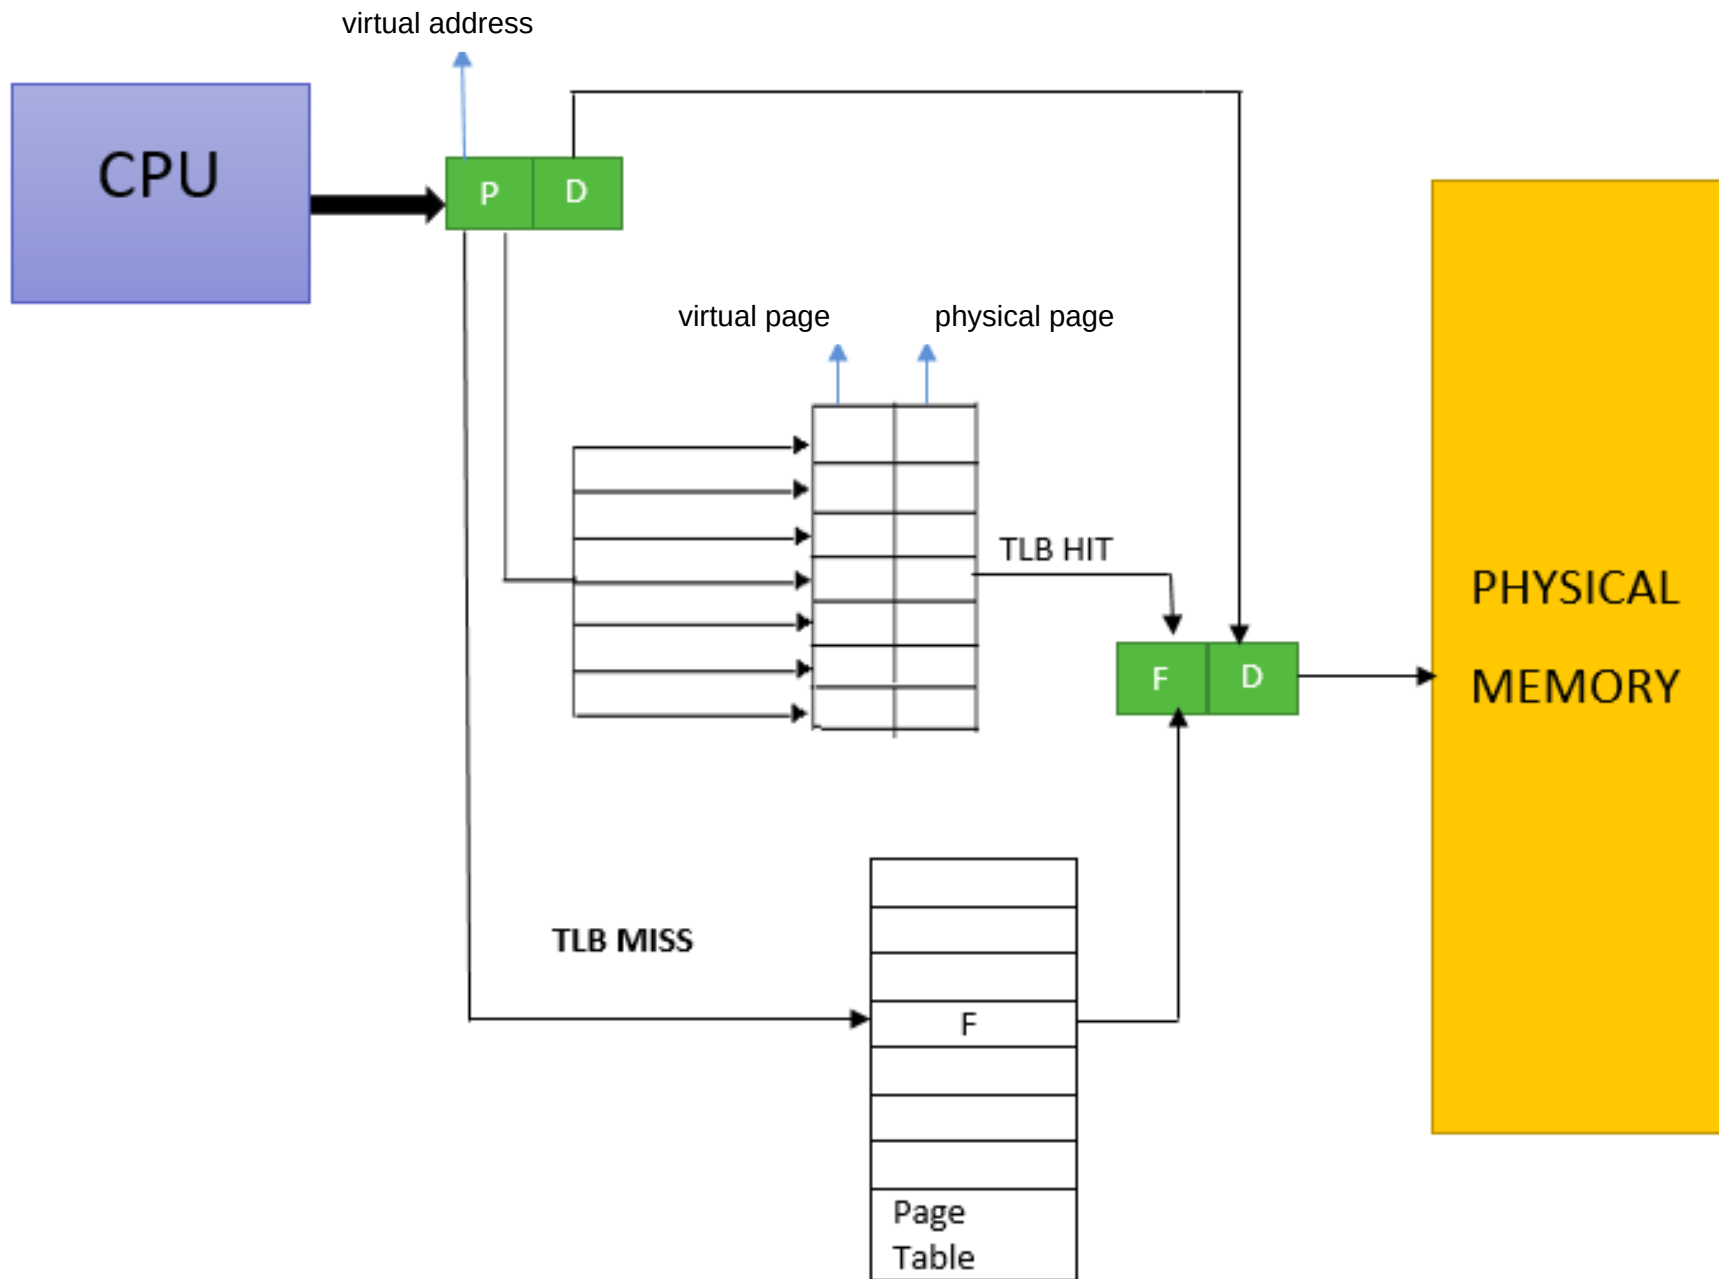


\*) 32 bits aligned to a 4-KByte boundary

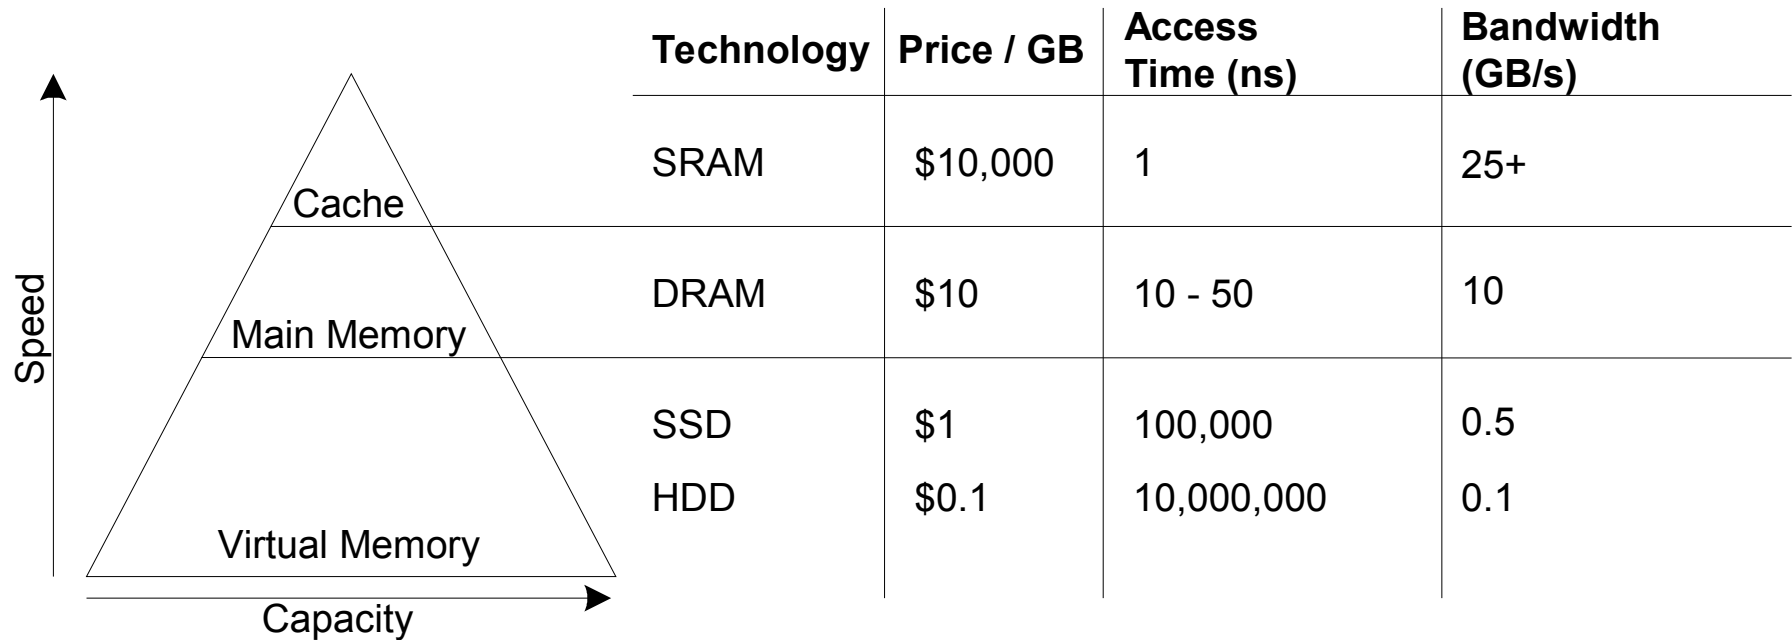


# TLB

- The TLB (Translation Lookaside Buffer) is a cache of the page table
- It maps virtual addresses to physical addresses
  - The tag is the virtual page
  - The value stored is the physical page
  - Also has valid, dirty, and protect bits
  - TLB is consulted before checking the page table
- It's usually a fully-associative cache.



# Memory Hierarchy



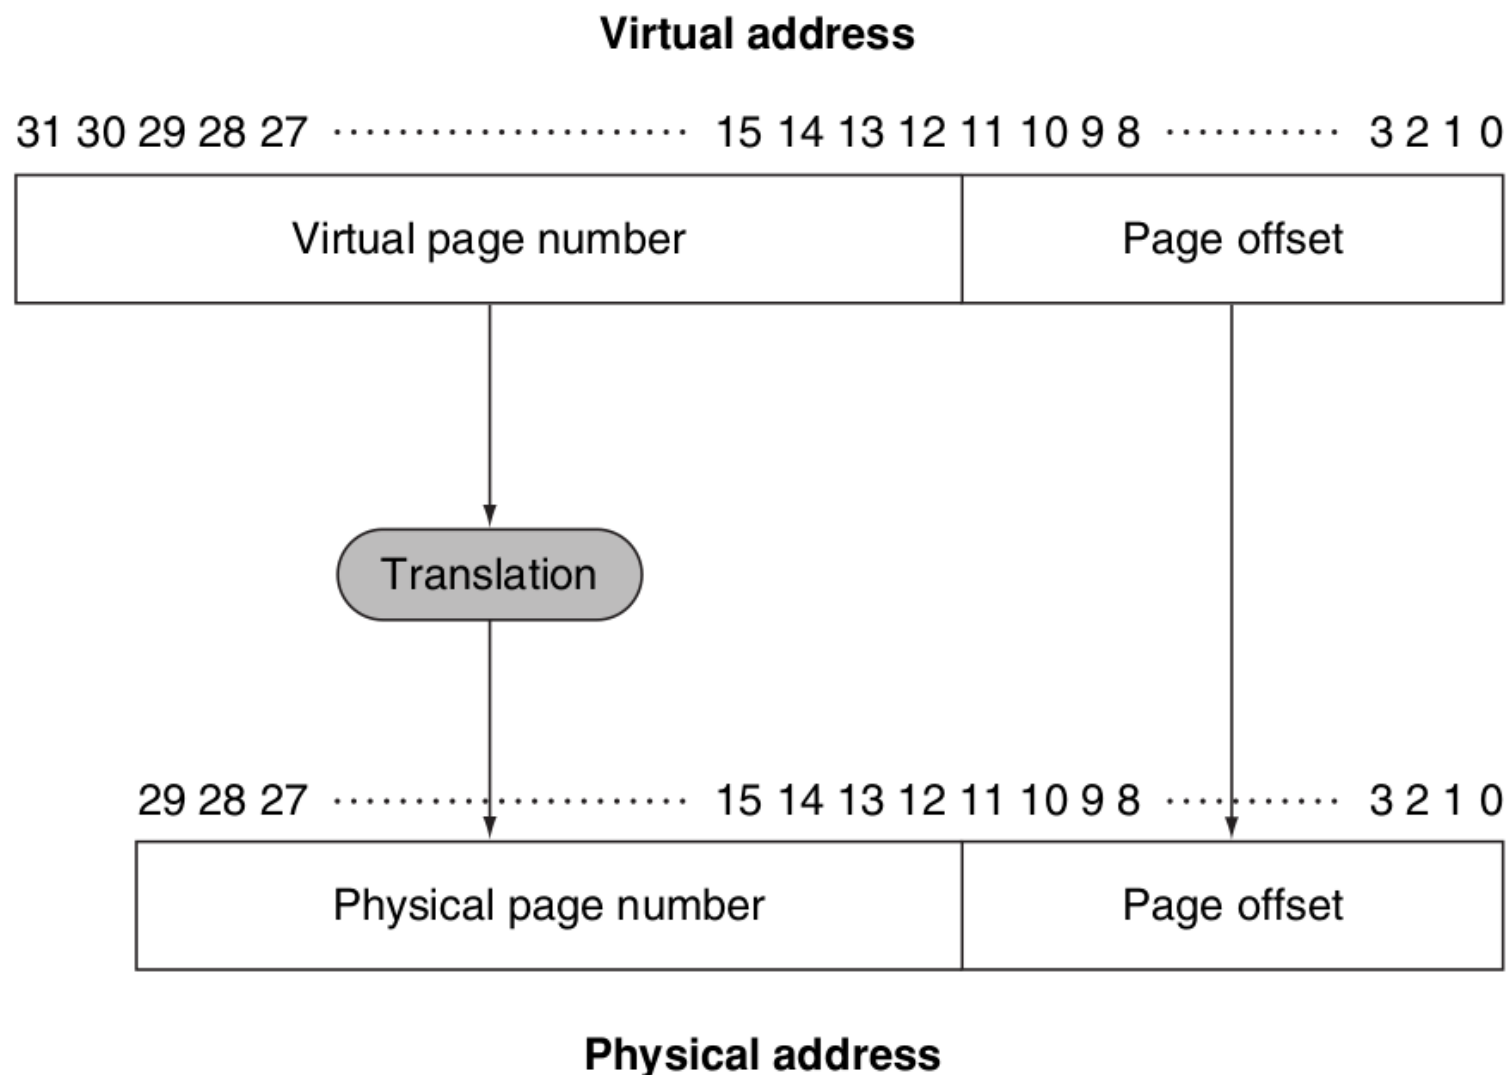
# Virtual Memory Example

- **System:**

- Virtual memory size: 4 GB =  $2^{32}$  bytes
- Physical memory size: 1 GB =  $2^{30}$  bytes
- Page size: 4 KB =  $2^{12}$  bytes

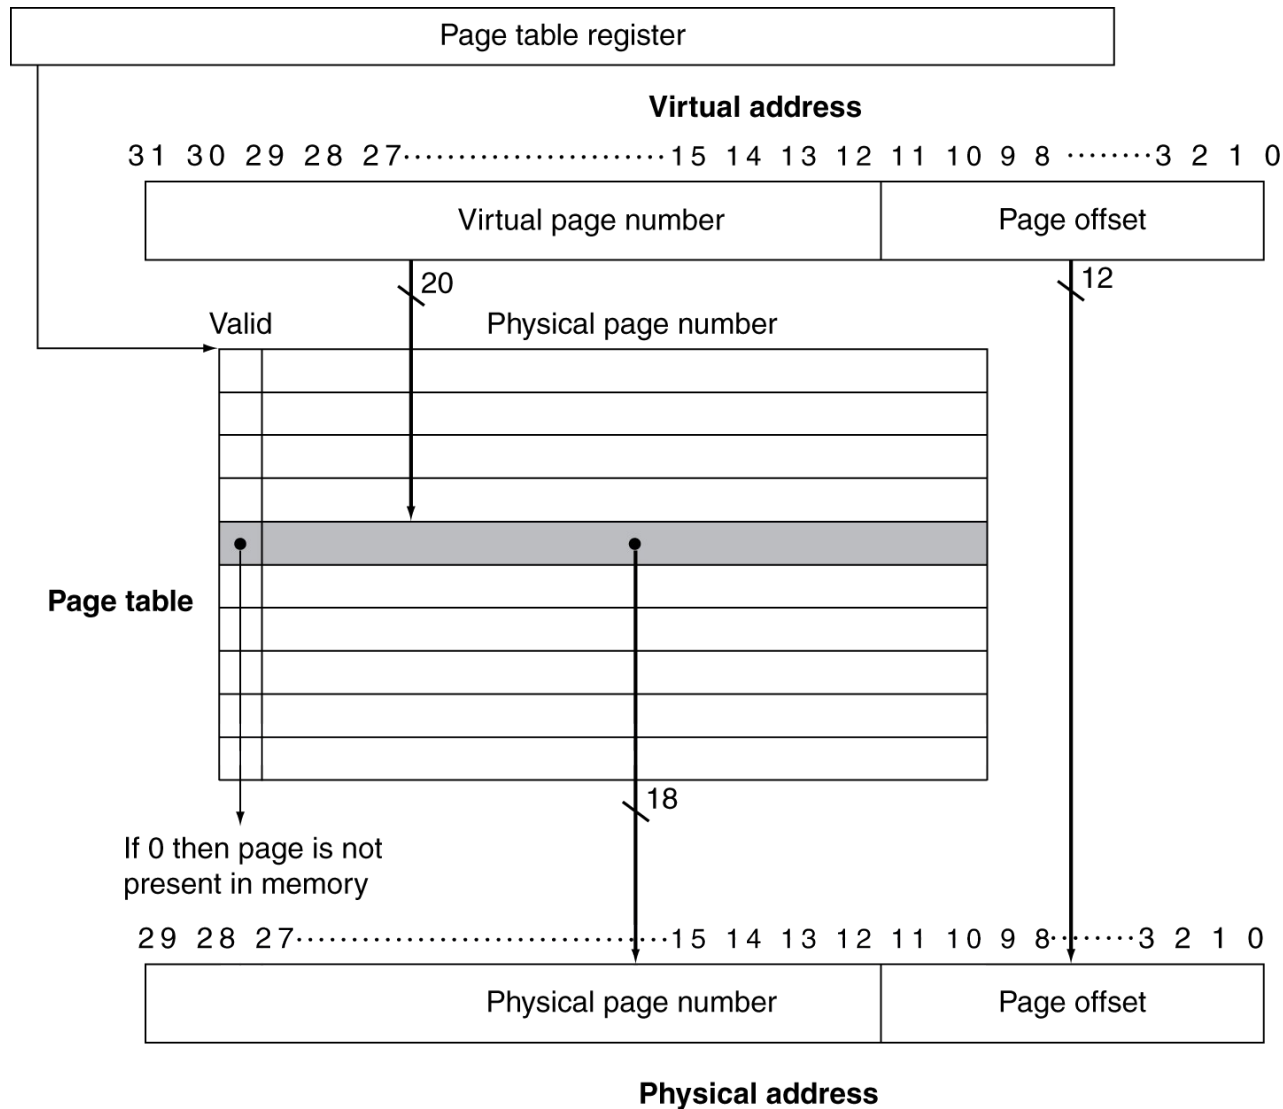
- **Organization:**

- Virtual address: 32 bits
- Physical address: 30 bits
- Page offset: 12 bits
- # Virtual pages =  $2^{32}/2^{12} = 2^{20}$  (VPN = 20 bits)
- # Physical pages =  $2^{30}/2^{12} = 2^{18}$  (PPN = 18 bits)

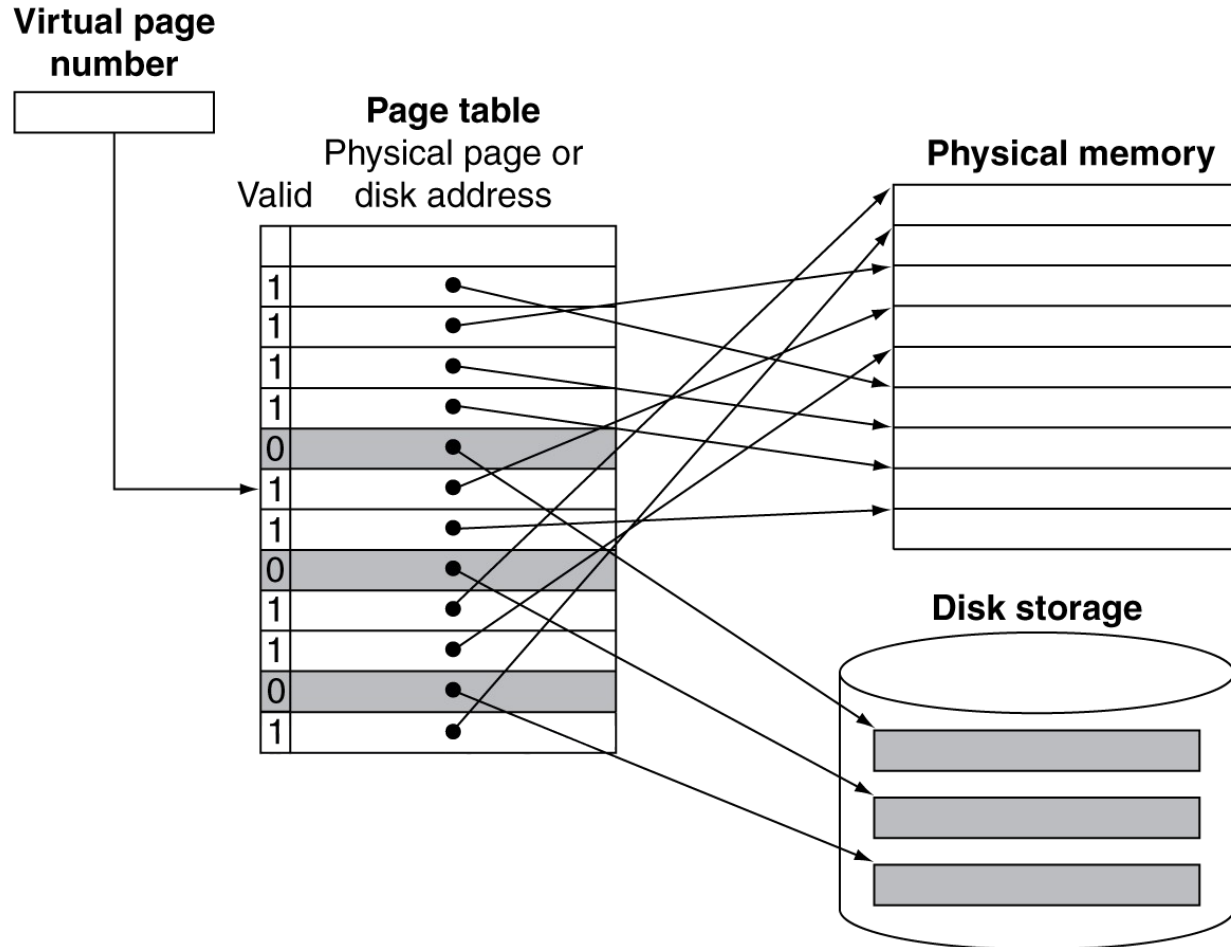


**FIGURE 5.26 Mapping from a virtual to a physical address.** The page size is  $2^{12} = 4$  KiB. The number of physical pages allowed in memory is  $2^{18}$ , since the physical page number has 18 bits in it. Thus, main memory can have at most 1 GiB, while the virtual address space is 4 GiB.

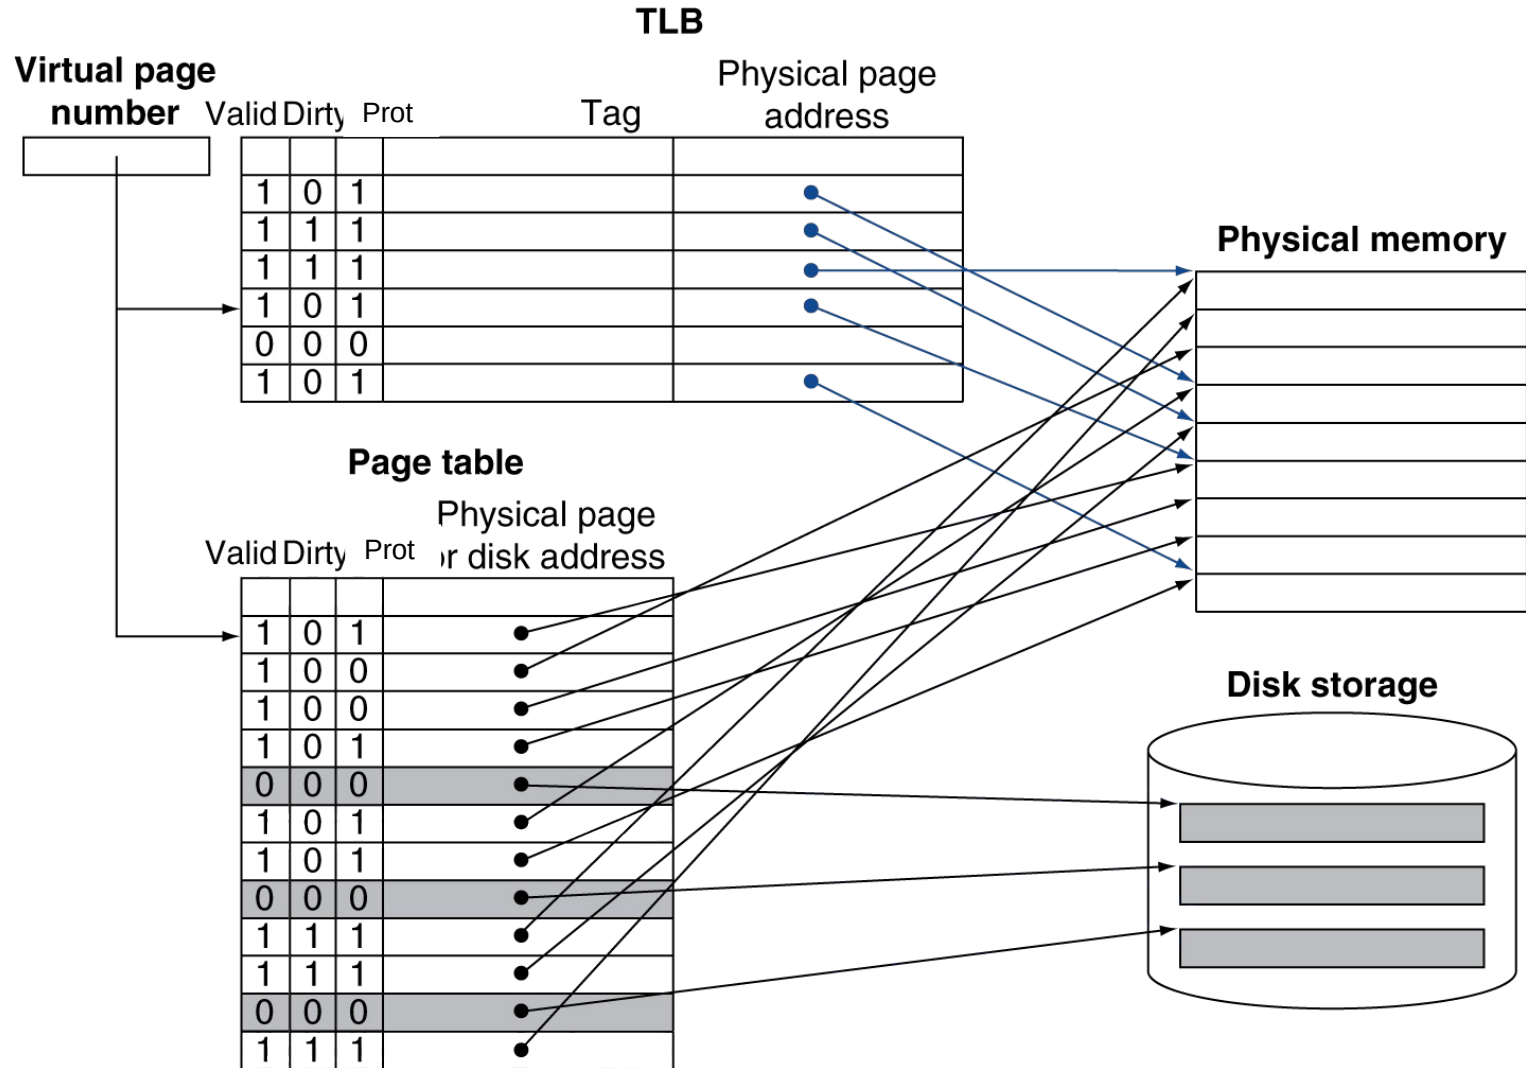
# Translation Using a Page Table



# Mapping Pages to Storage



# Fast Translation Using a TLB





# Hierarchical page tables


Modern 64-bit Intel-based systems have 4 or 5 levels of page tables.

Each layer has a 9-bit index.

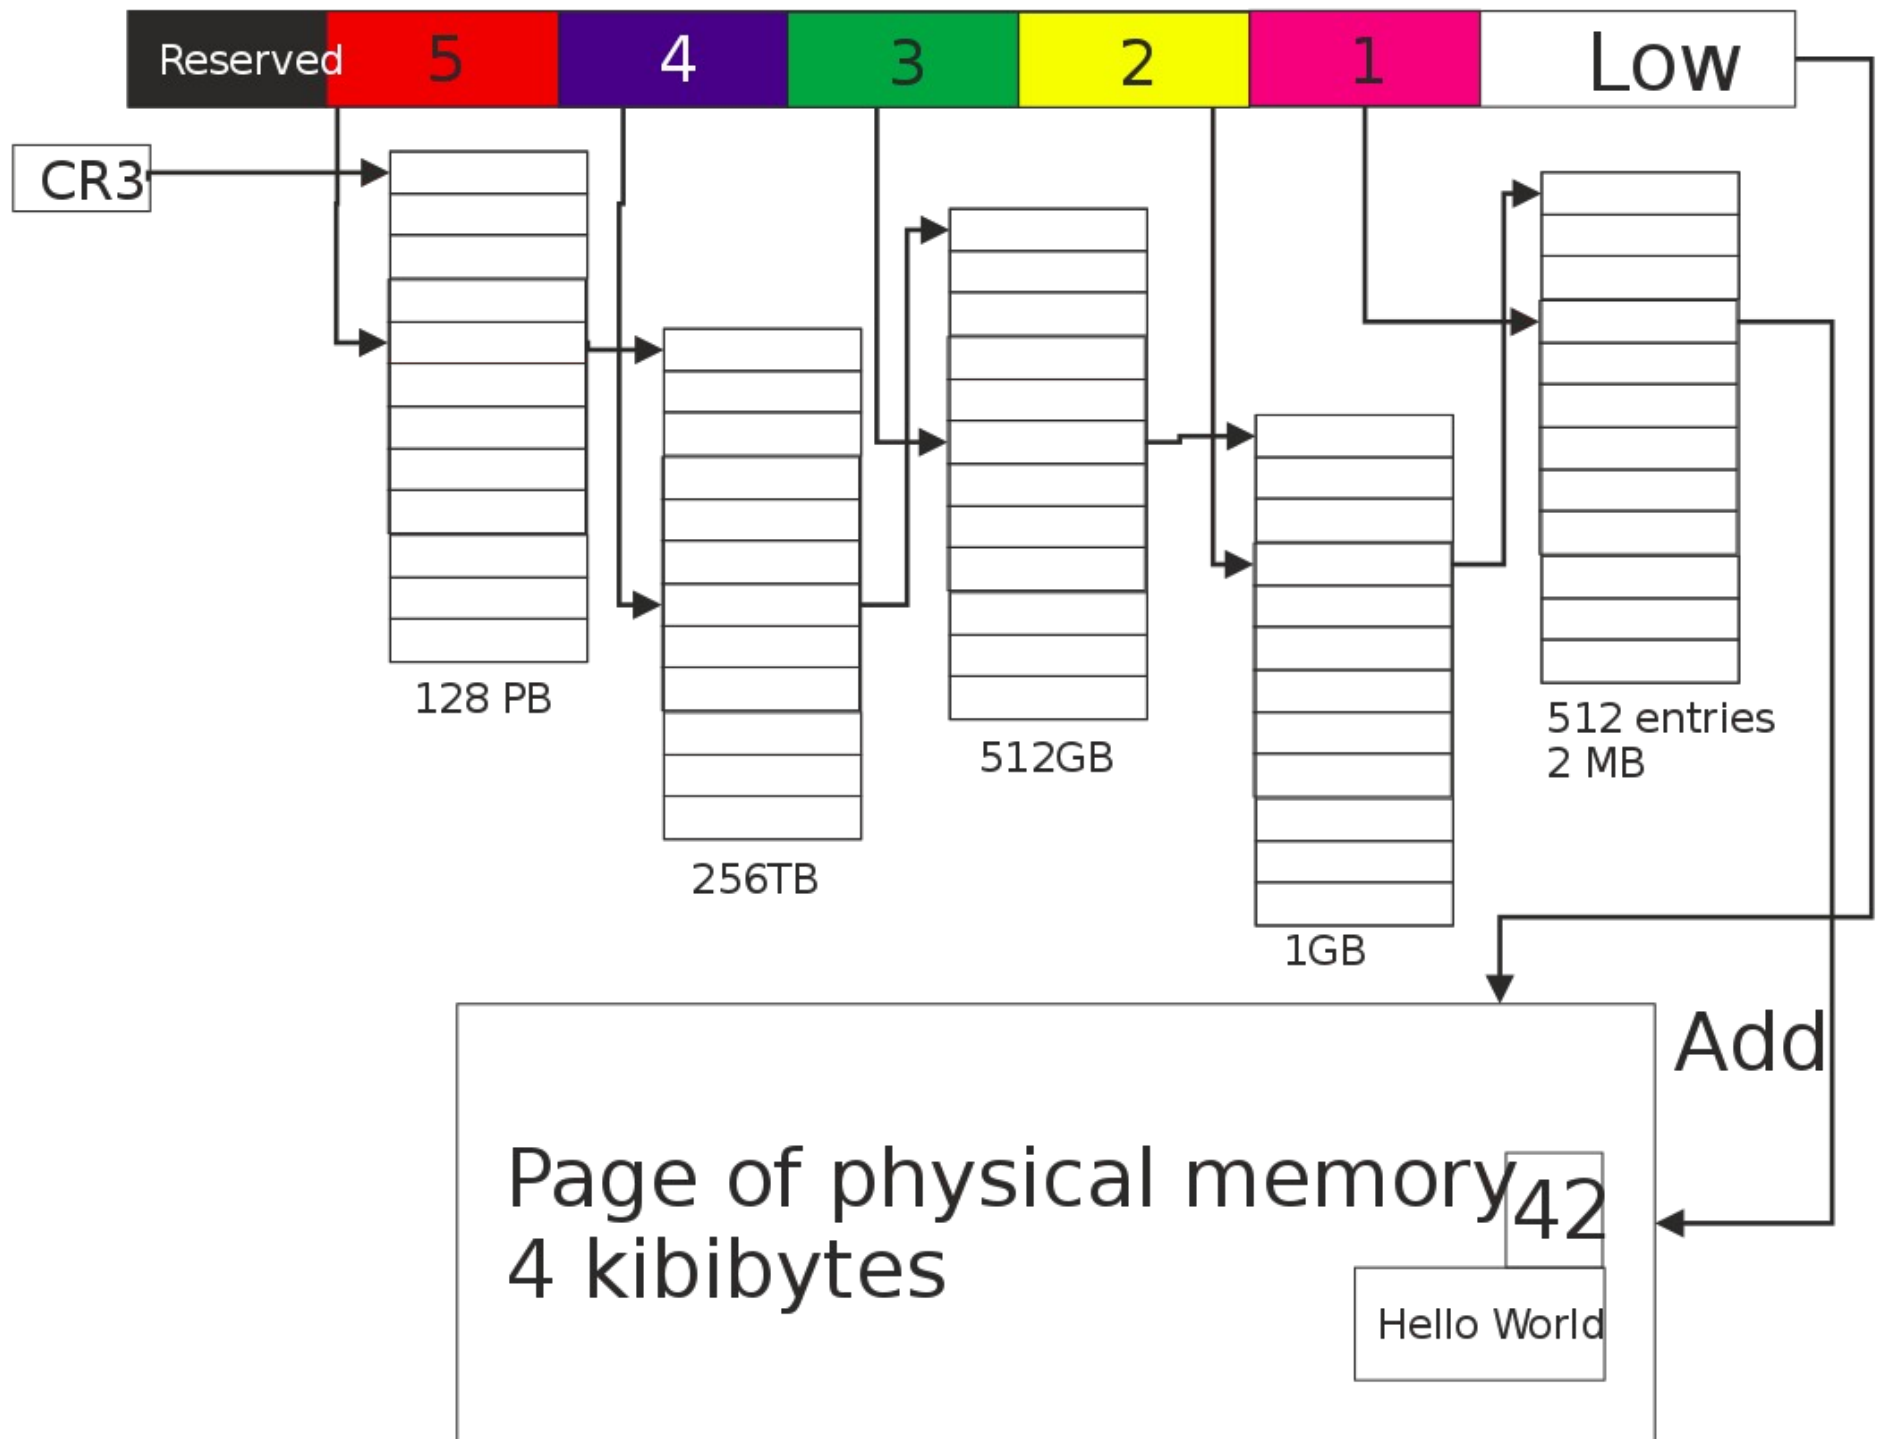
4-level page tables allow virtual address of 48 bits, address space of 256 TiB.

5-level page tables allow virtual address of 57 bits, address space of 128 PiB.

Typically, low virtual address are for userland, high virtual addresses are reserved for kernel.



# Virtual Address



Note that the most significant bits of virtual addresses are unused. Rather than simply be zero, they are *sign extended*, that is, they must all have the same value as the most significant used bit of the virtual address.